

CS61c Fall 2021

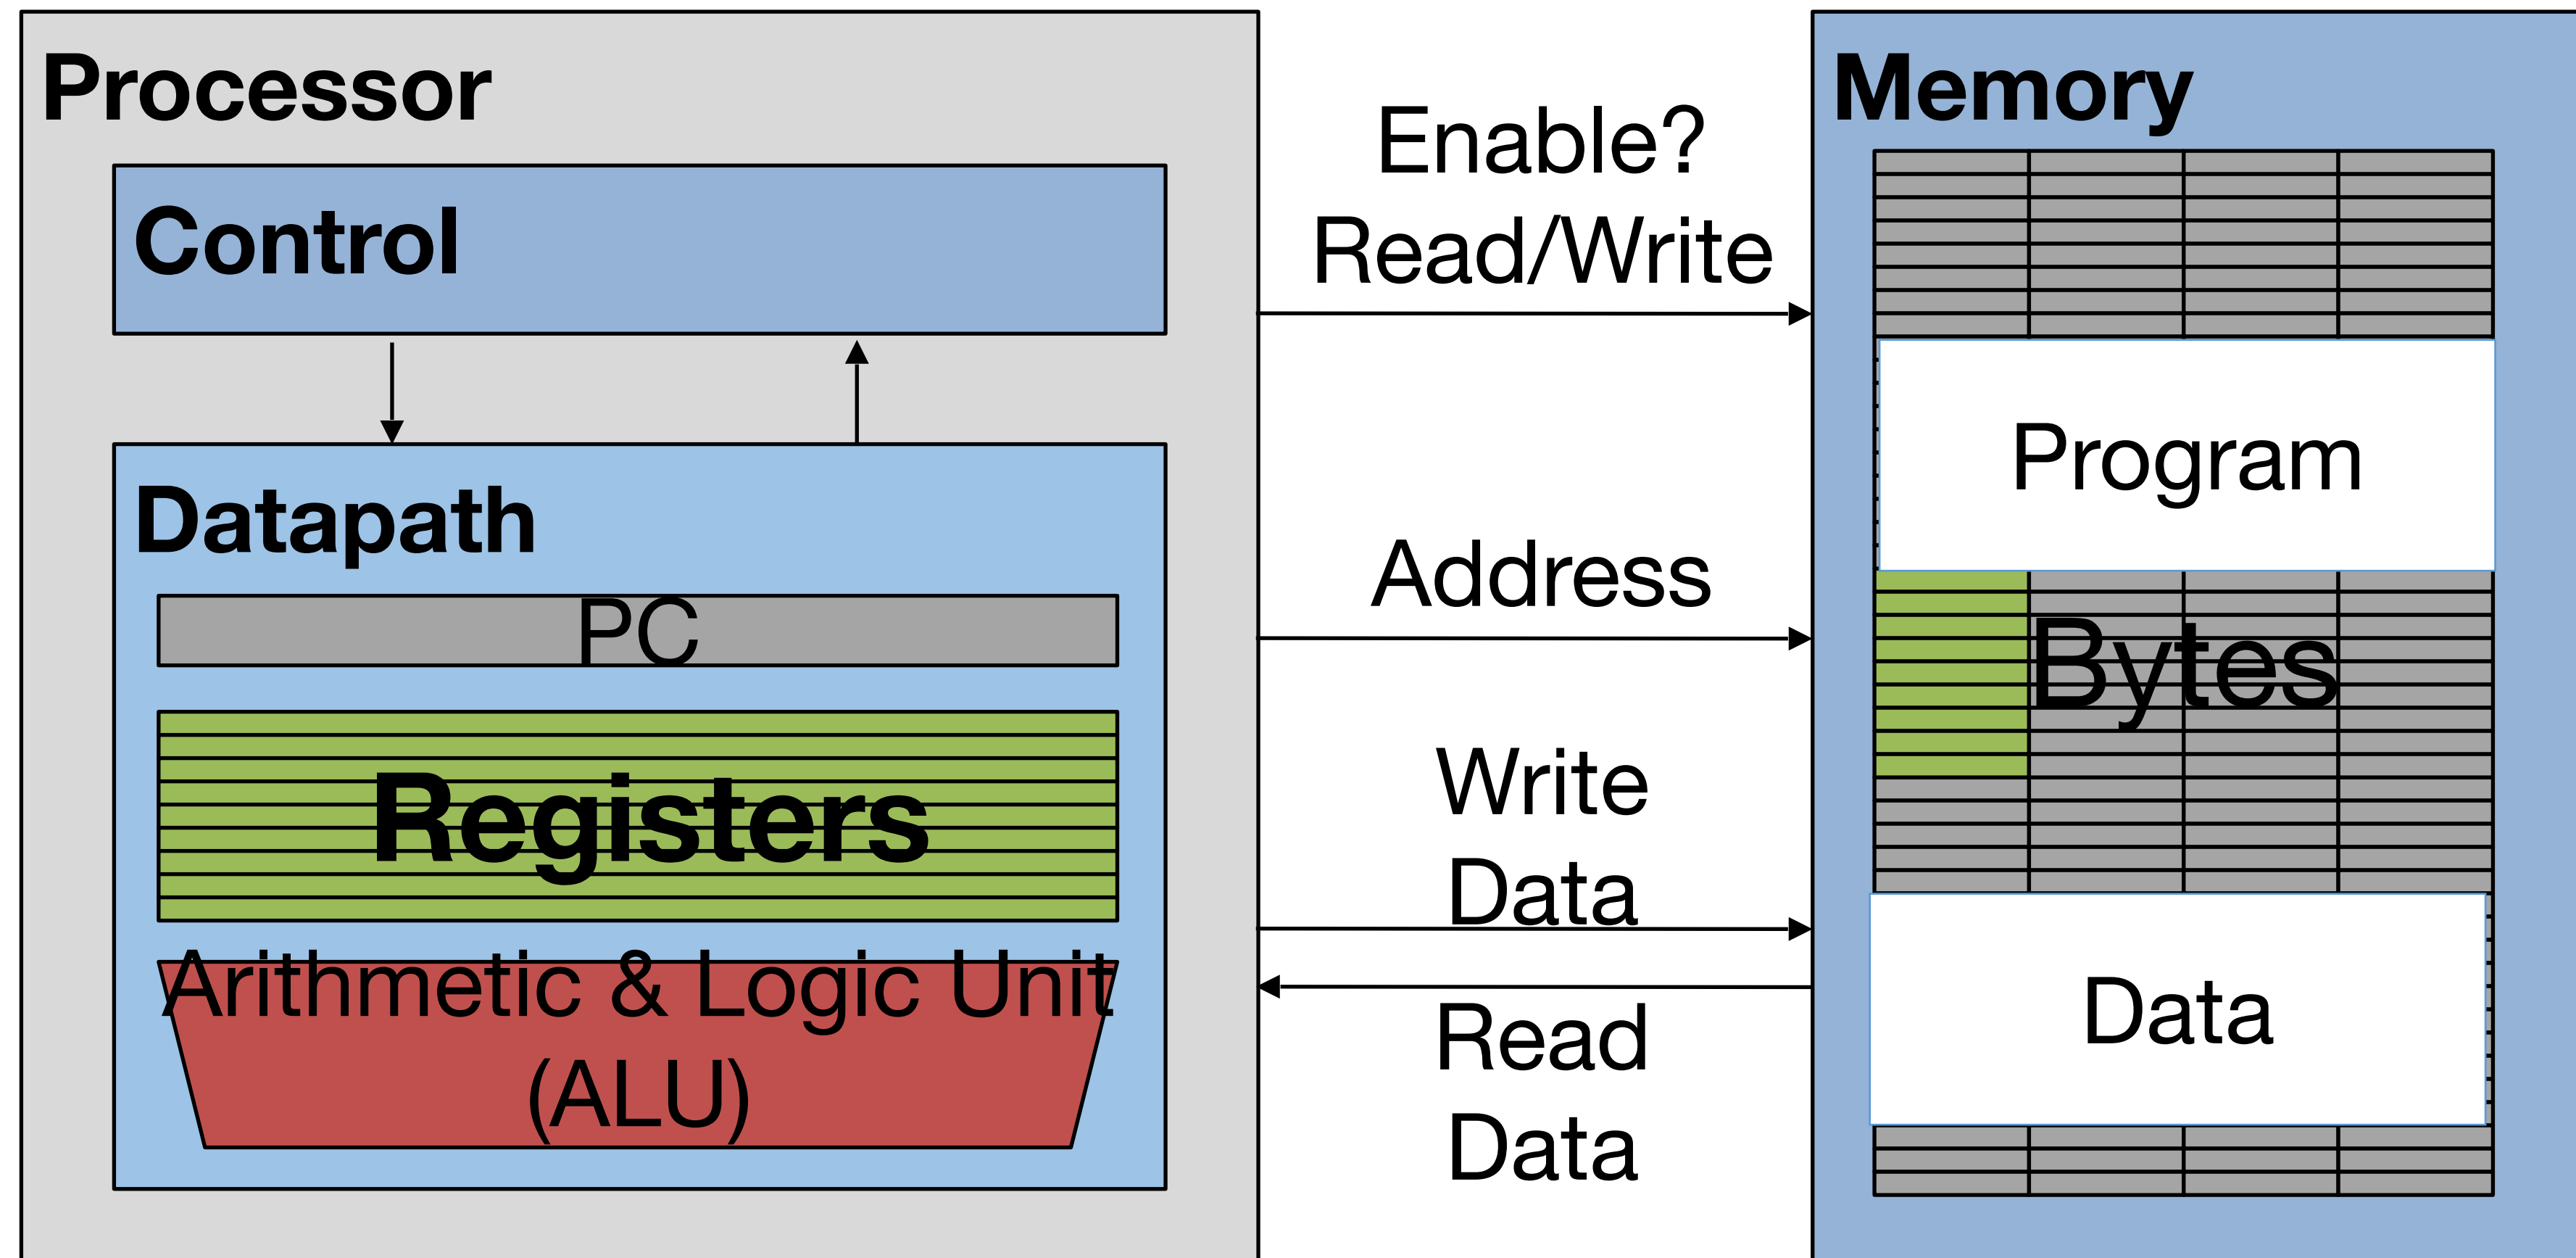
Lecture 14

Pipelining RISC-V

Review

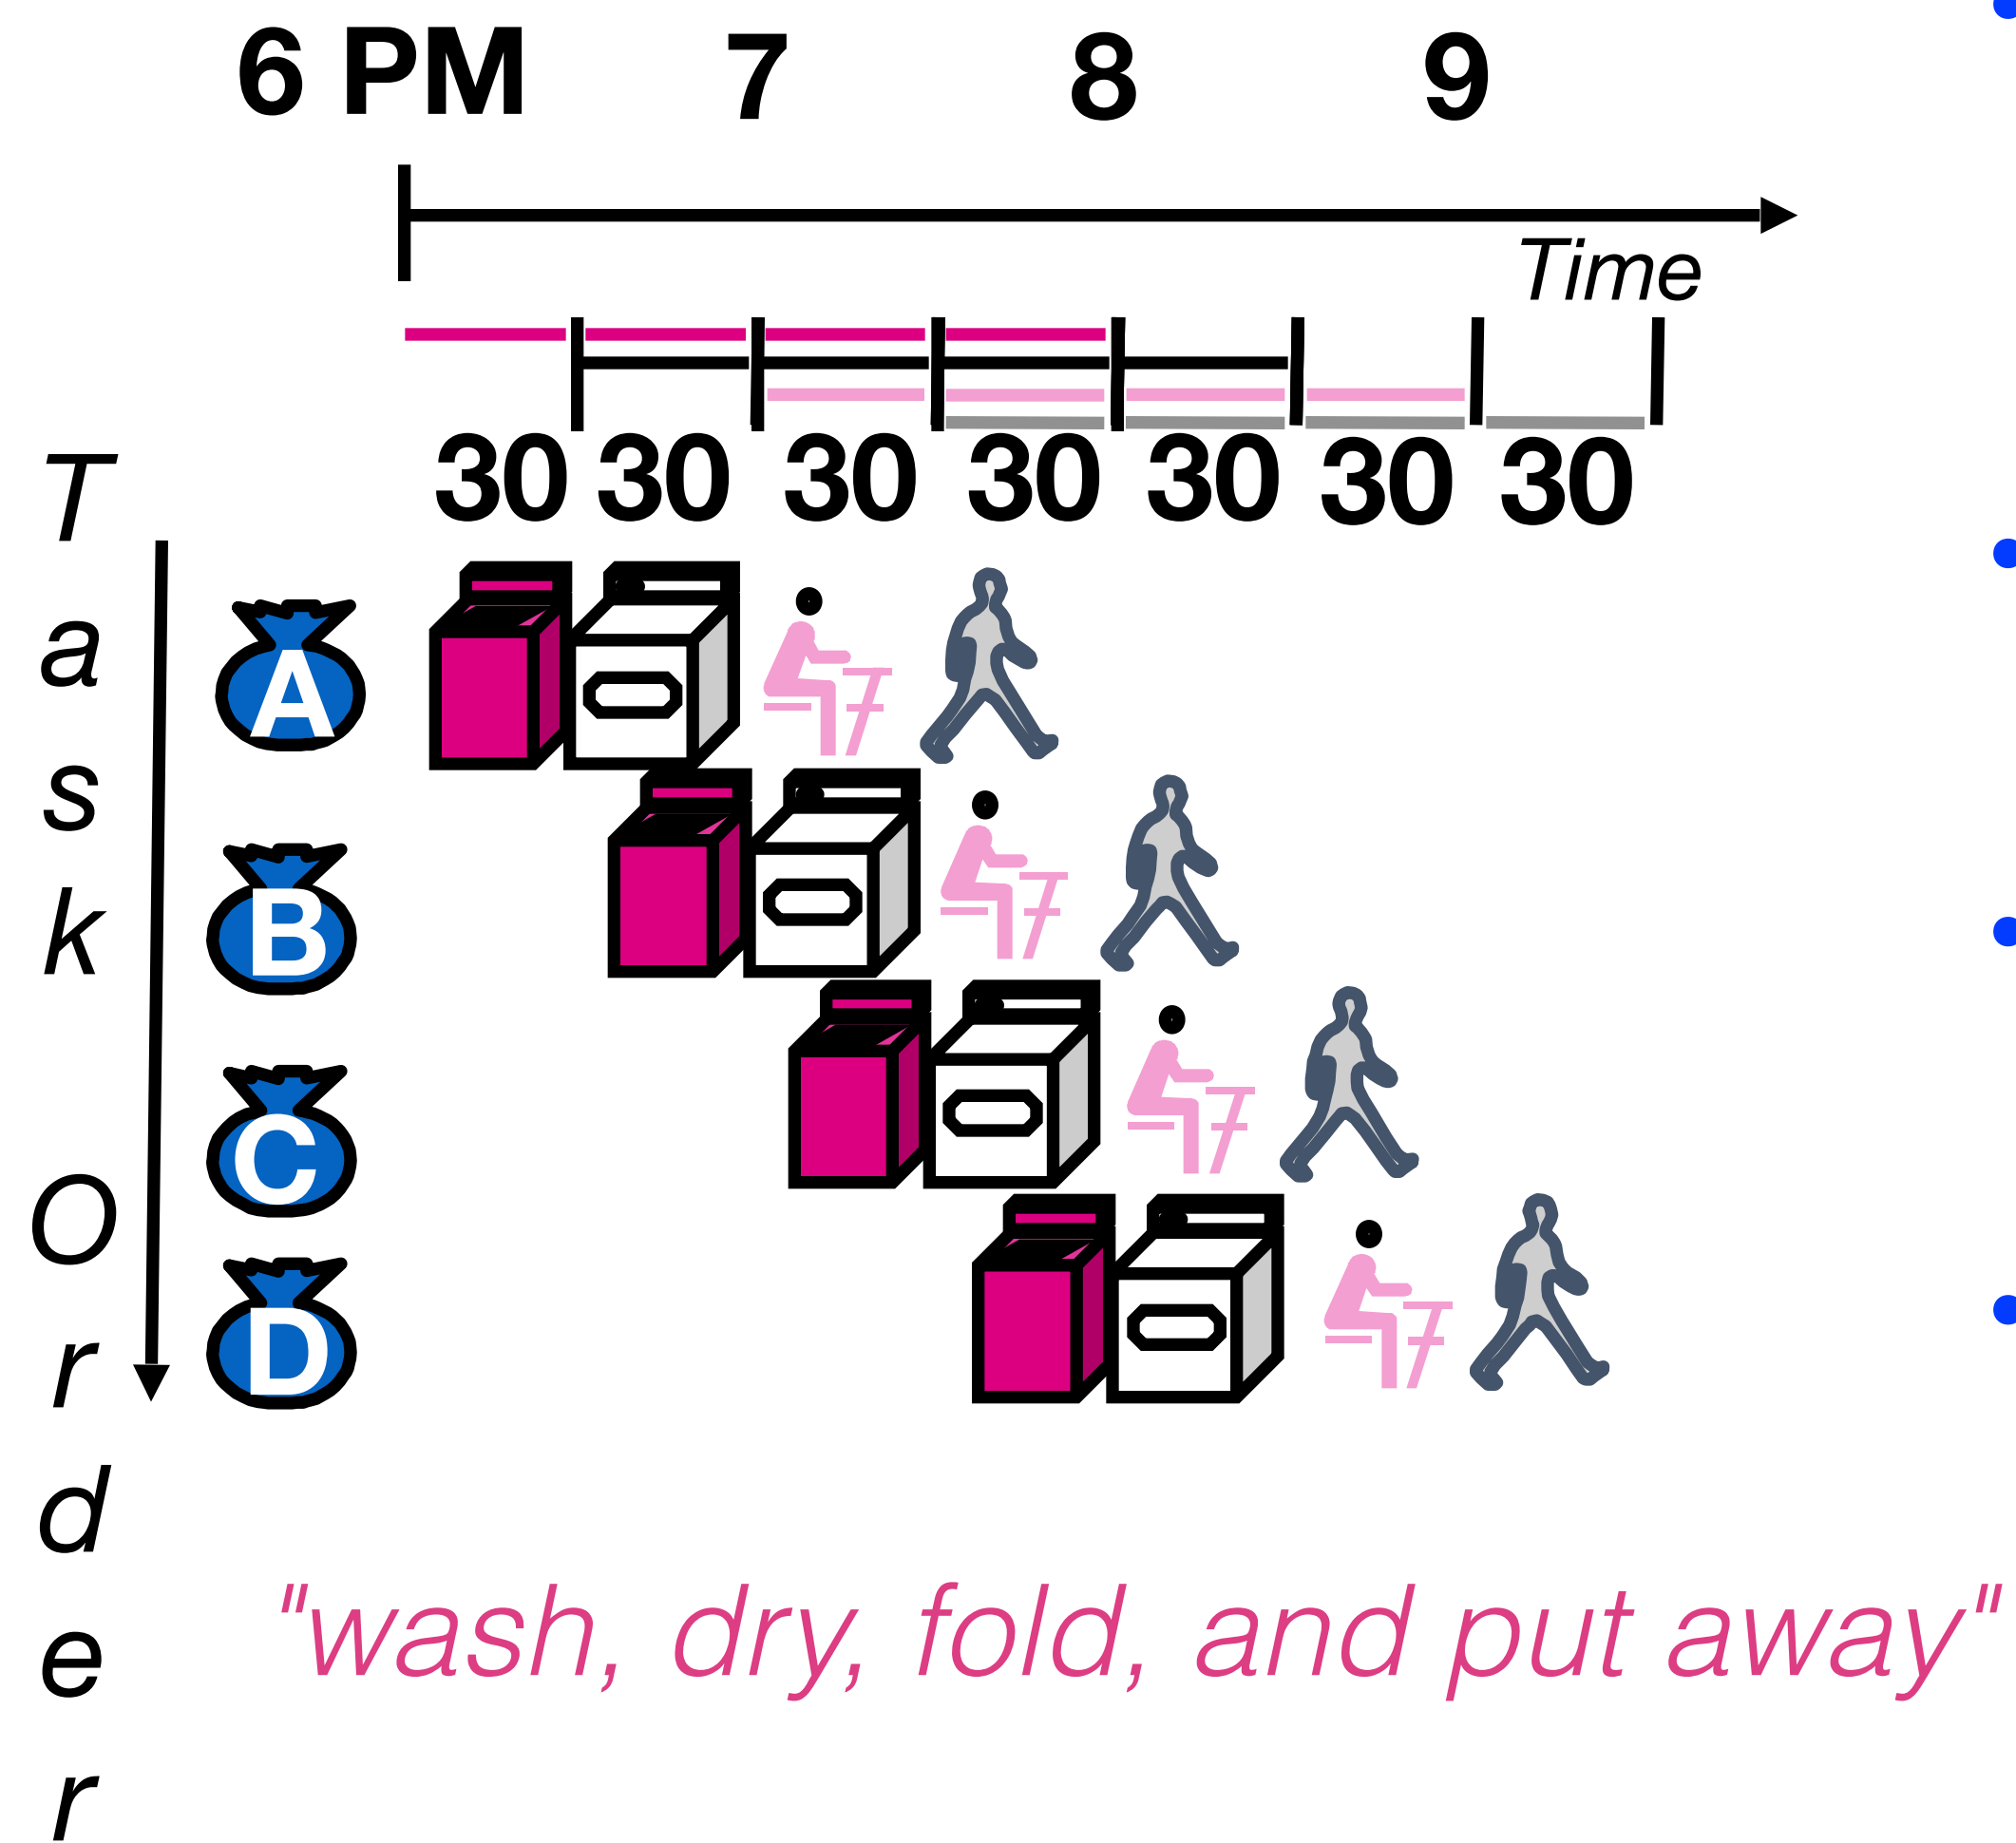
- Controller
 - Tells universal datapath how to execute each instruction
- Instruction timing
 - Set by instruction complexity, architecture, technology
 - Pipelining increases clock frequency, “instructions per second”
 - But does not reduce time to complete instruction
- Performance measures
 - Different measures depending on objective
 - Response time
 - Jobs / second
 - Energy Efficiency (joules/operation, joules/instruction)

Processor



Processor-Memory Interface

Pipelining Overview (Review)

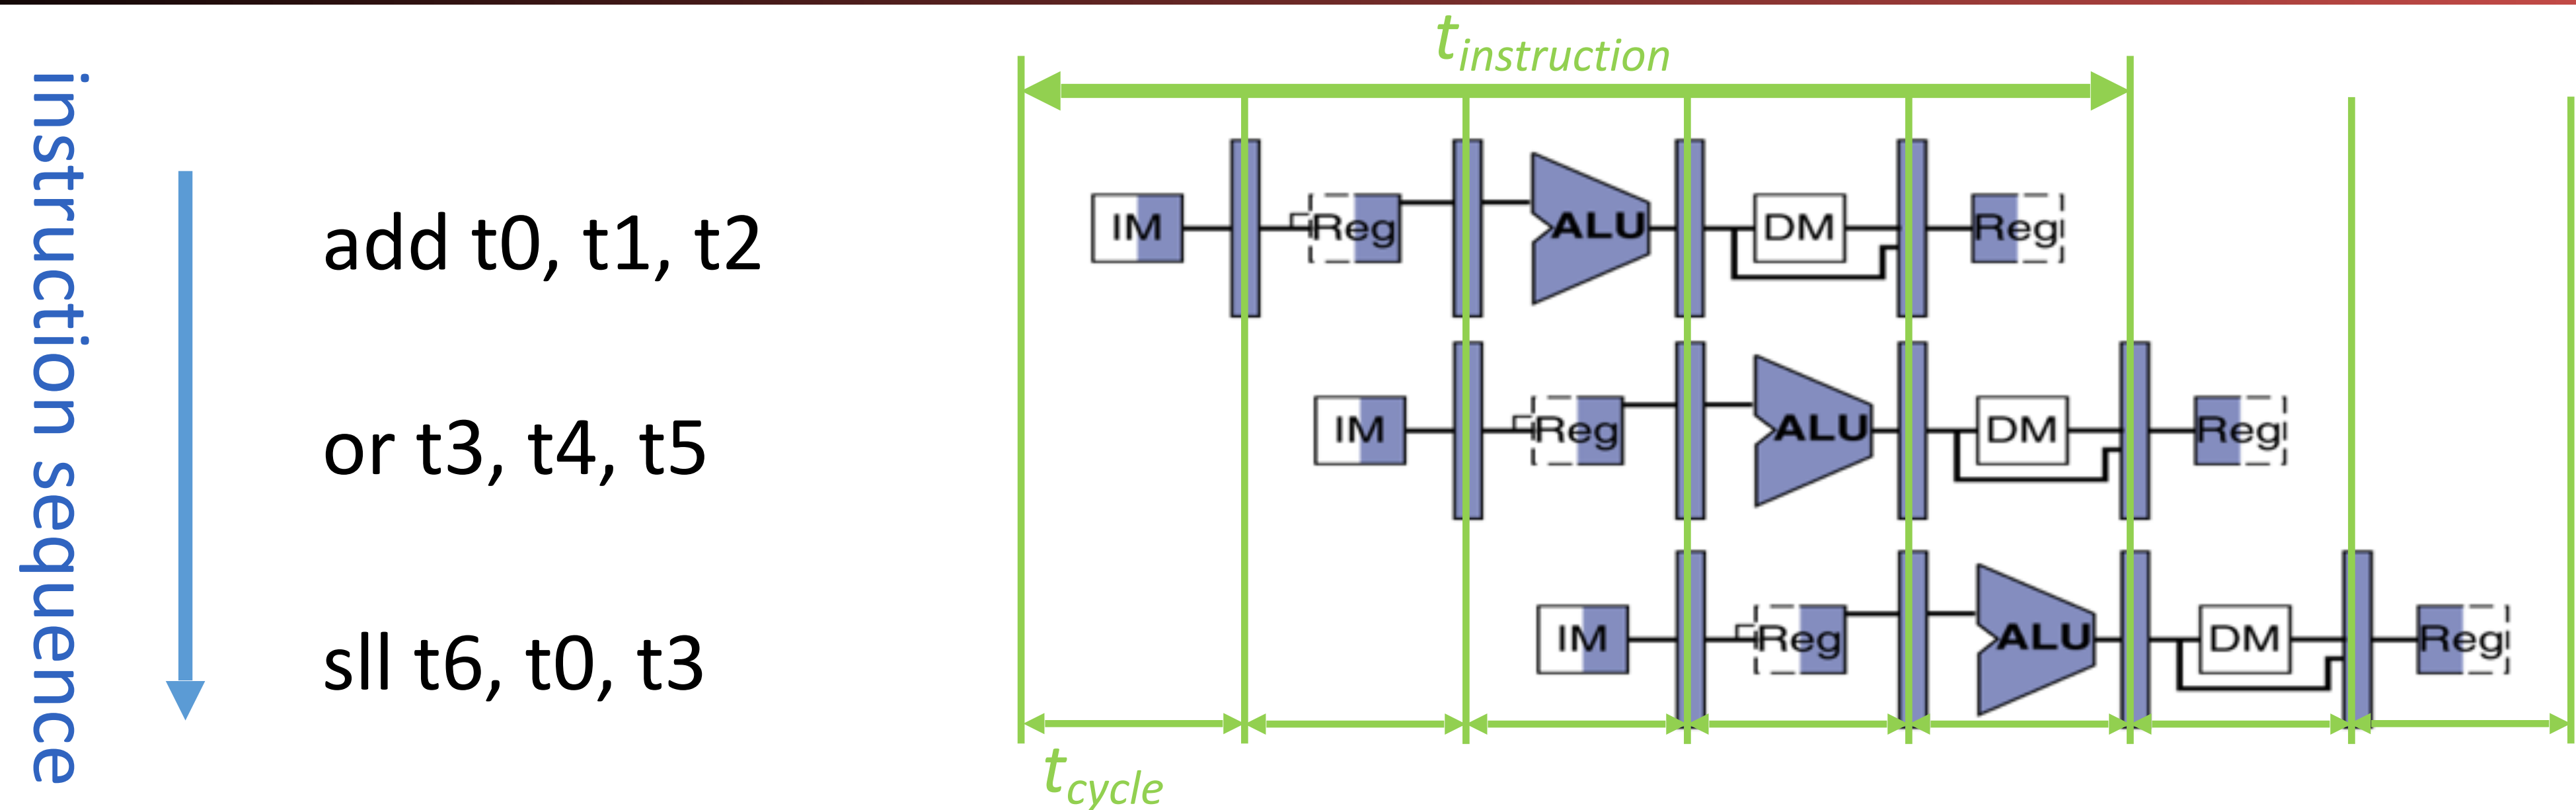


- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- Multiple tasks operating simultaneously using different resources
- Potential speedup = Number pipe stages
- Time to "fill" pipeline and time to "drain" it reduces speedup: 2.3X v. 4X in this example
- With lots of laundry, approaches 4X

Agenda

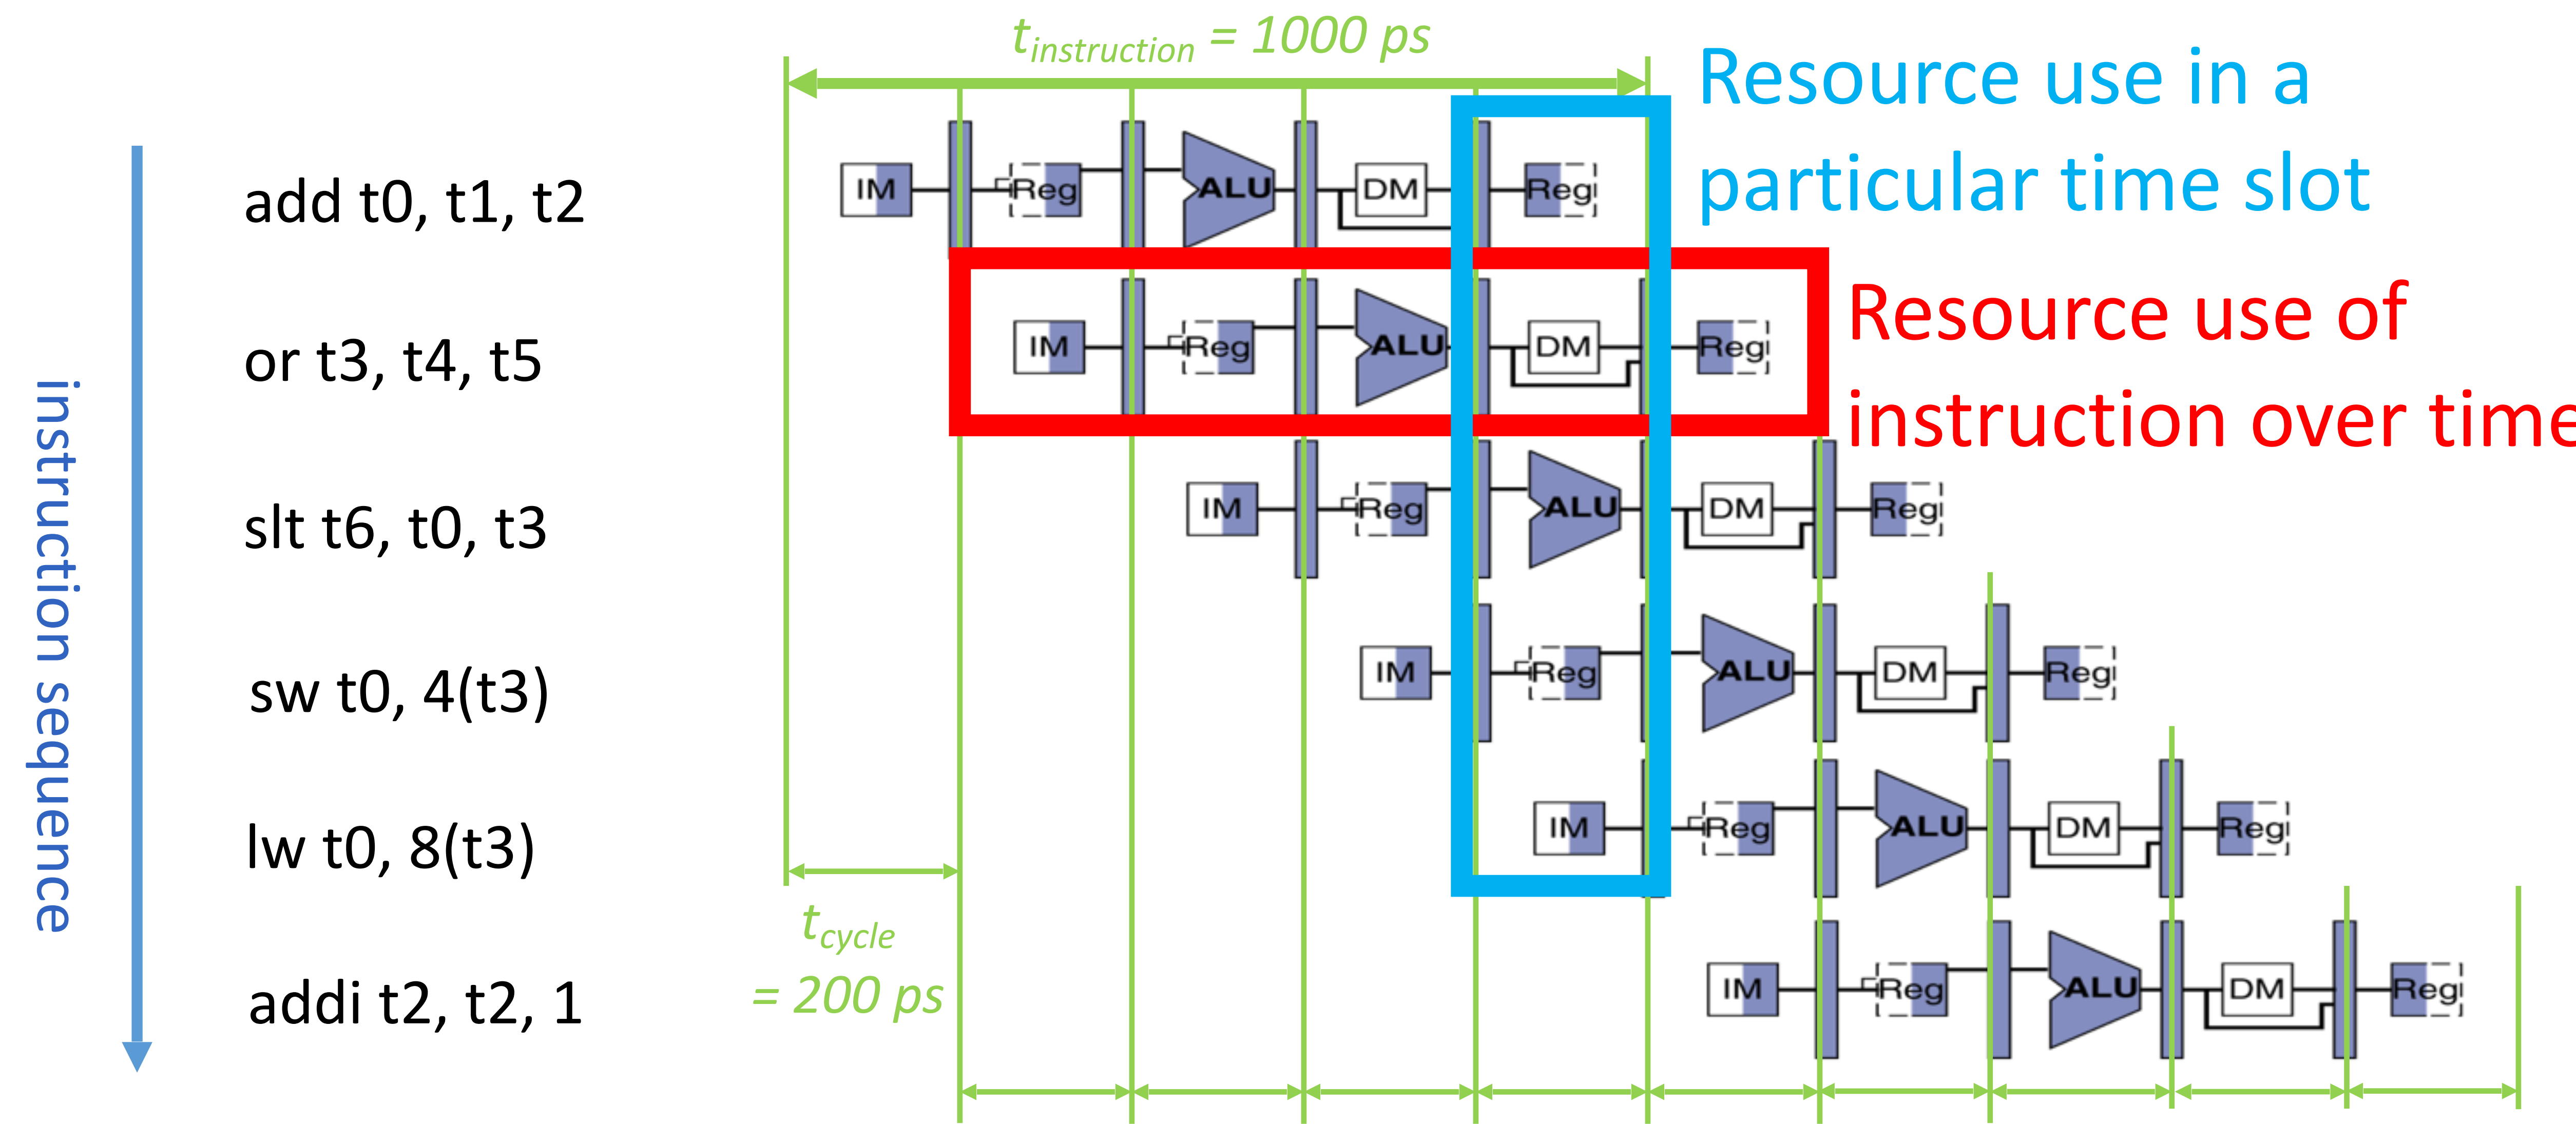
- **RISC-V Pipeline**
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Recap: Pipelining with RISC-V

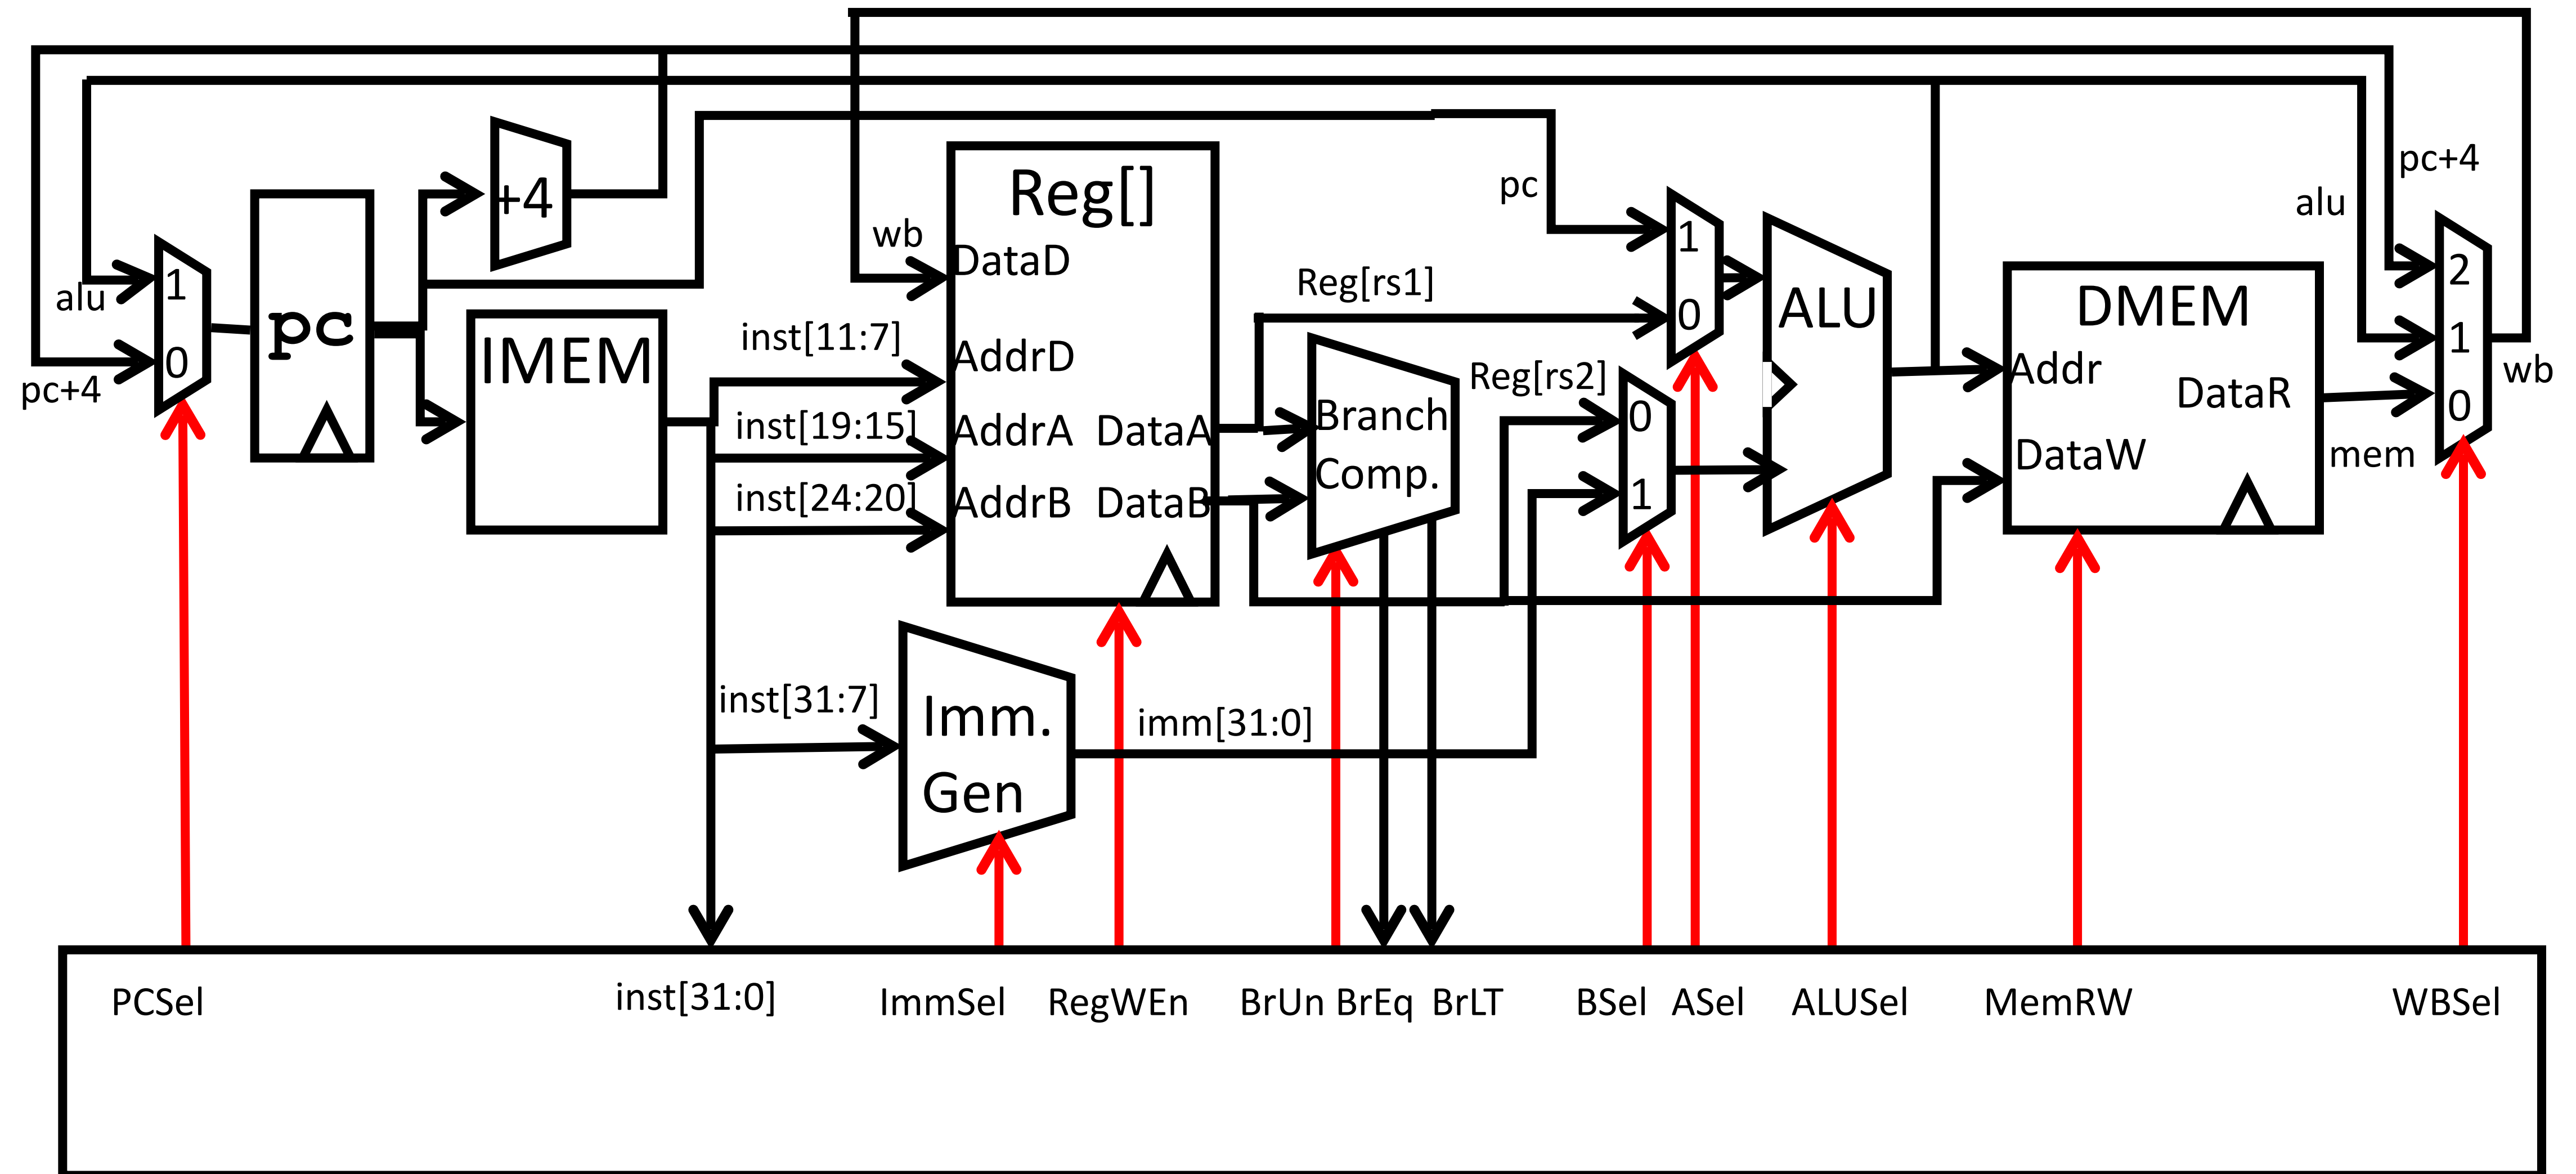


	Single Cycle	Pipelining
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, f_s	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Relative speed	1 x	4 x

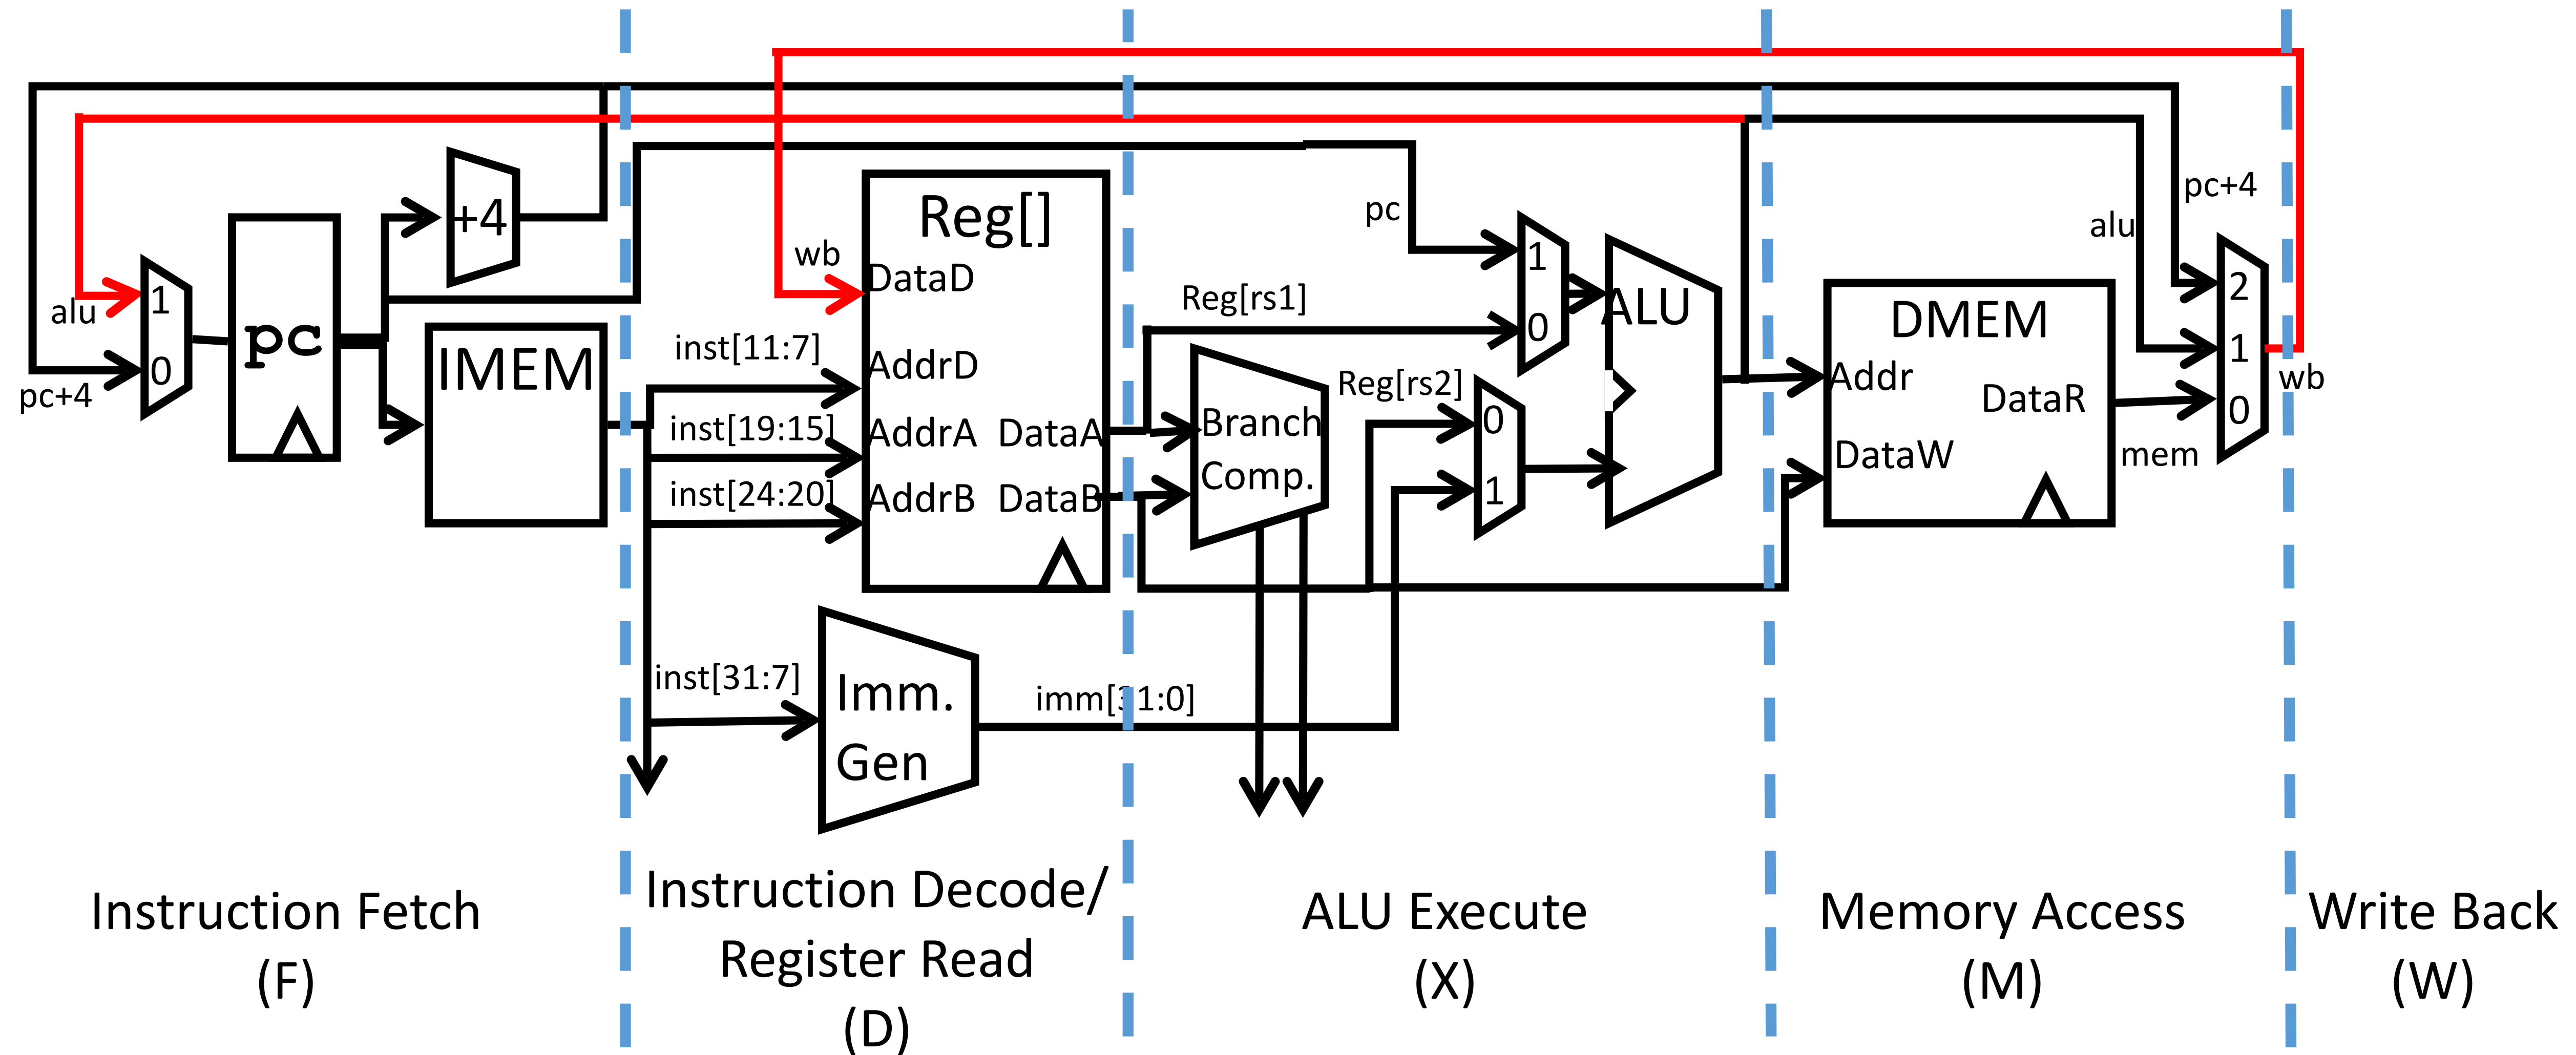
RISC-V Pipeline



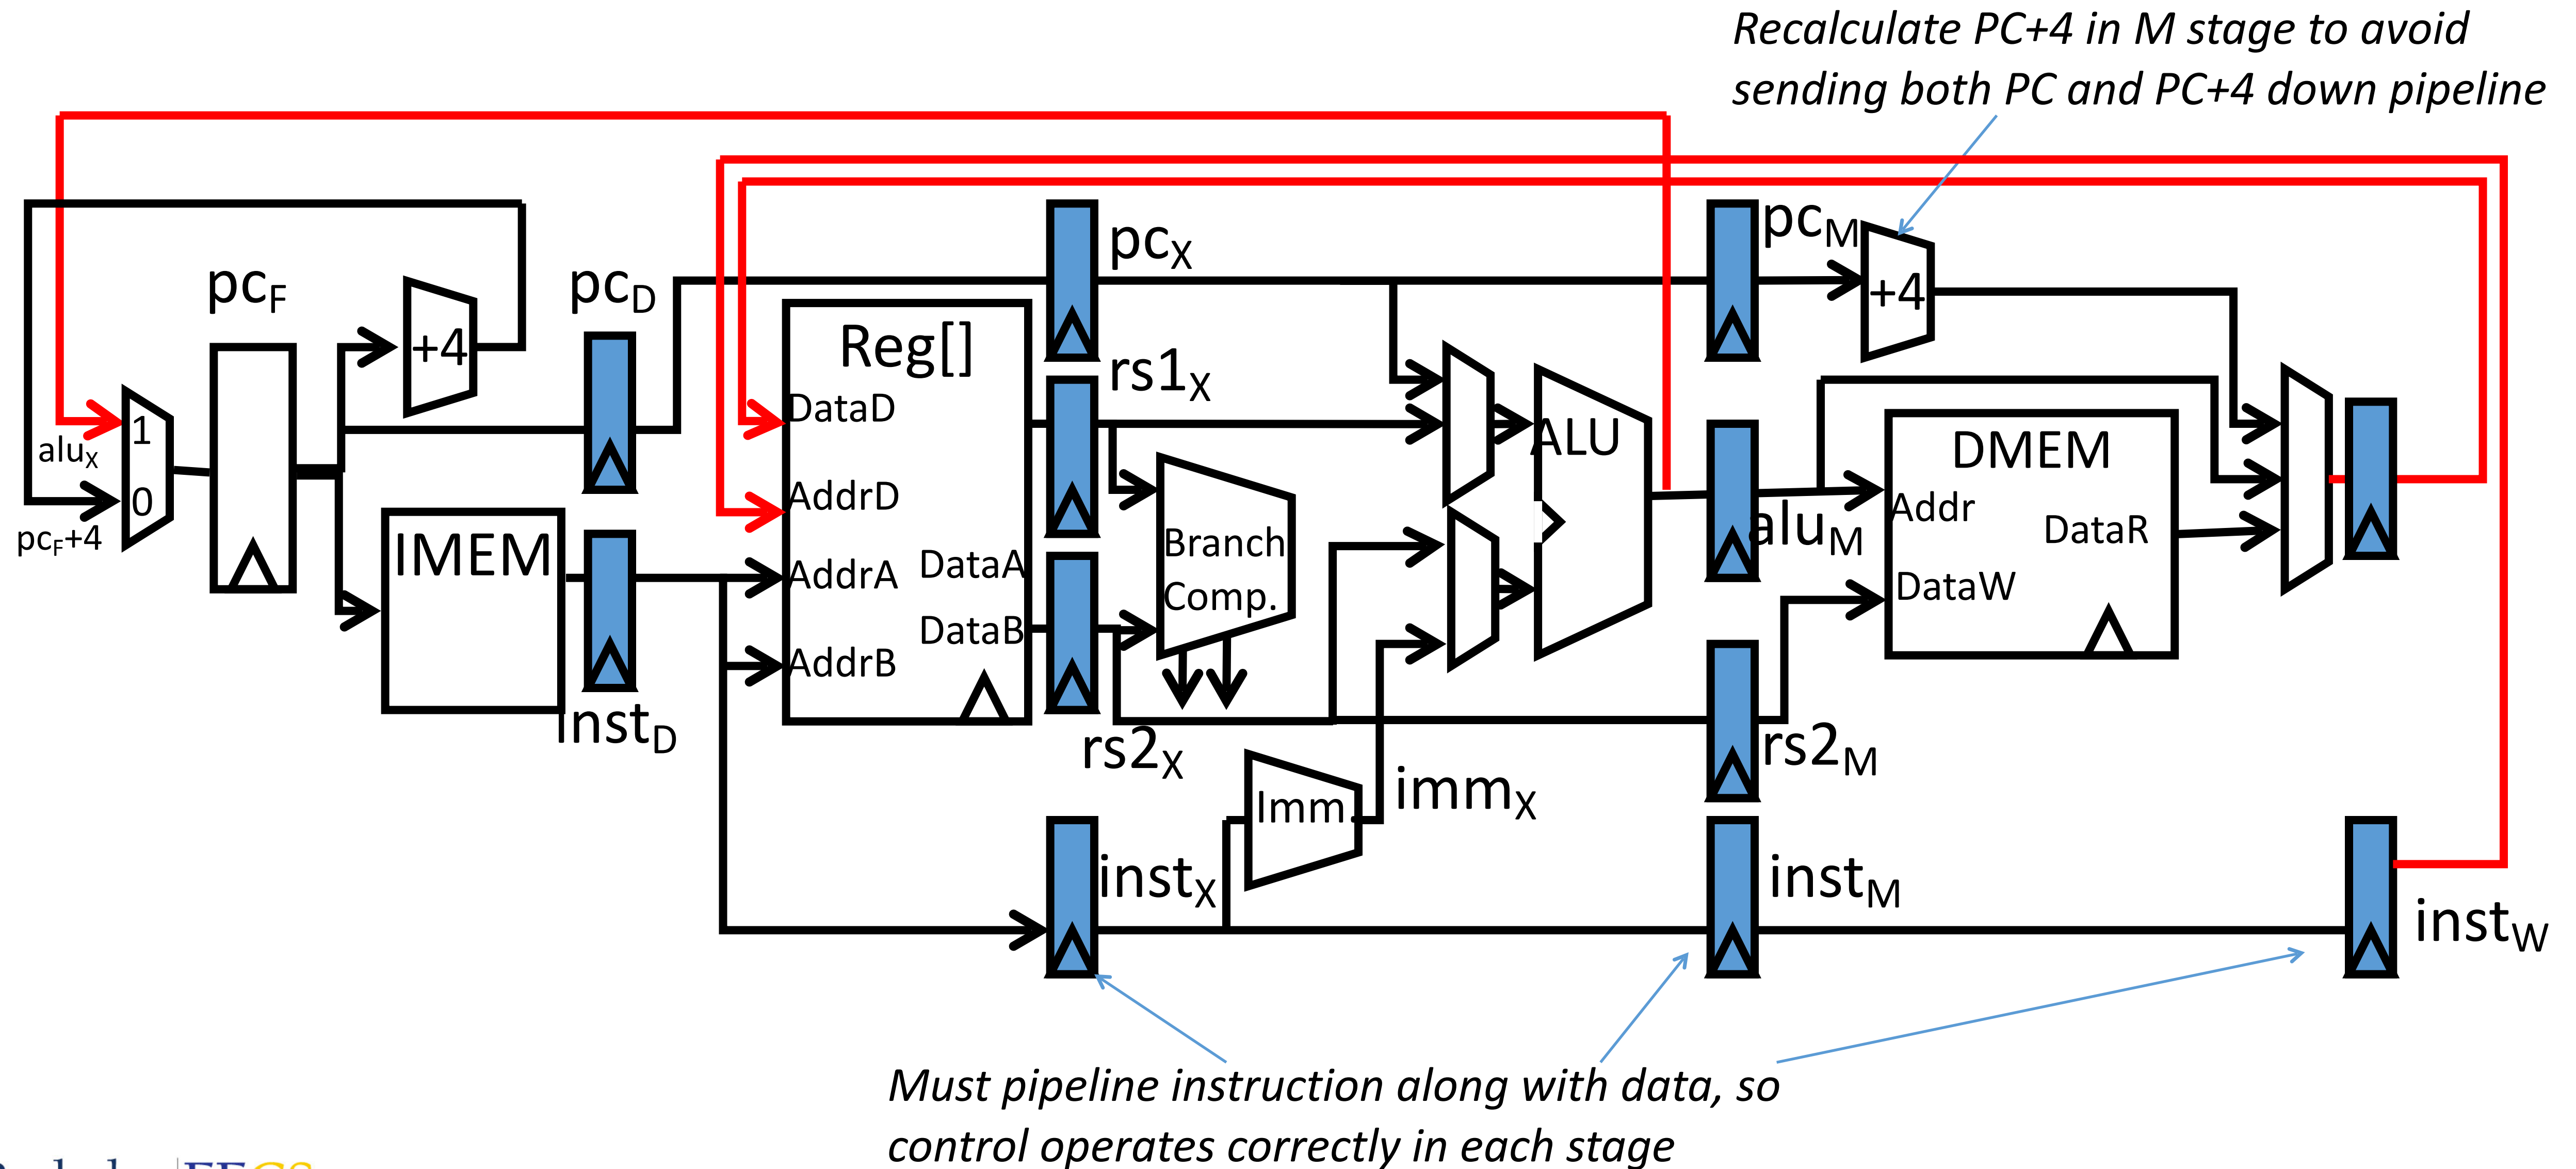
Single-Cycle RISC-V RV32I Datapath



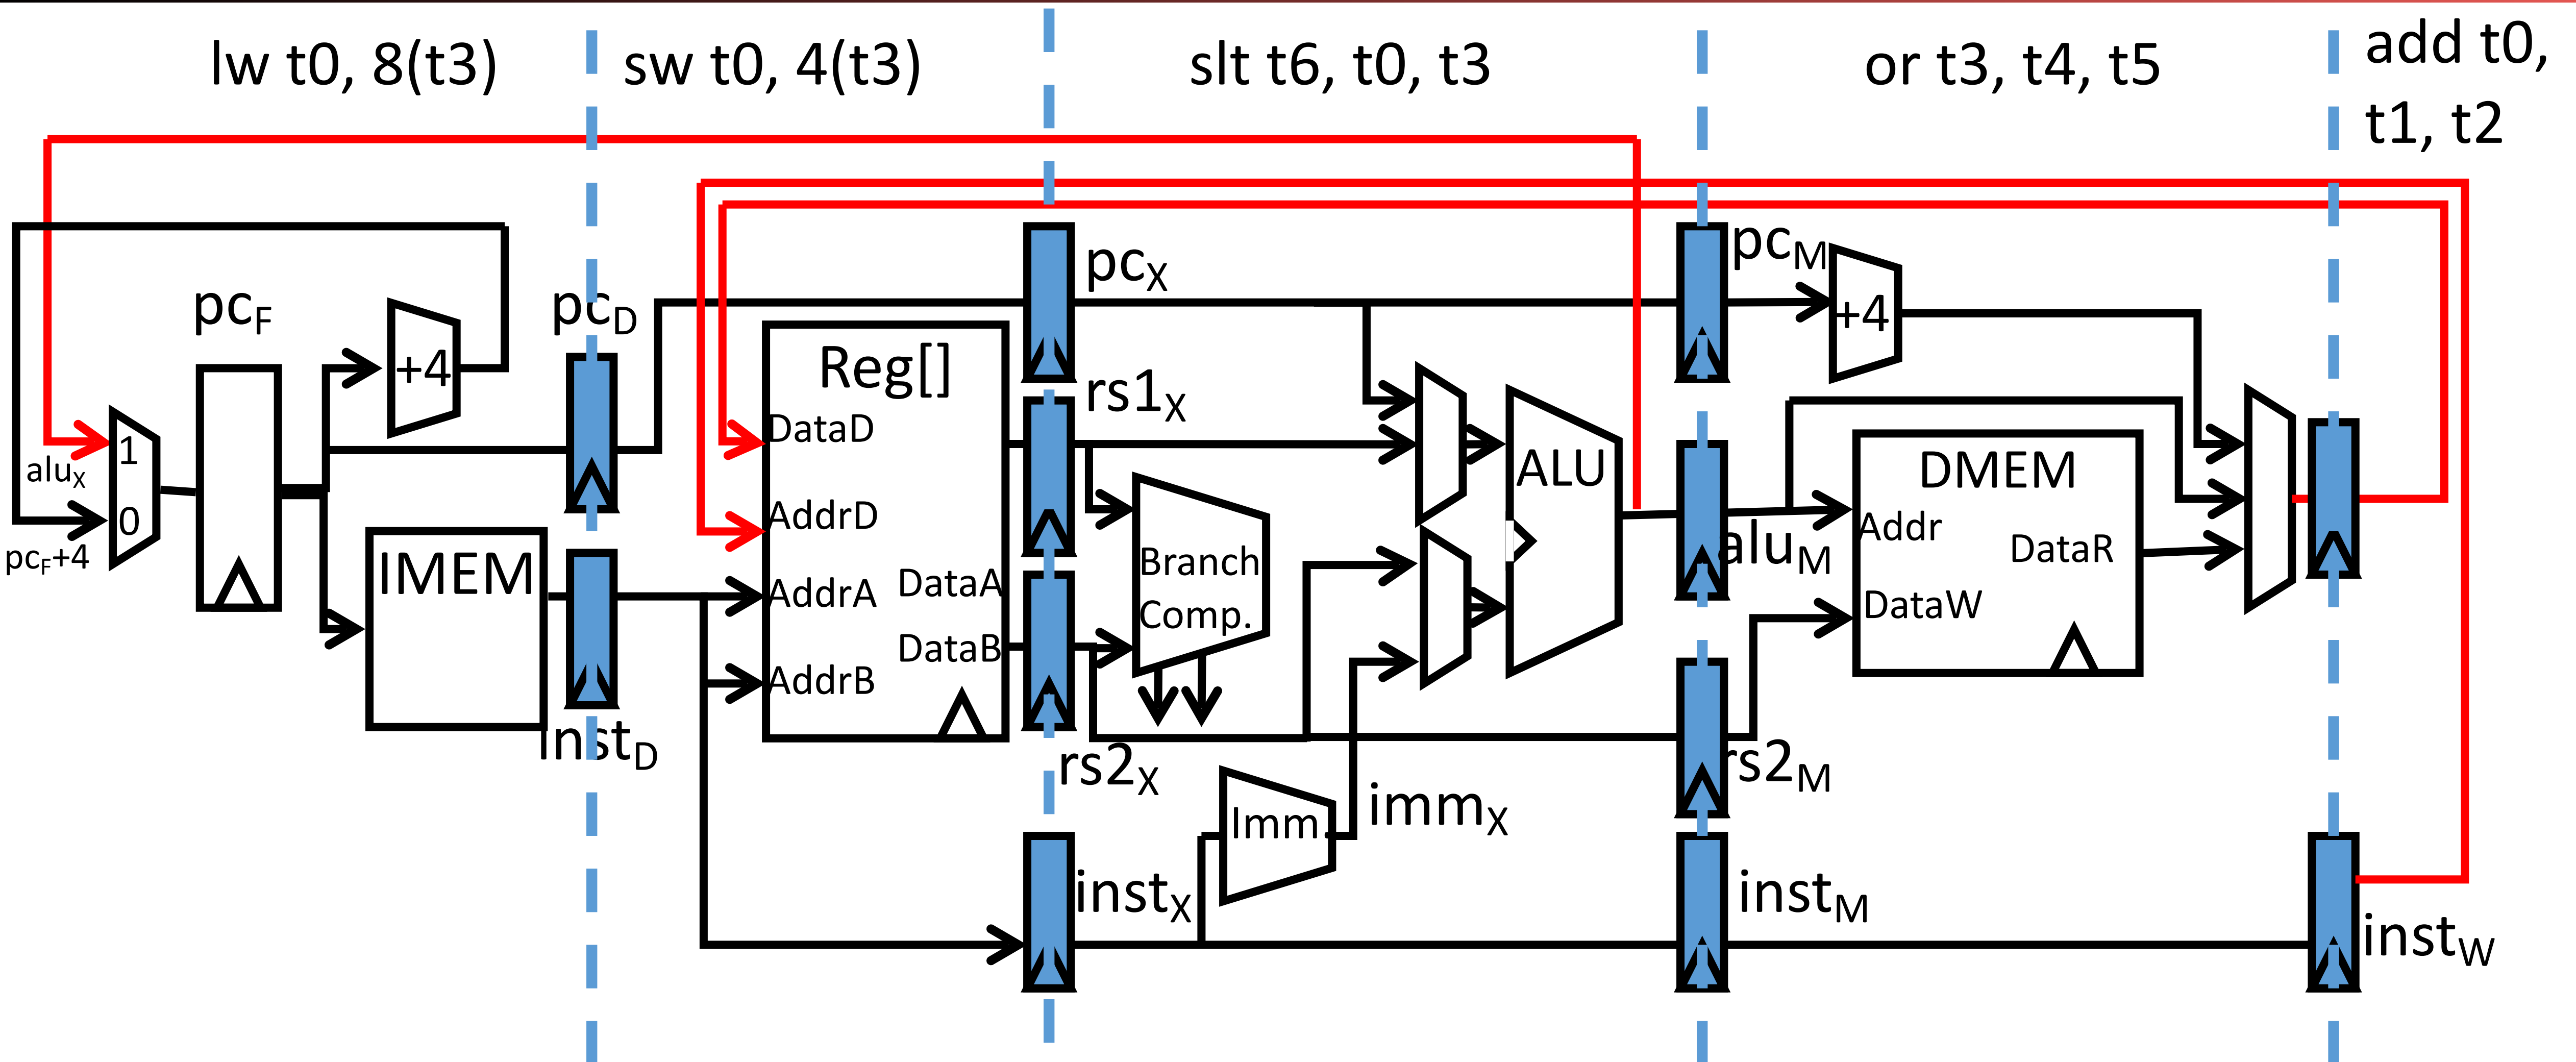
Pipelining RISC-V RV32I Datapath



Pipelined RISC-V RV32I Datapath



Each stage operates on different instruction



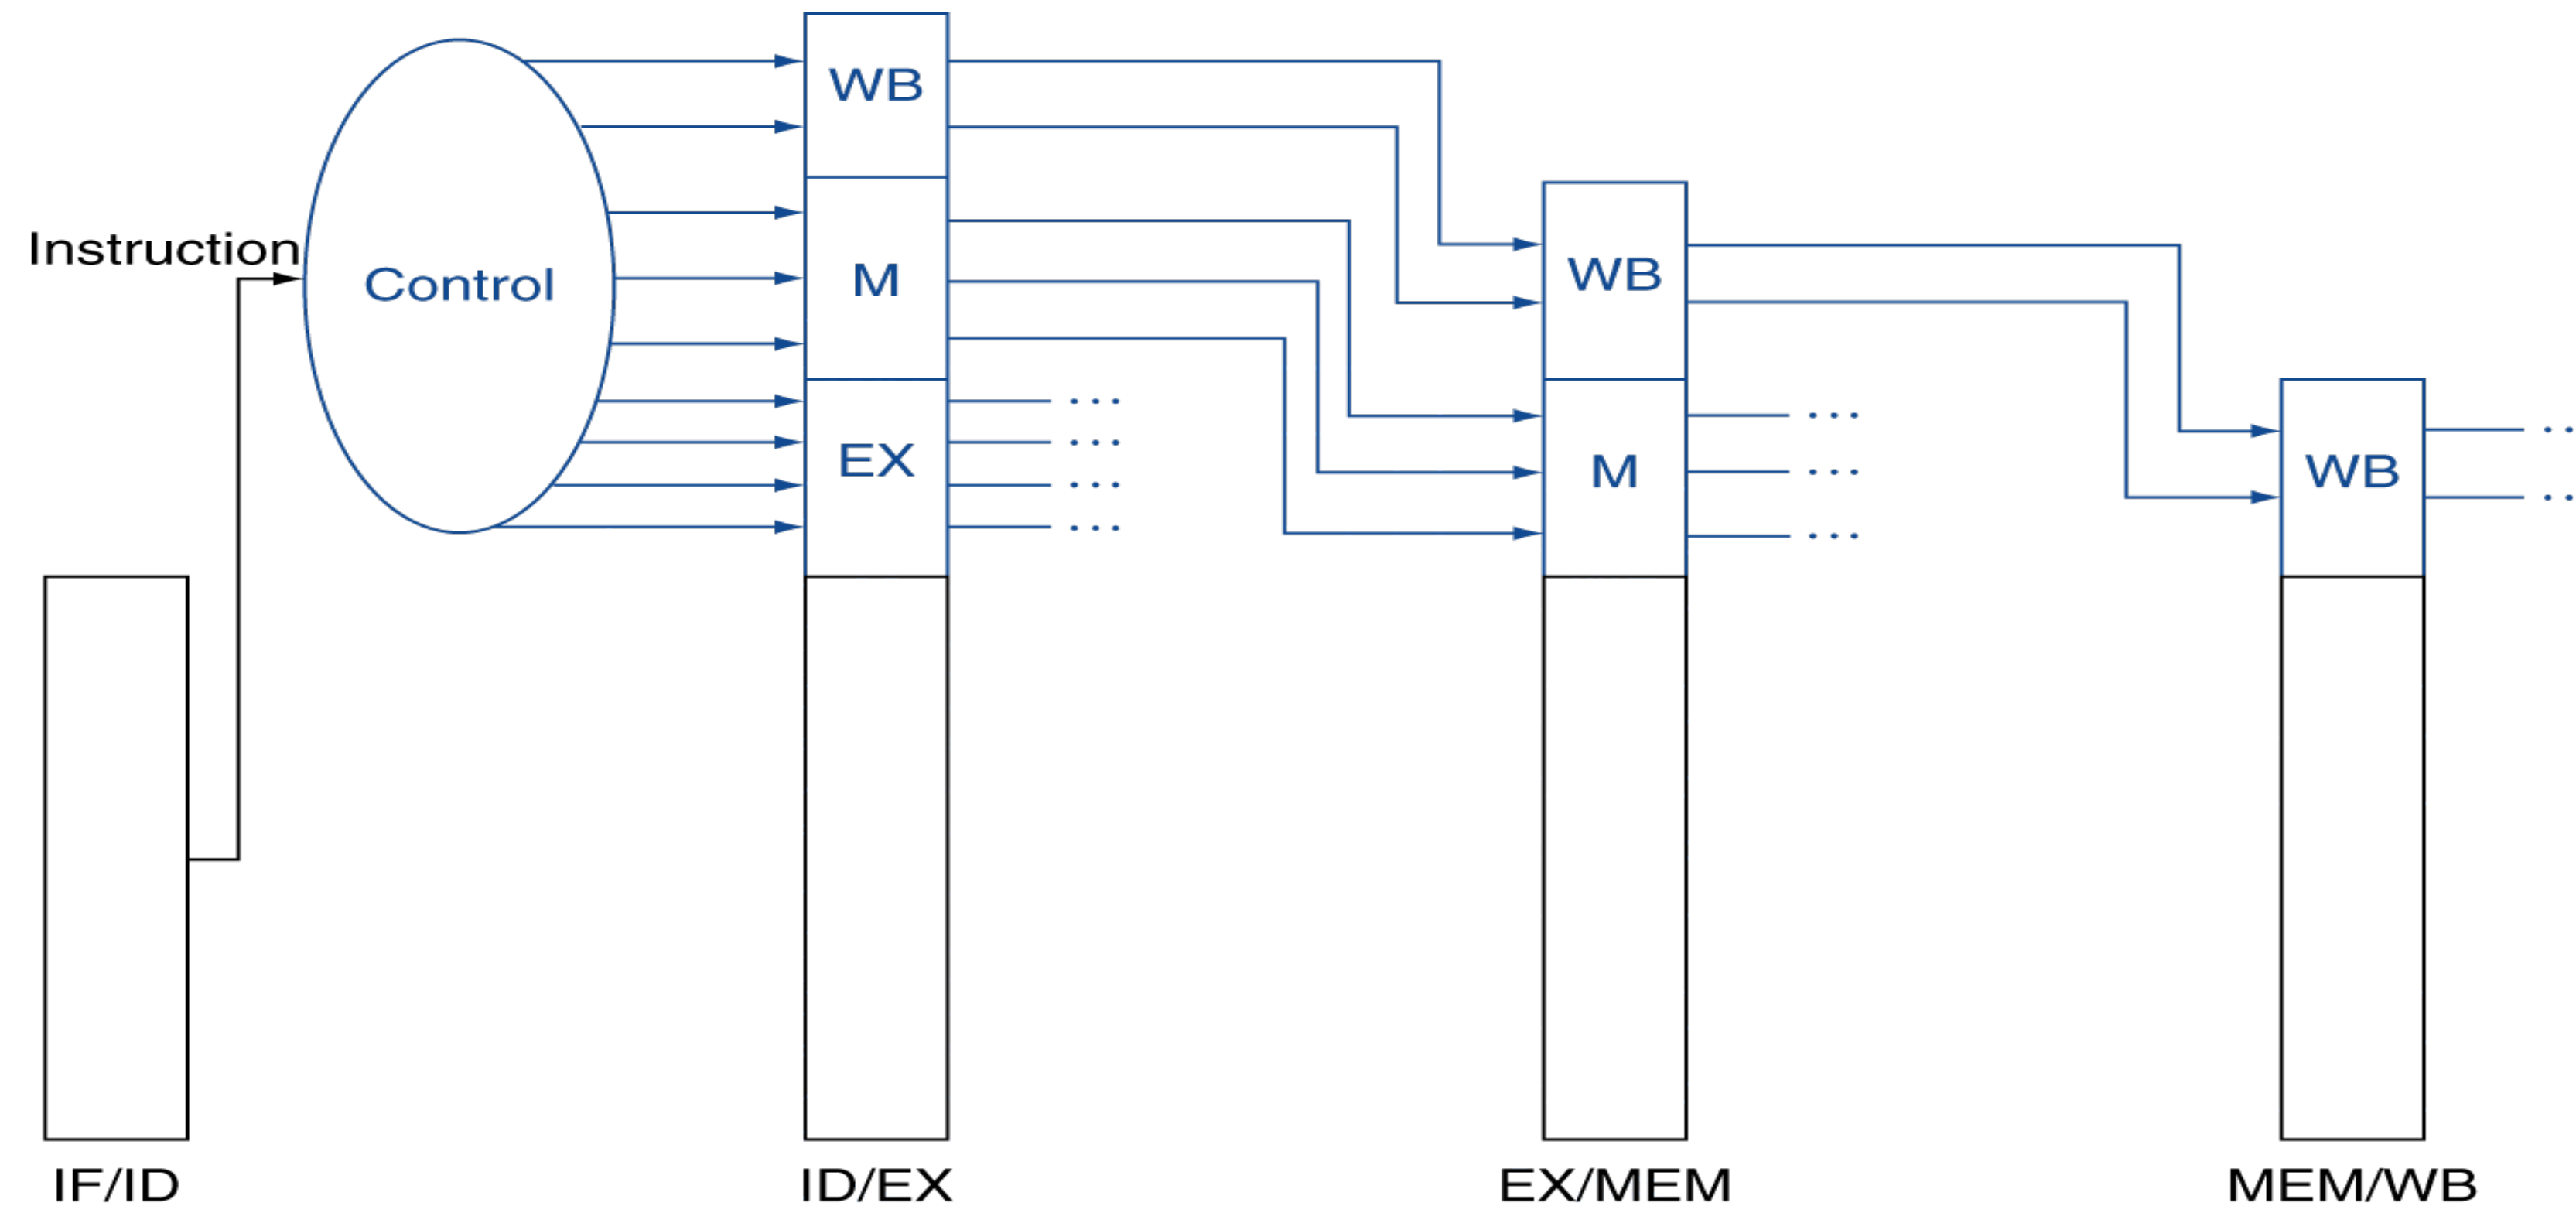
Pipeline registers separate stages, hold data for each instruction in flight

Agenda

- RISC-V Pipeline
- **Pipeline Control**
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages



Question:

1 / throughput

$$\frac{\textit{time}}{\textit{program}} = \frac{\textit{instructions}}{\textit{program}} \cdot \frac{\textit{cycles}}{\textit{instruction}} \cdot \frac{\textit{time}}{\textit{cycle}}$$

Pipelining the single-cycle processor can increase processor performance by:

	Instructions /program	Cycles/ instruction	Time/cycle
A	decrease	decrease	same
B	same	increase	decrease
C	same	same	decrease
D	increase	decrease	increase

Hazards Ahead



Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - **Structural** \Rightarrow Two instructions compete for hardware
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Structural Hazard

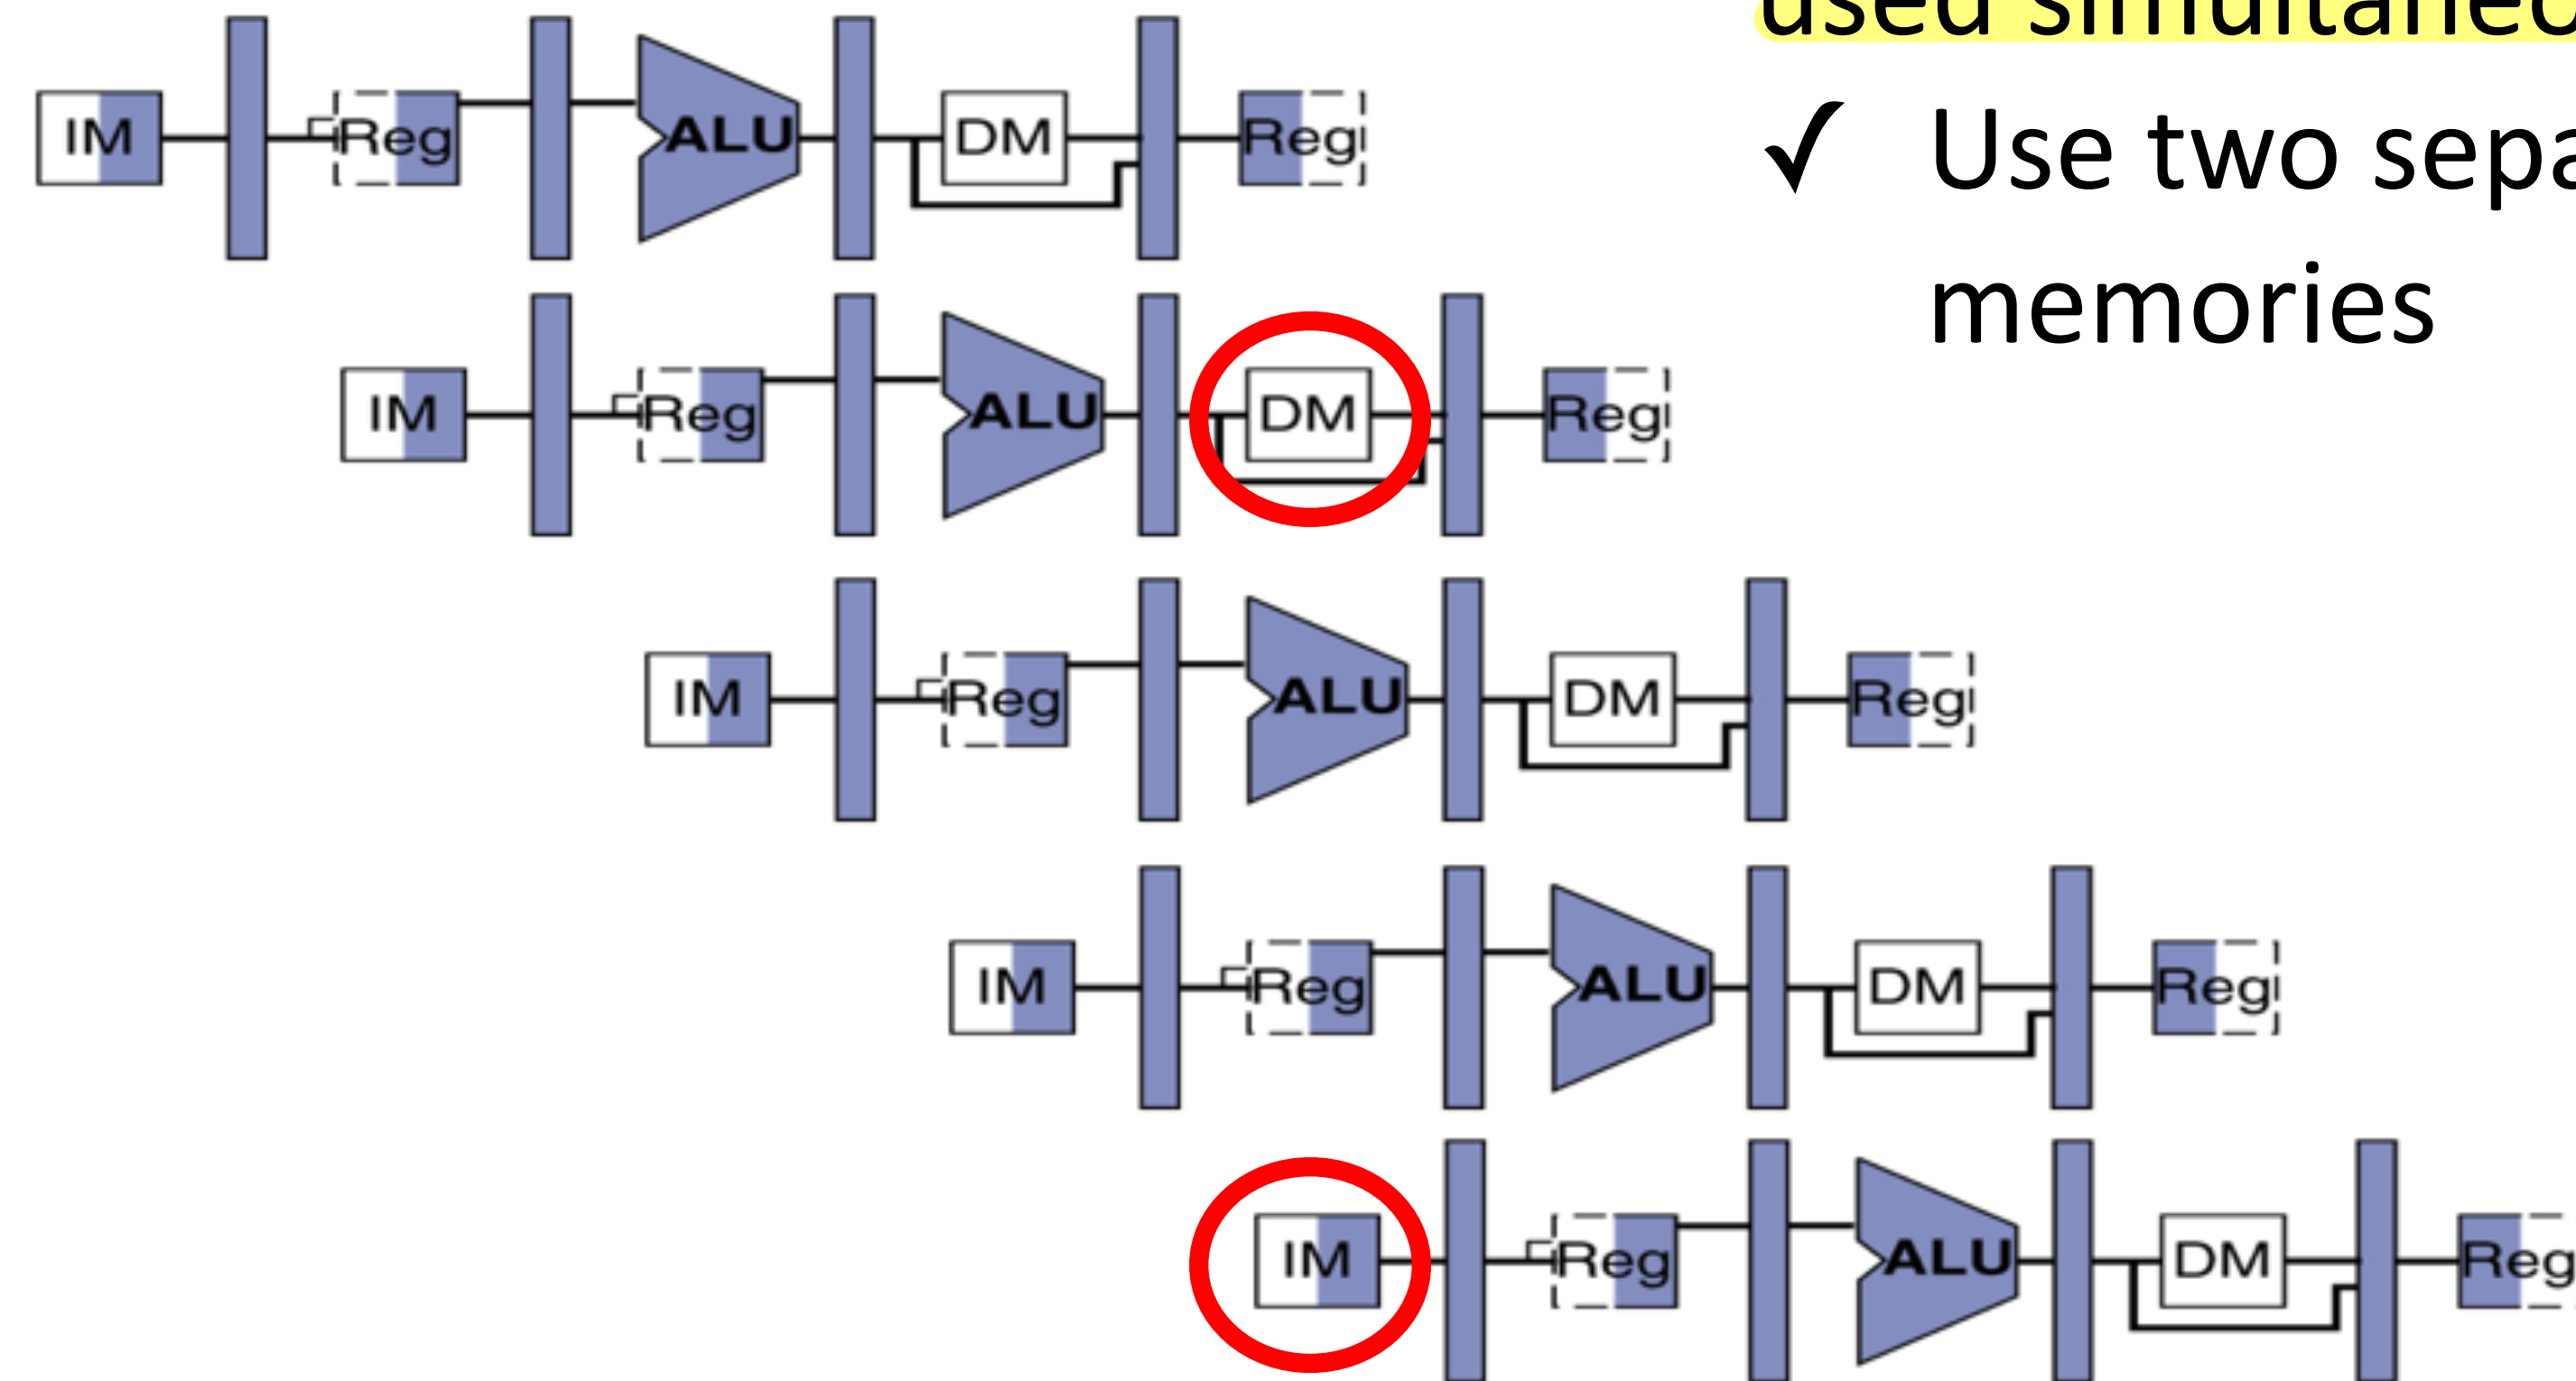
- **Problem:** Two or more instructions in the pipeline compete for access to a single physical resource
- **Solution 1:** Instructions take turns to use resource, some instructions have to stall
- **Solution 2:** Add more hardware to machine
 - *Can always solve a structural hazard by adding more hardware*

Regfile Structural Hazards

- Each instruction:
 - can read up to two operands in decode stage
 - can write one value in writeback stage
 - *therefore two different instructions might be accessing the register file on the same cycle!*
- Avoid structural hazard by having separate “ports”
 - two independent read ports and one independent write port
- *Reads from one instruction and writes from another can happen simultaneously*

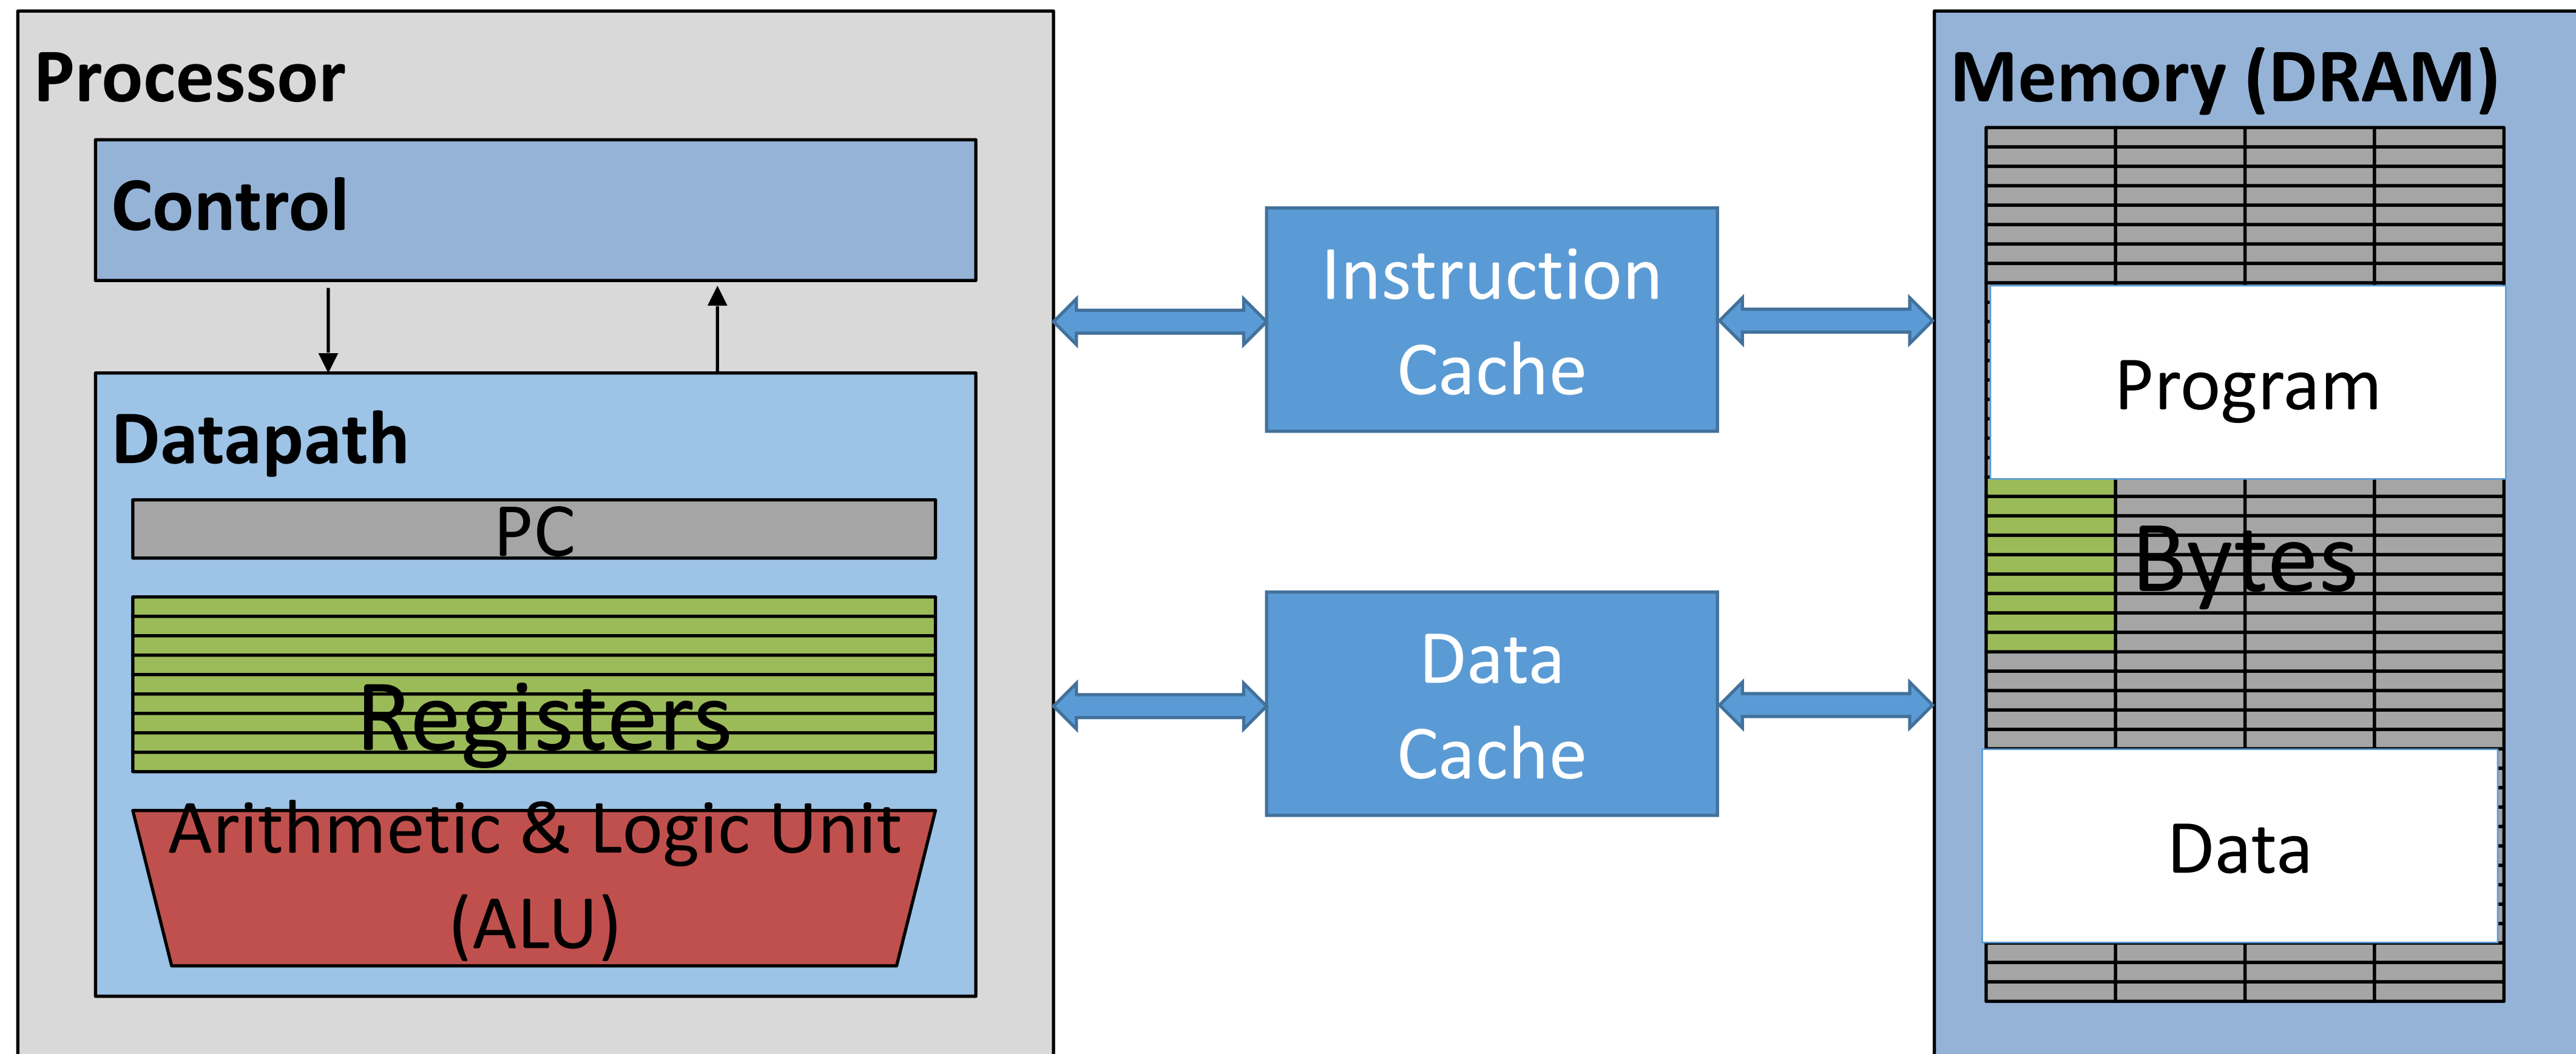
Structural Hazard: Memory Access

instruction sequence
↓
add t0, t1, t2
sw t0, 4(t5)
slt t6, t0, t3
or t3, t4, t5
lw t0, 8(t3)



- Instruction and data memory used simultaneously
✓ Use two separate memories

Instruction and Data Caches



Caches: relatively small and fast “buffer” memories

Structural Hazards – Summary

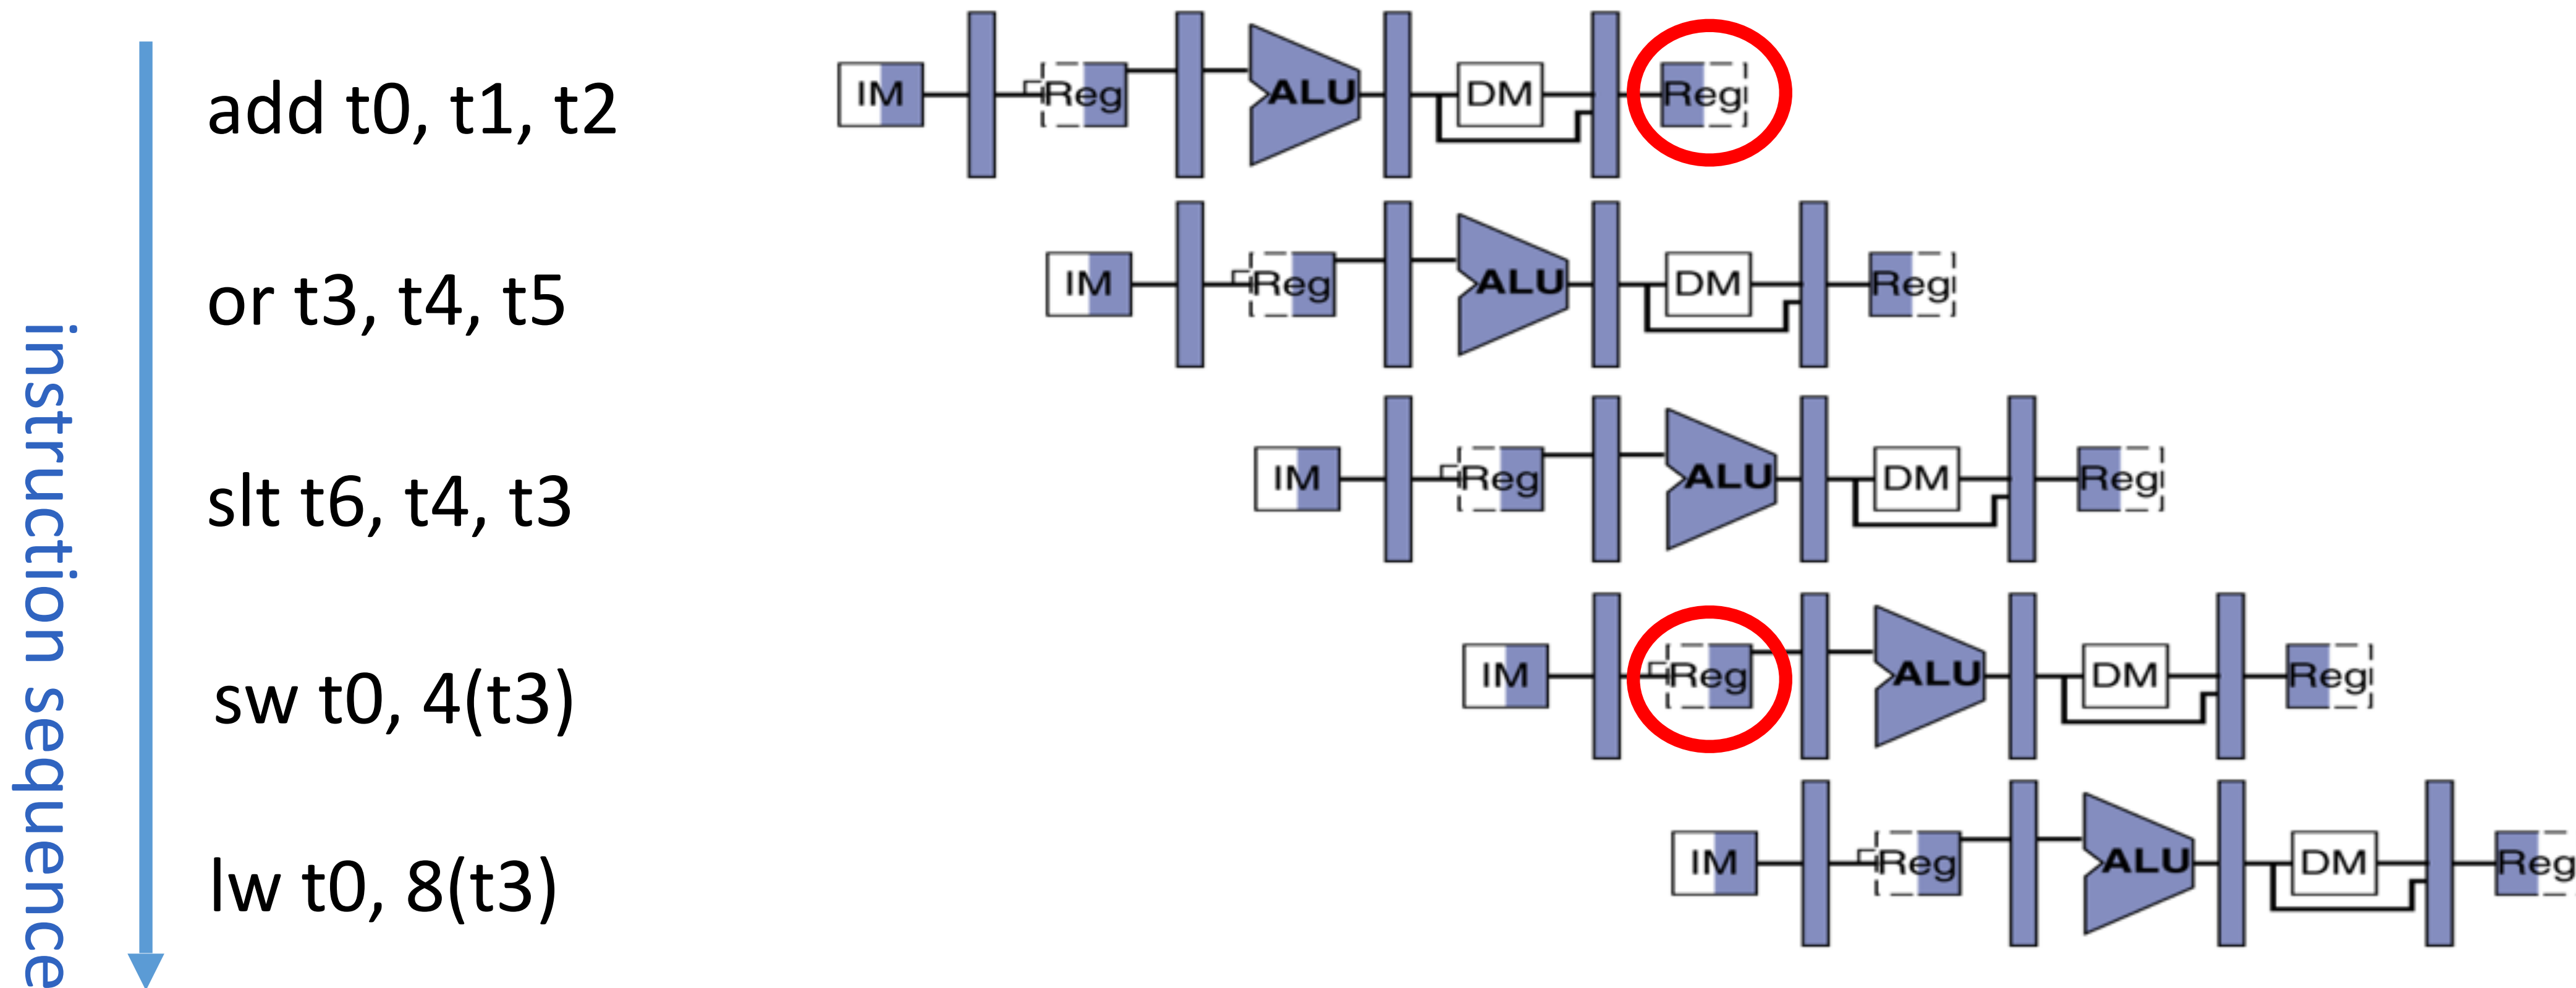
- Conflict for use of a resource
- In RISC-V pipeline with a single memory
 - Load/store requires data access
 - Without separate memories, instruction fetch would have to stall for that cycle
 - All other operations in pipeline would have to wait
- Pipelined datapaths require separate instruction/data memories
 - Or at least separate instruction/data **caches**
- **Multi-ported register file**
- RISC ISAs (including RISC-V) designed to avoid structural hazards
 - e.g. at most one memory access/instruction
 - **limited operands per instruction**

Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - **Data**
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Data Hazard: Register Access

- Separate ports, but what if write to same value as read?
- Does **sw** in the example fetch the old or new value?



Register Access Policy

instruction sequence

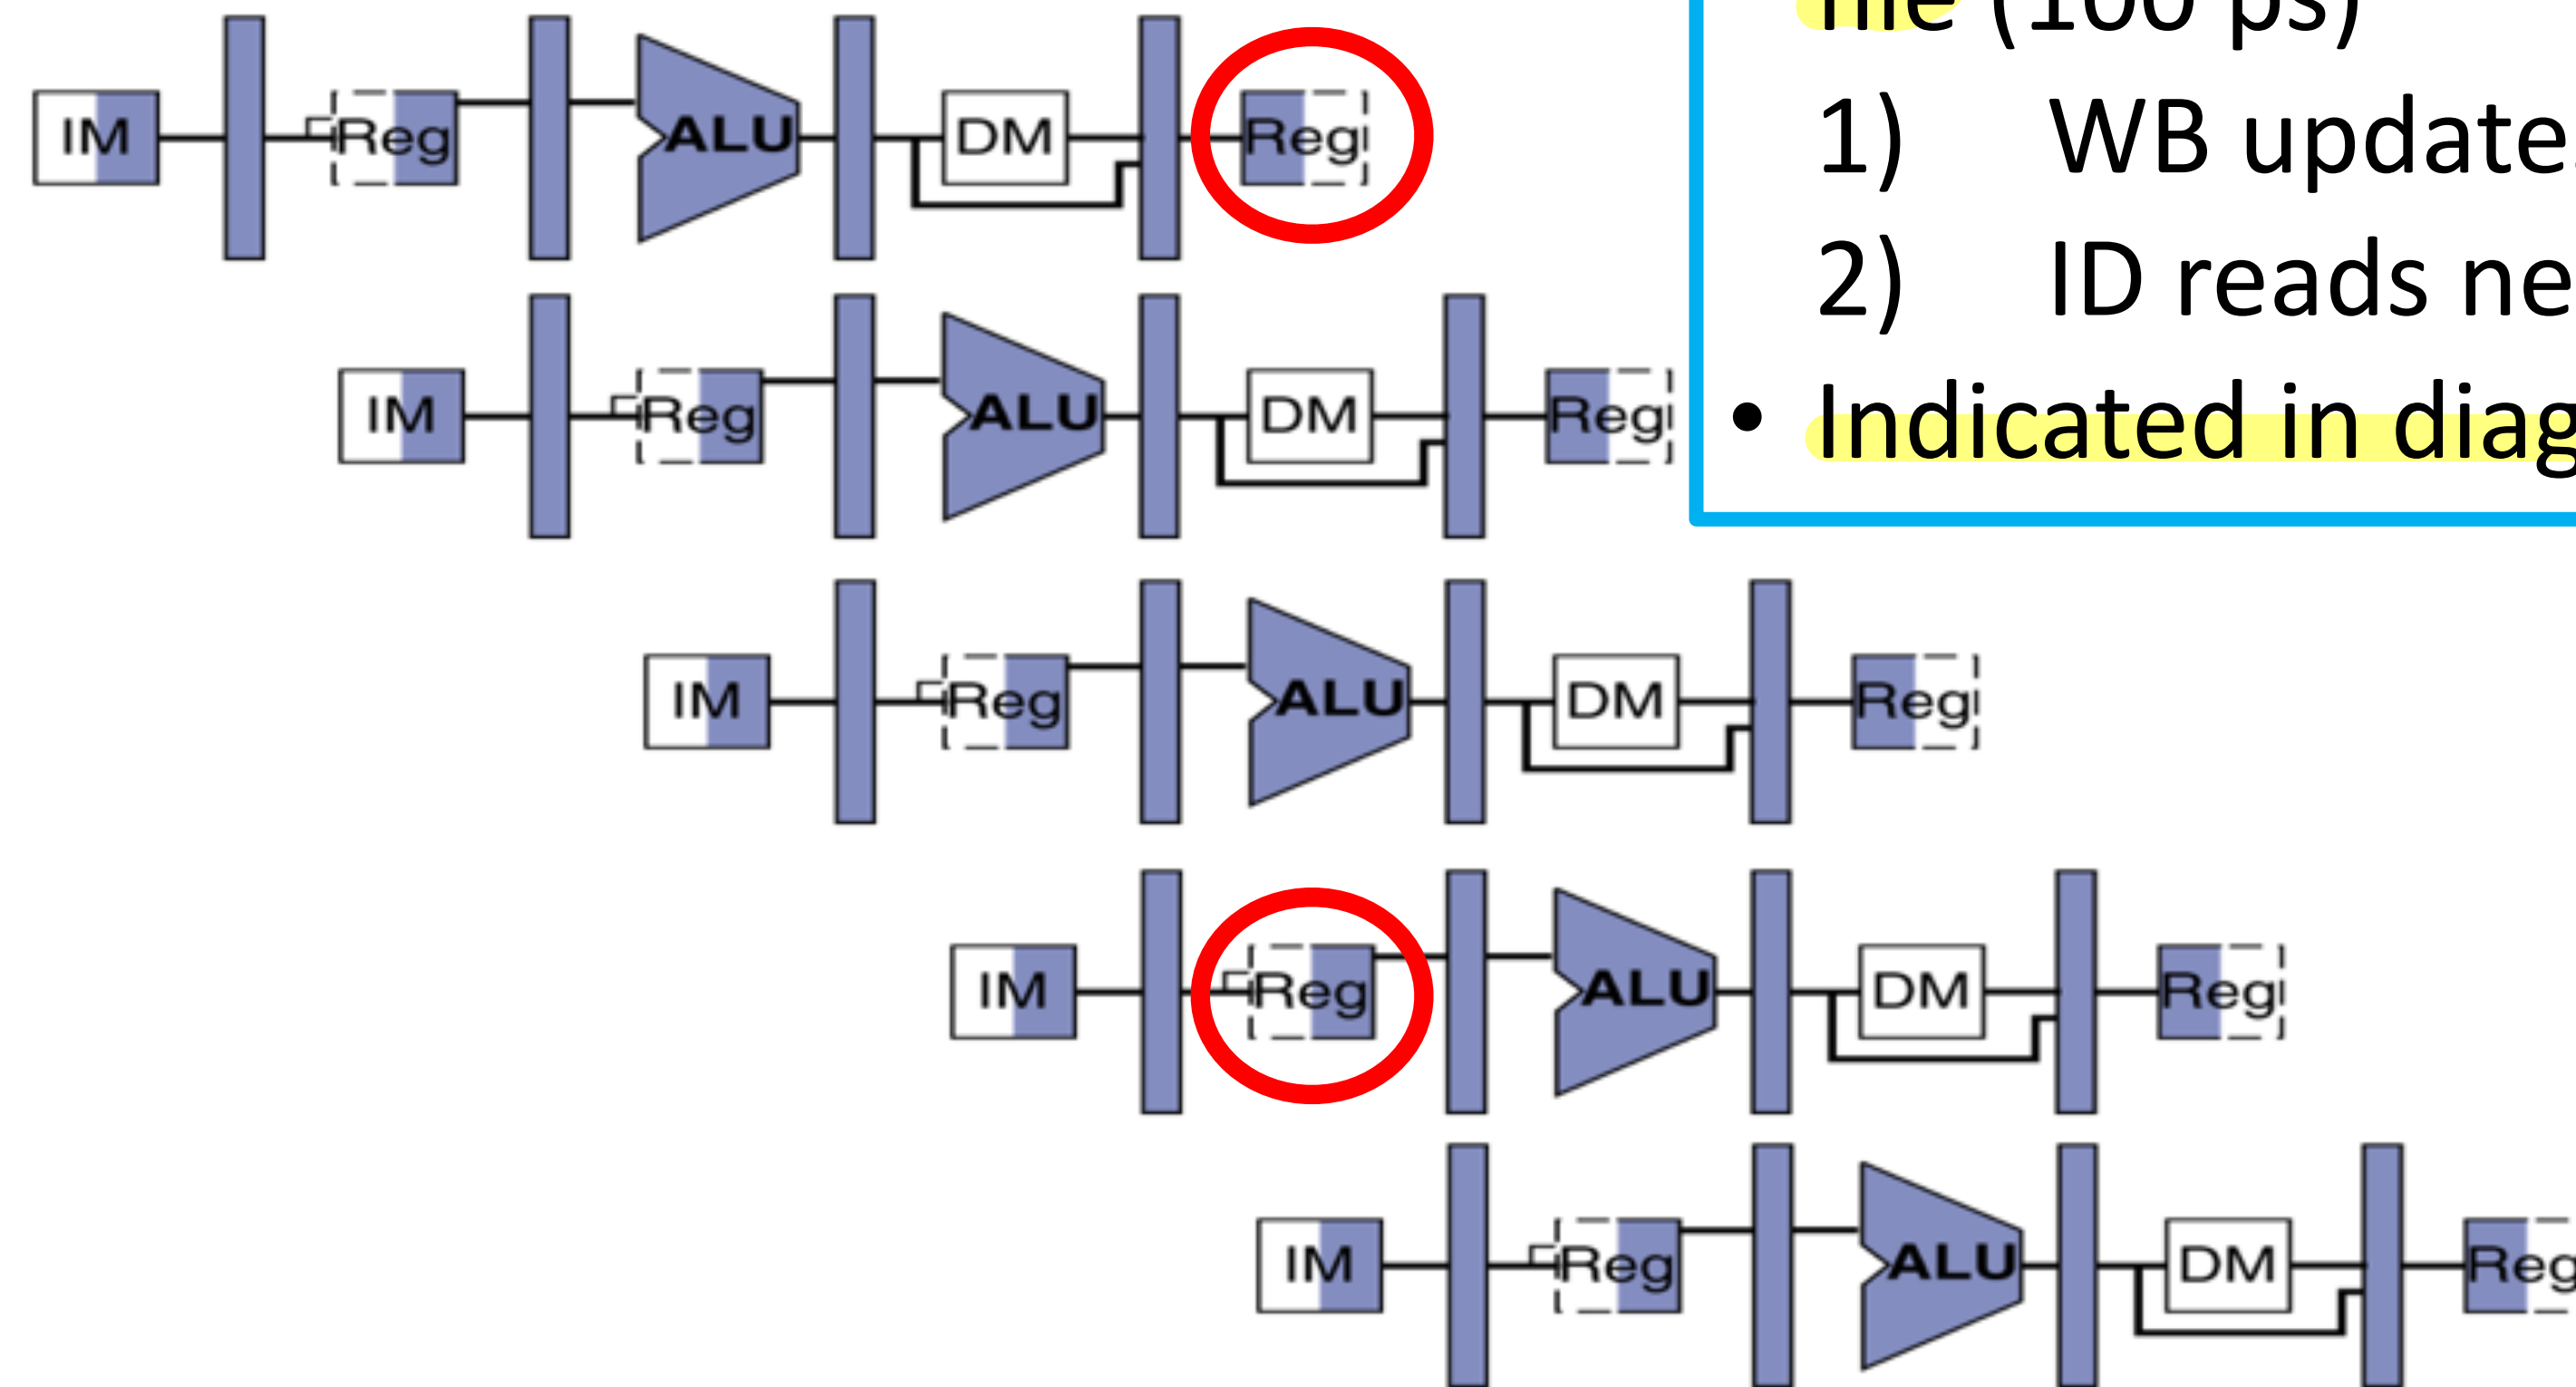
add t0, t1, t2

or t3, t4, t5

slt t6, t4, t3

sw t0, 4(t3)

lw t0, 8(t3)



- Exploit high speed of register file (100 ps)
 - 1) WB updates value
 - 2) ID reads new value
- Indicated in diagram by shading

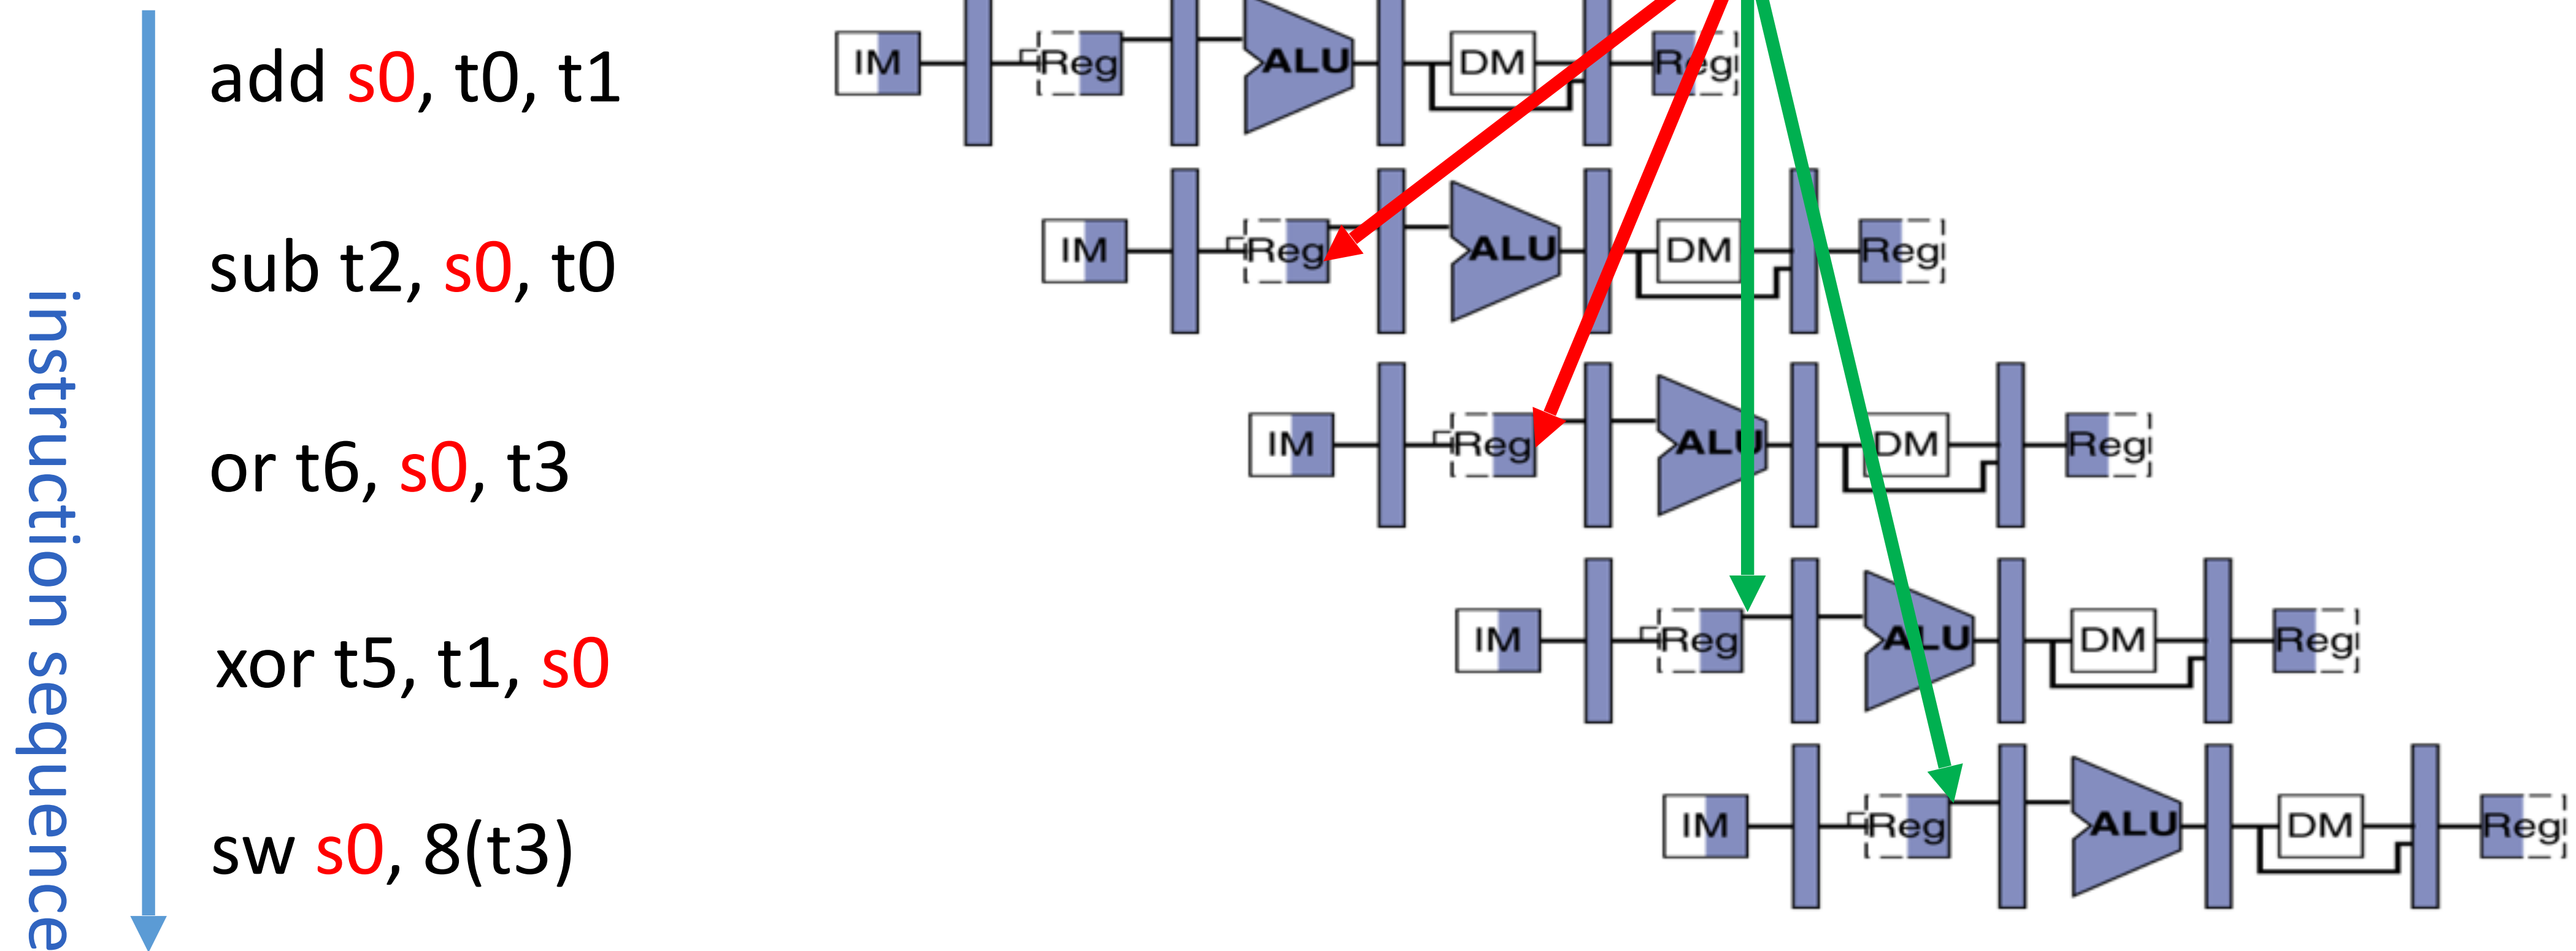
Might not always be possible to write then read in same cycle, especially in high-frequency designs.

Data Hazard: ALU Result

s0 holds "5" then add instr changes s0 to "9"

Value of s0

5	5	5	5	5/9	9	9	9	9
---	---	---	---	-----	---	---	---	---

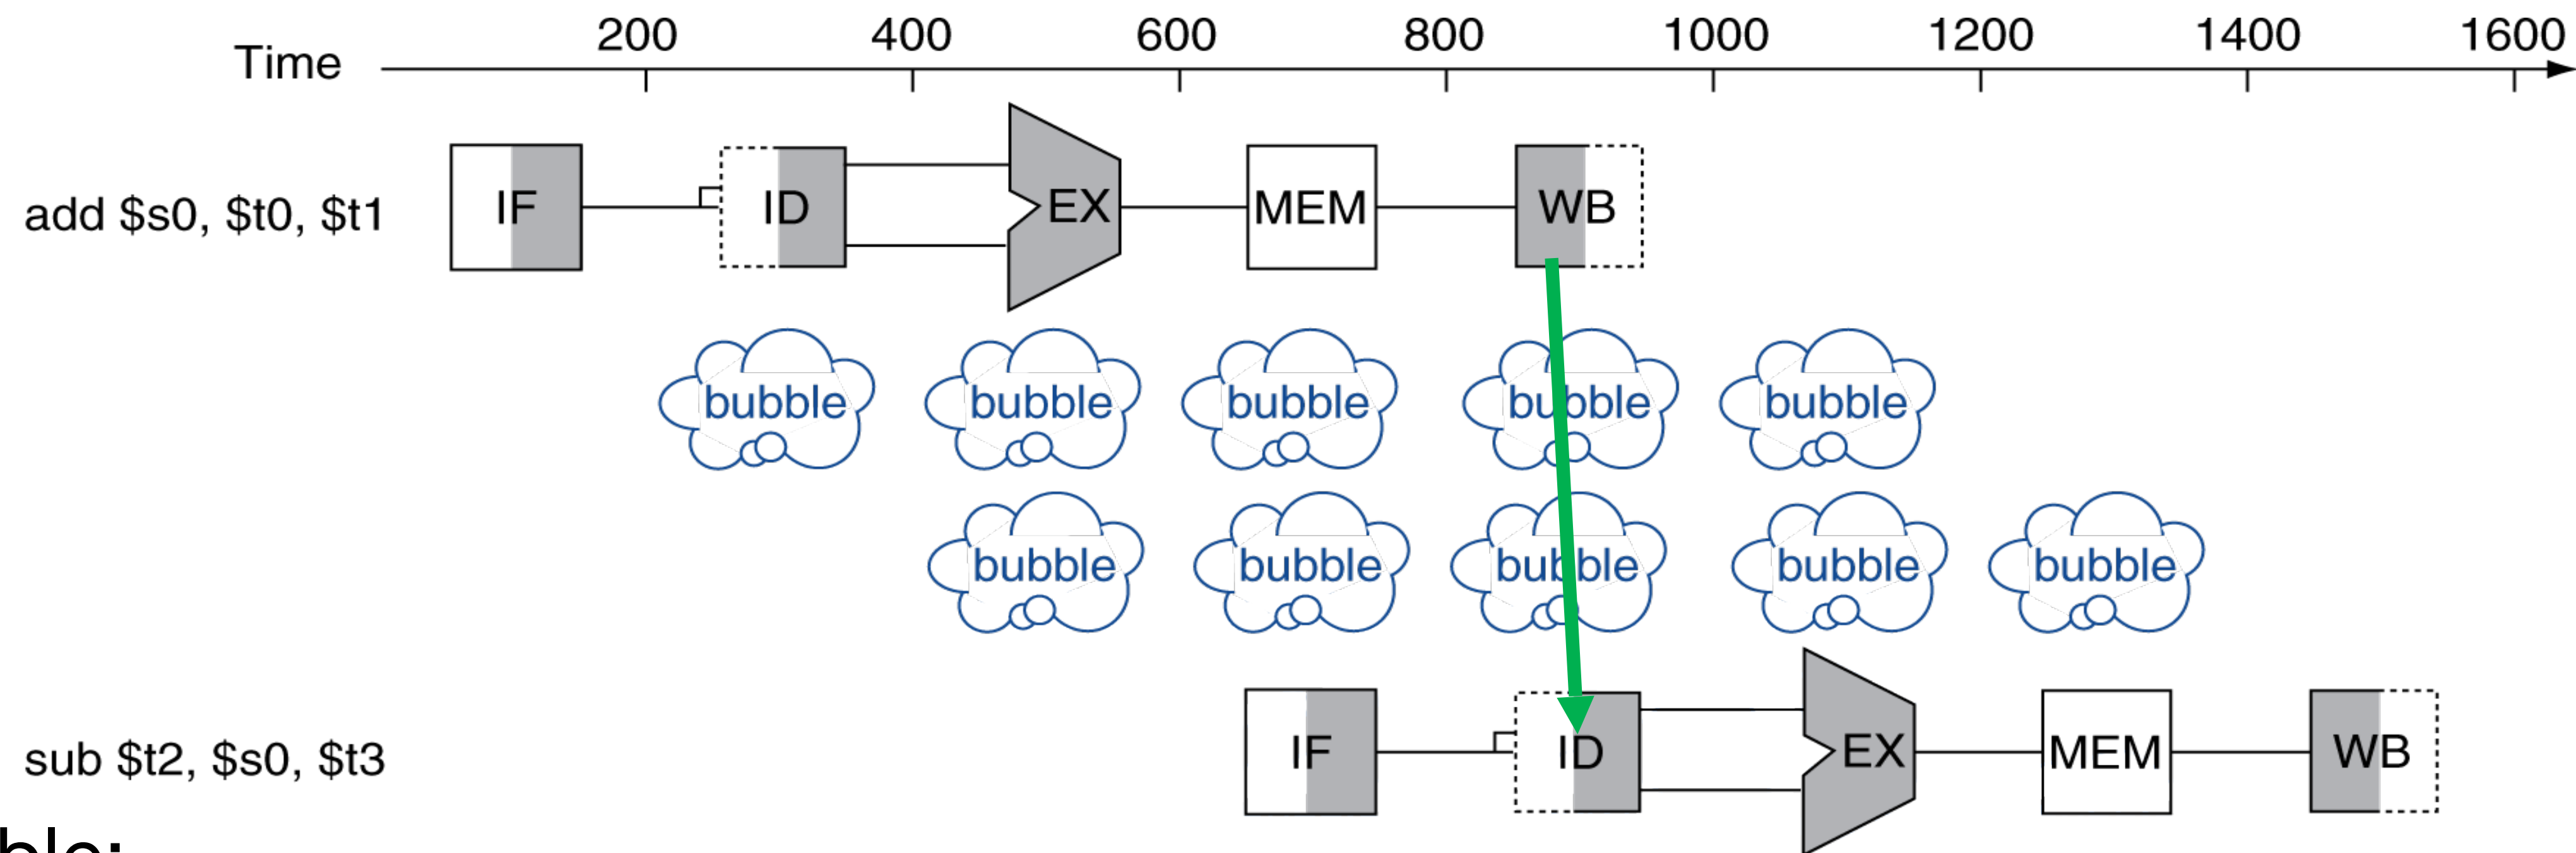


Without some fix, **sub** and **or** will calculate wrong result!

Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction

- add **s0**, t0, t1
- sub t2, **s0**, t3

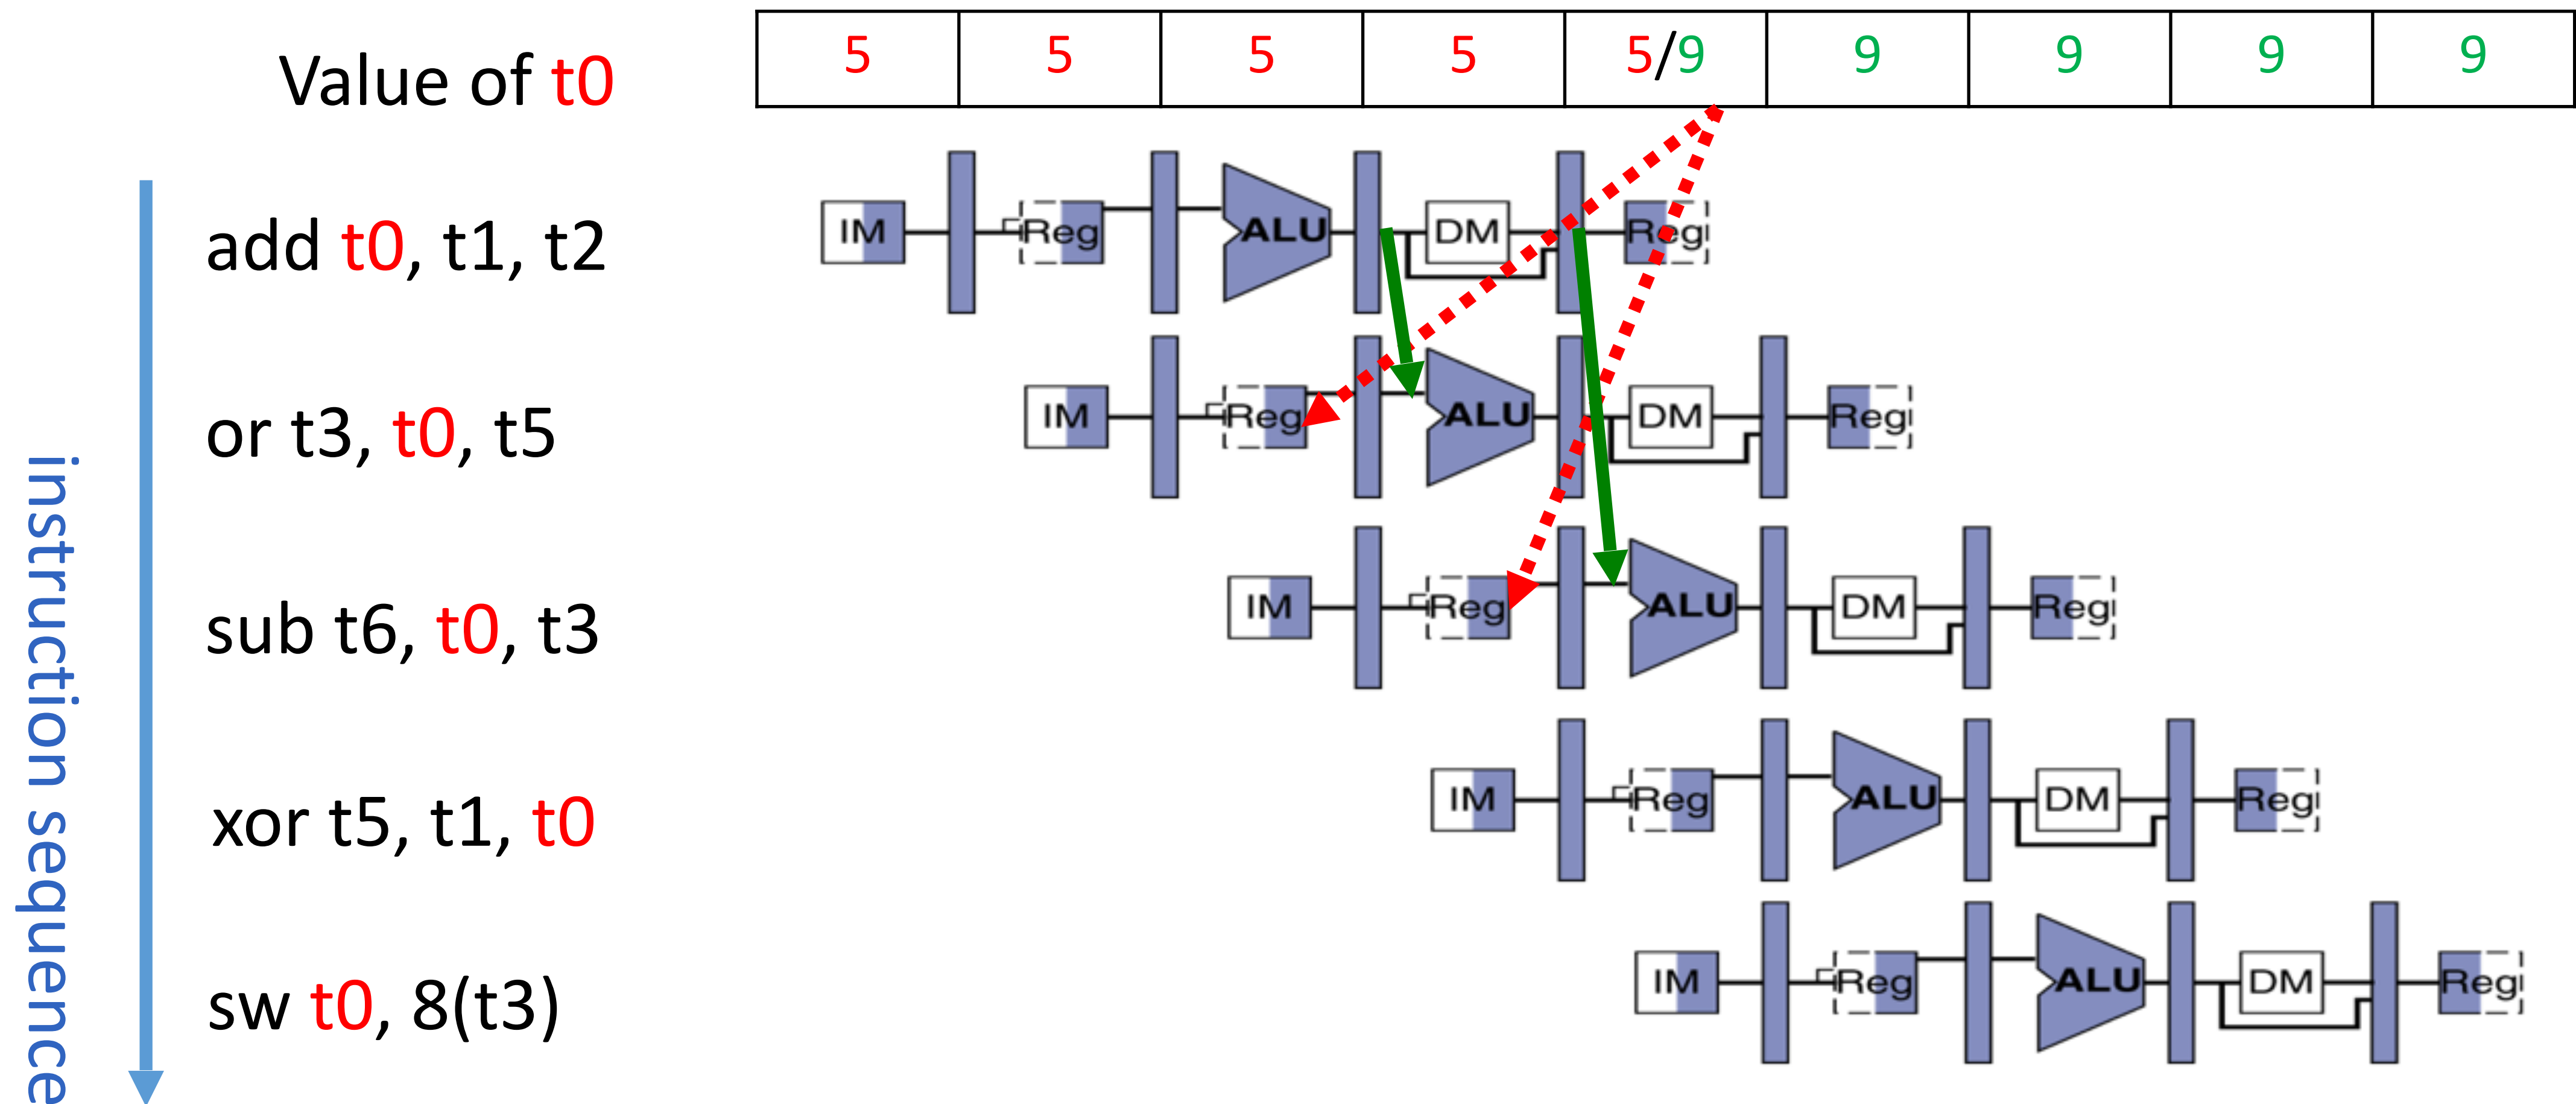


- Bubble:
 - stall the dependent instruction
 - effectively NOP: affected pipeline stages do “nothing”

Stalls and Performance

- Stalls reduce performance
 - But stalls might be required to get correct results
- Compiler could try to arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

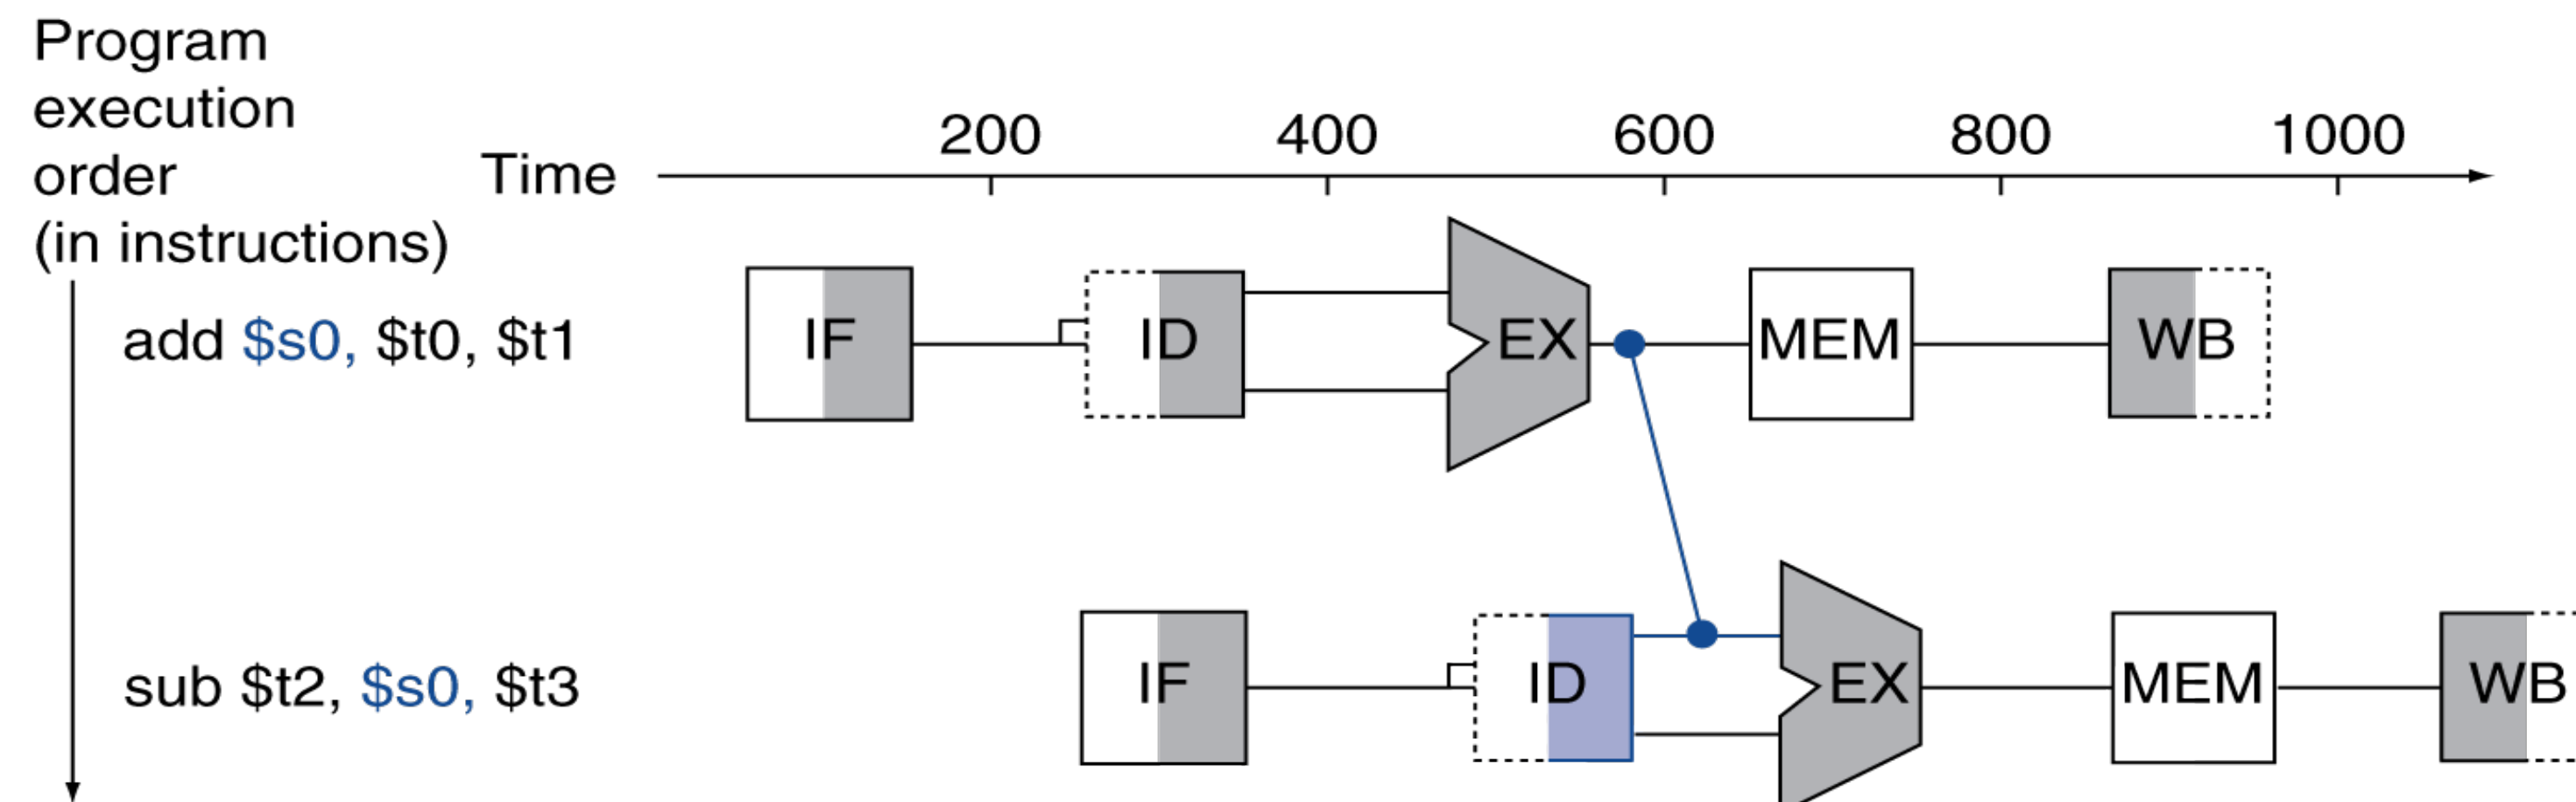
Solution 2: Forwarding



**Forwarding: grab operand from pipeline stage,
rather than register file**

Forwarding (aka Bypassing)

- Access result before it is stored in a register
- Requires extra connections in the datapath

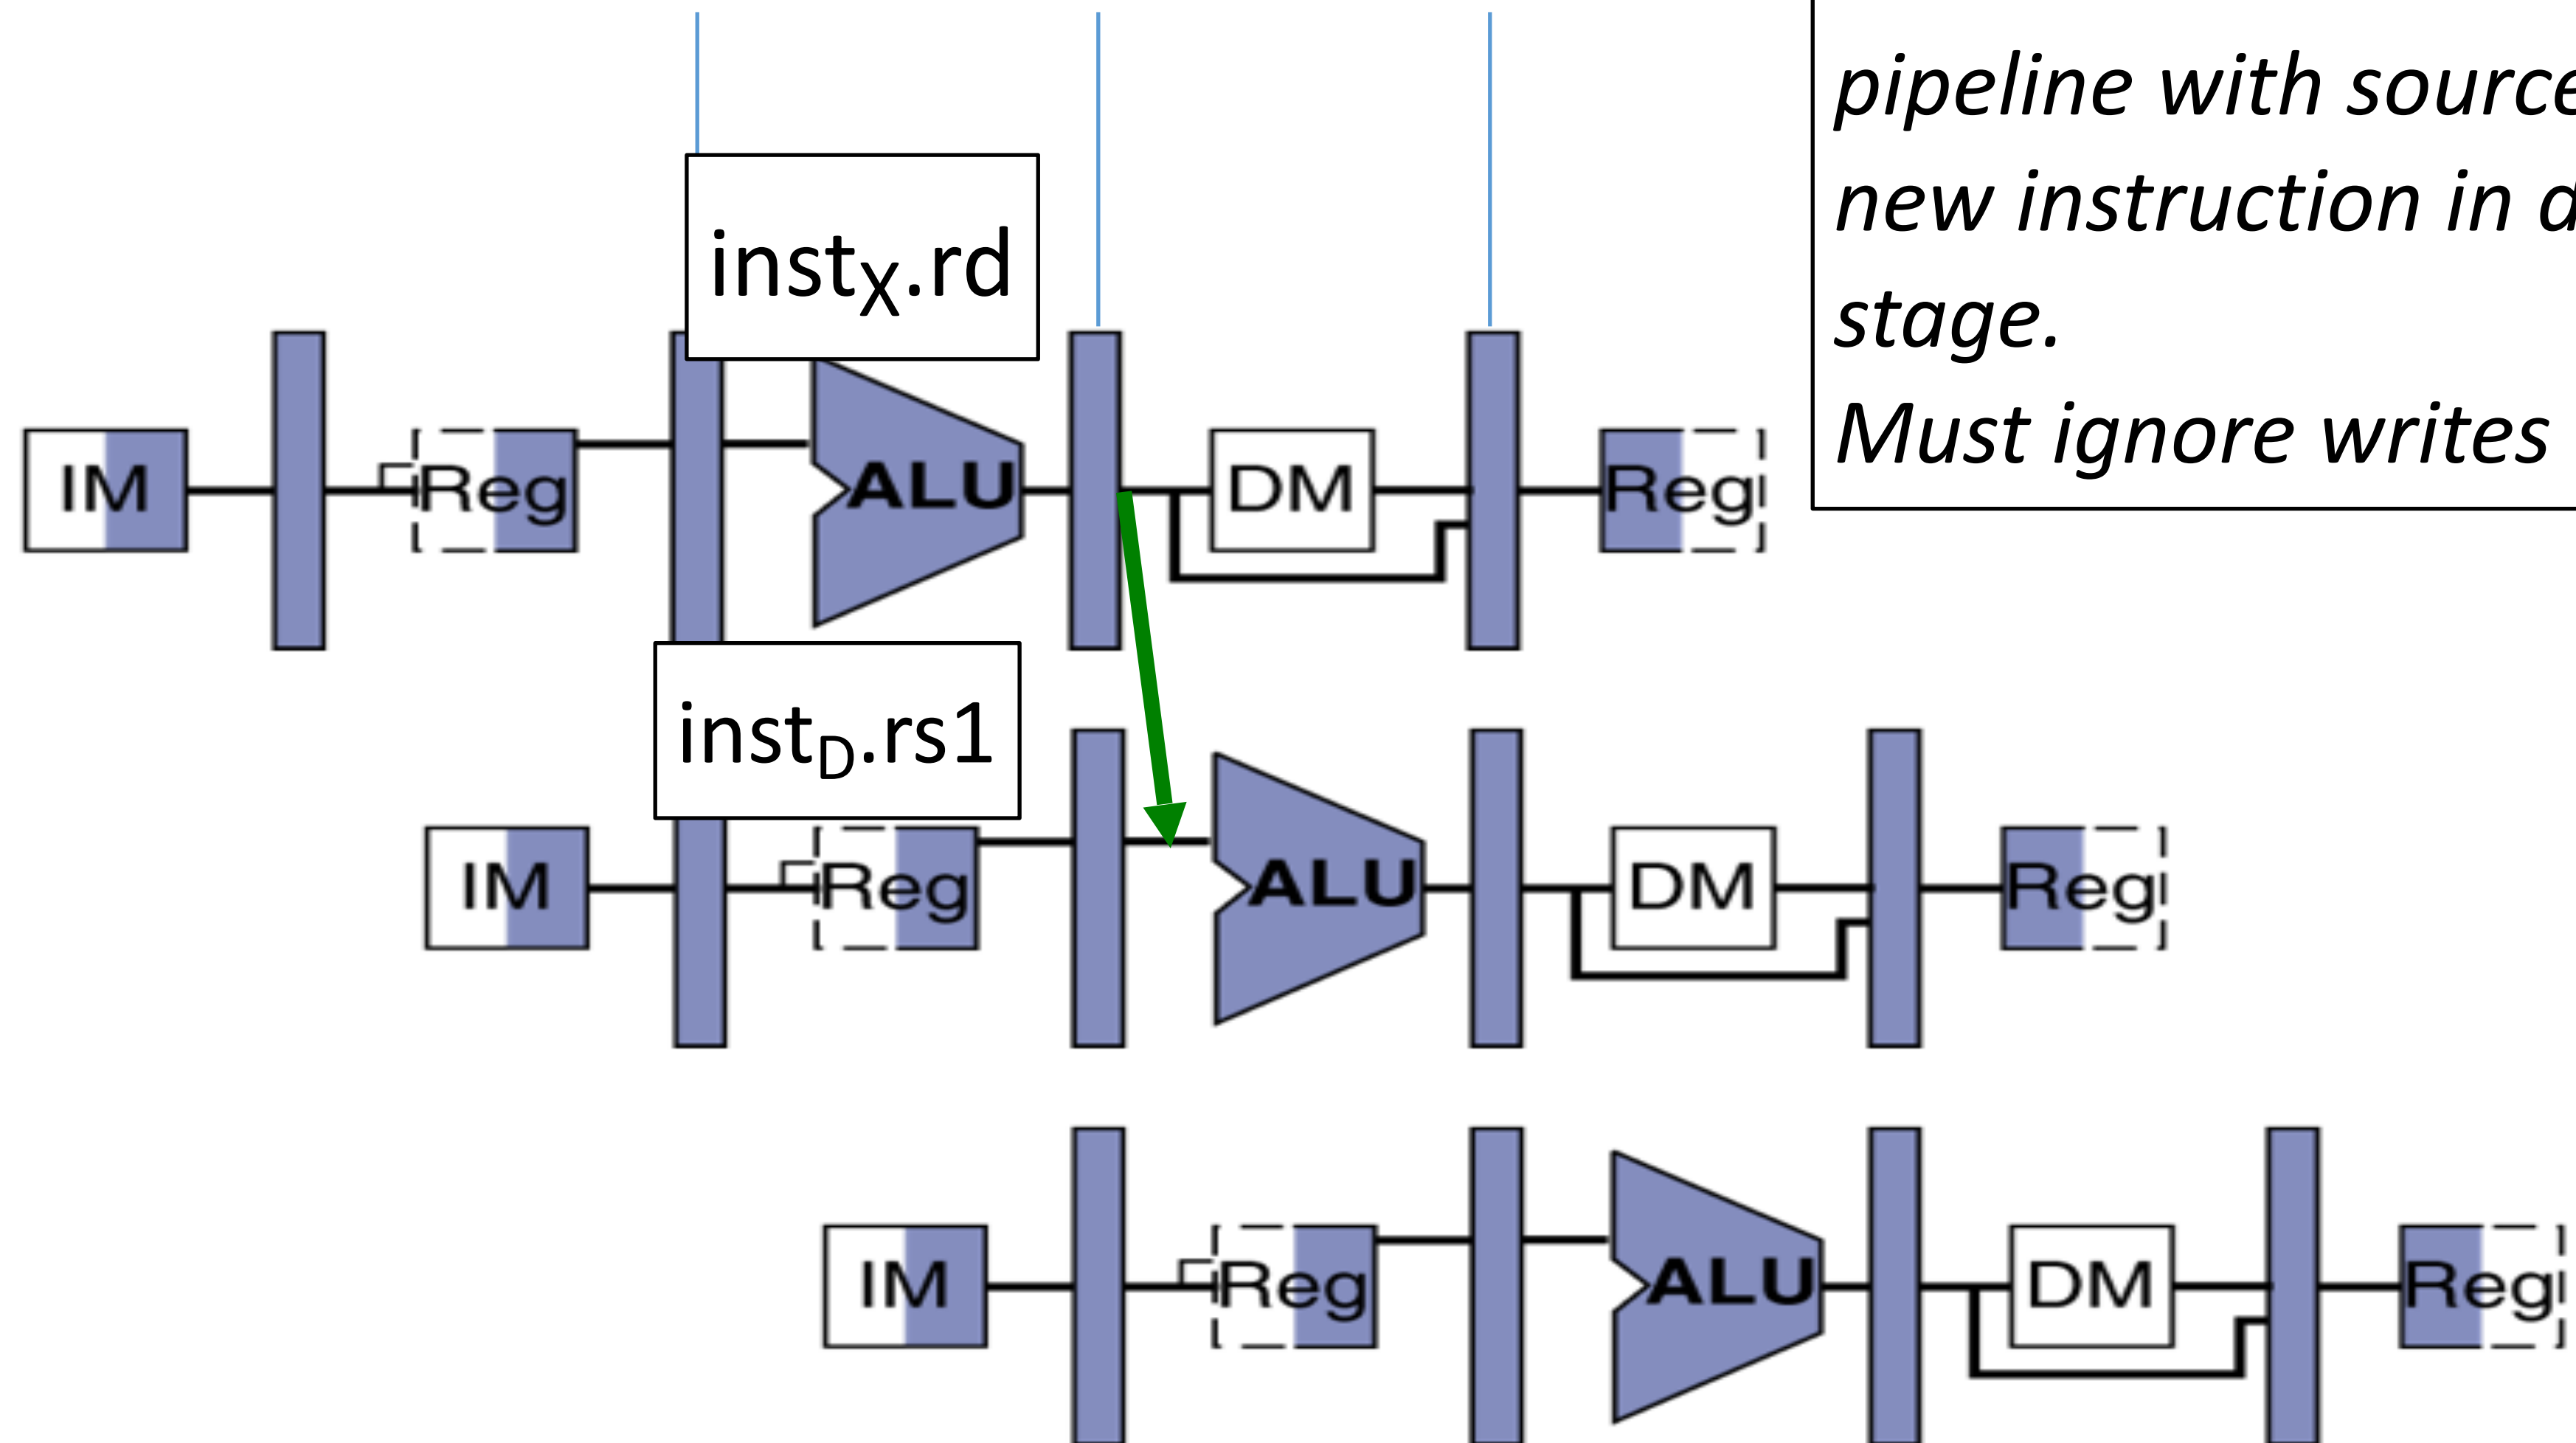


1) Detect Need for Forwarding (example)

add **t0**, t1, t2

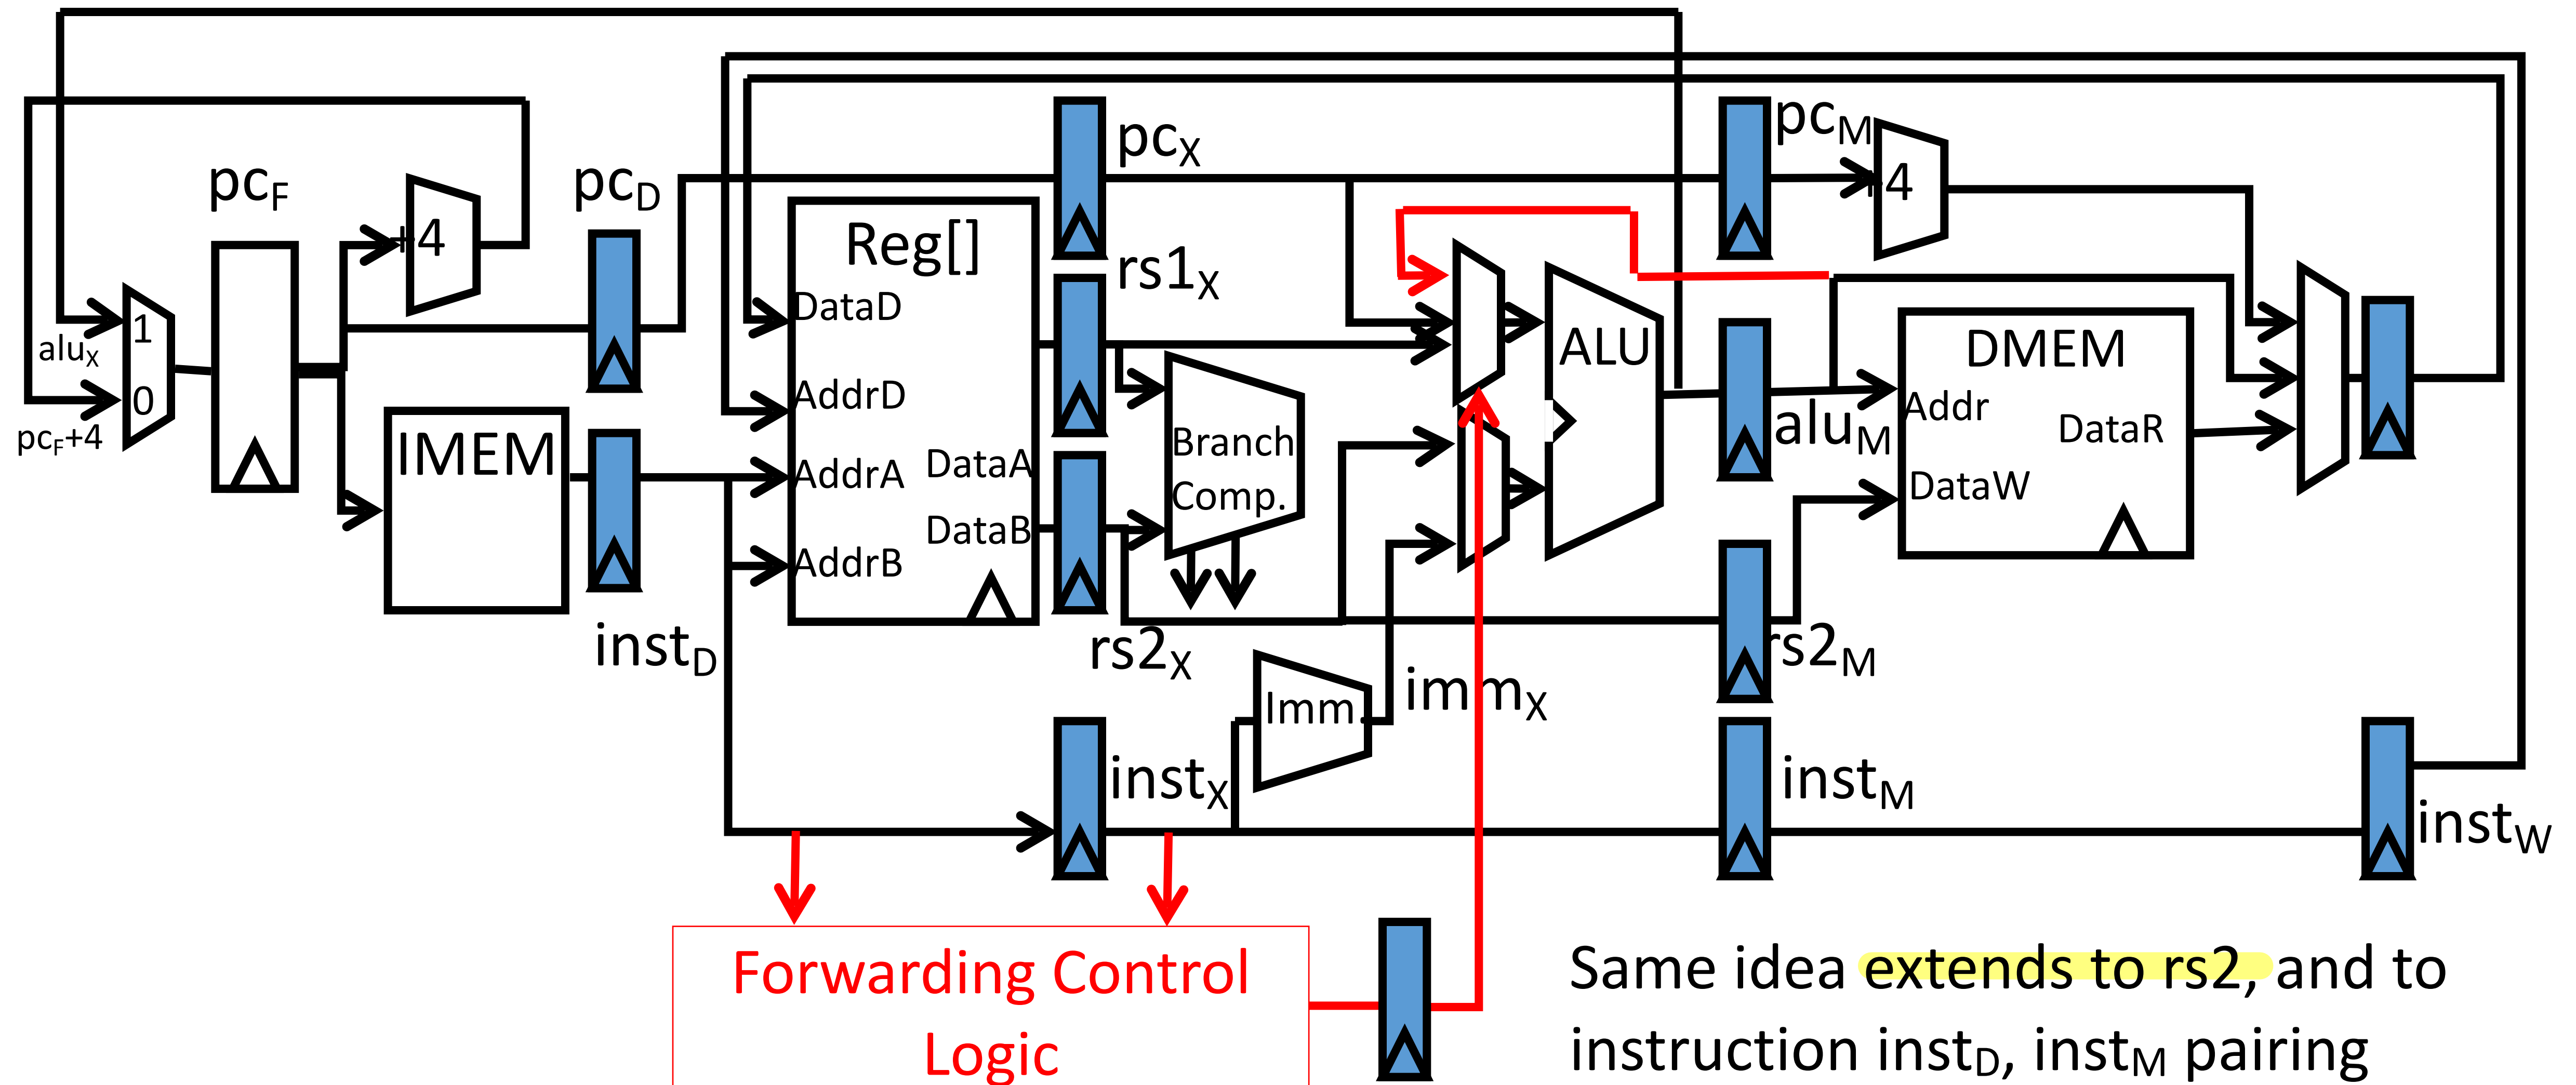
or t3, **t0**, t5

sub t6, **t0**, t3



*Compare destination of
older instructions in
pipeline with sources of
new instruction in decode
stage.
Must ignore writes to x0!*

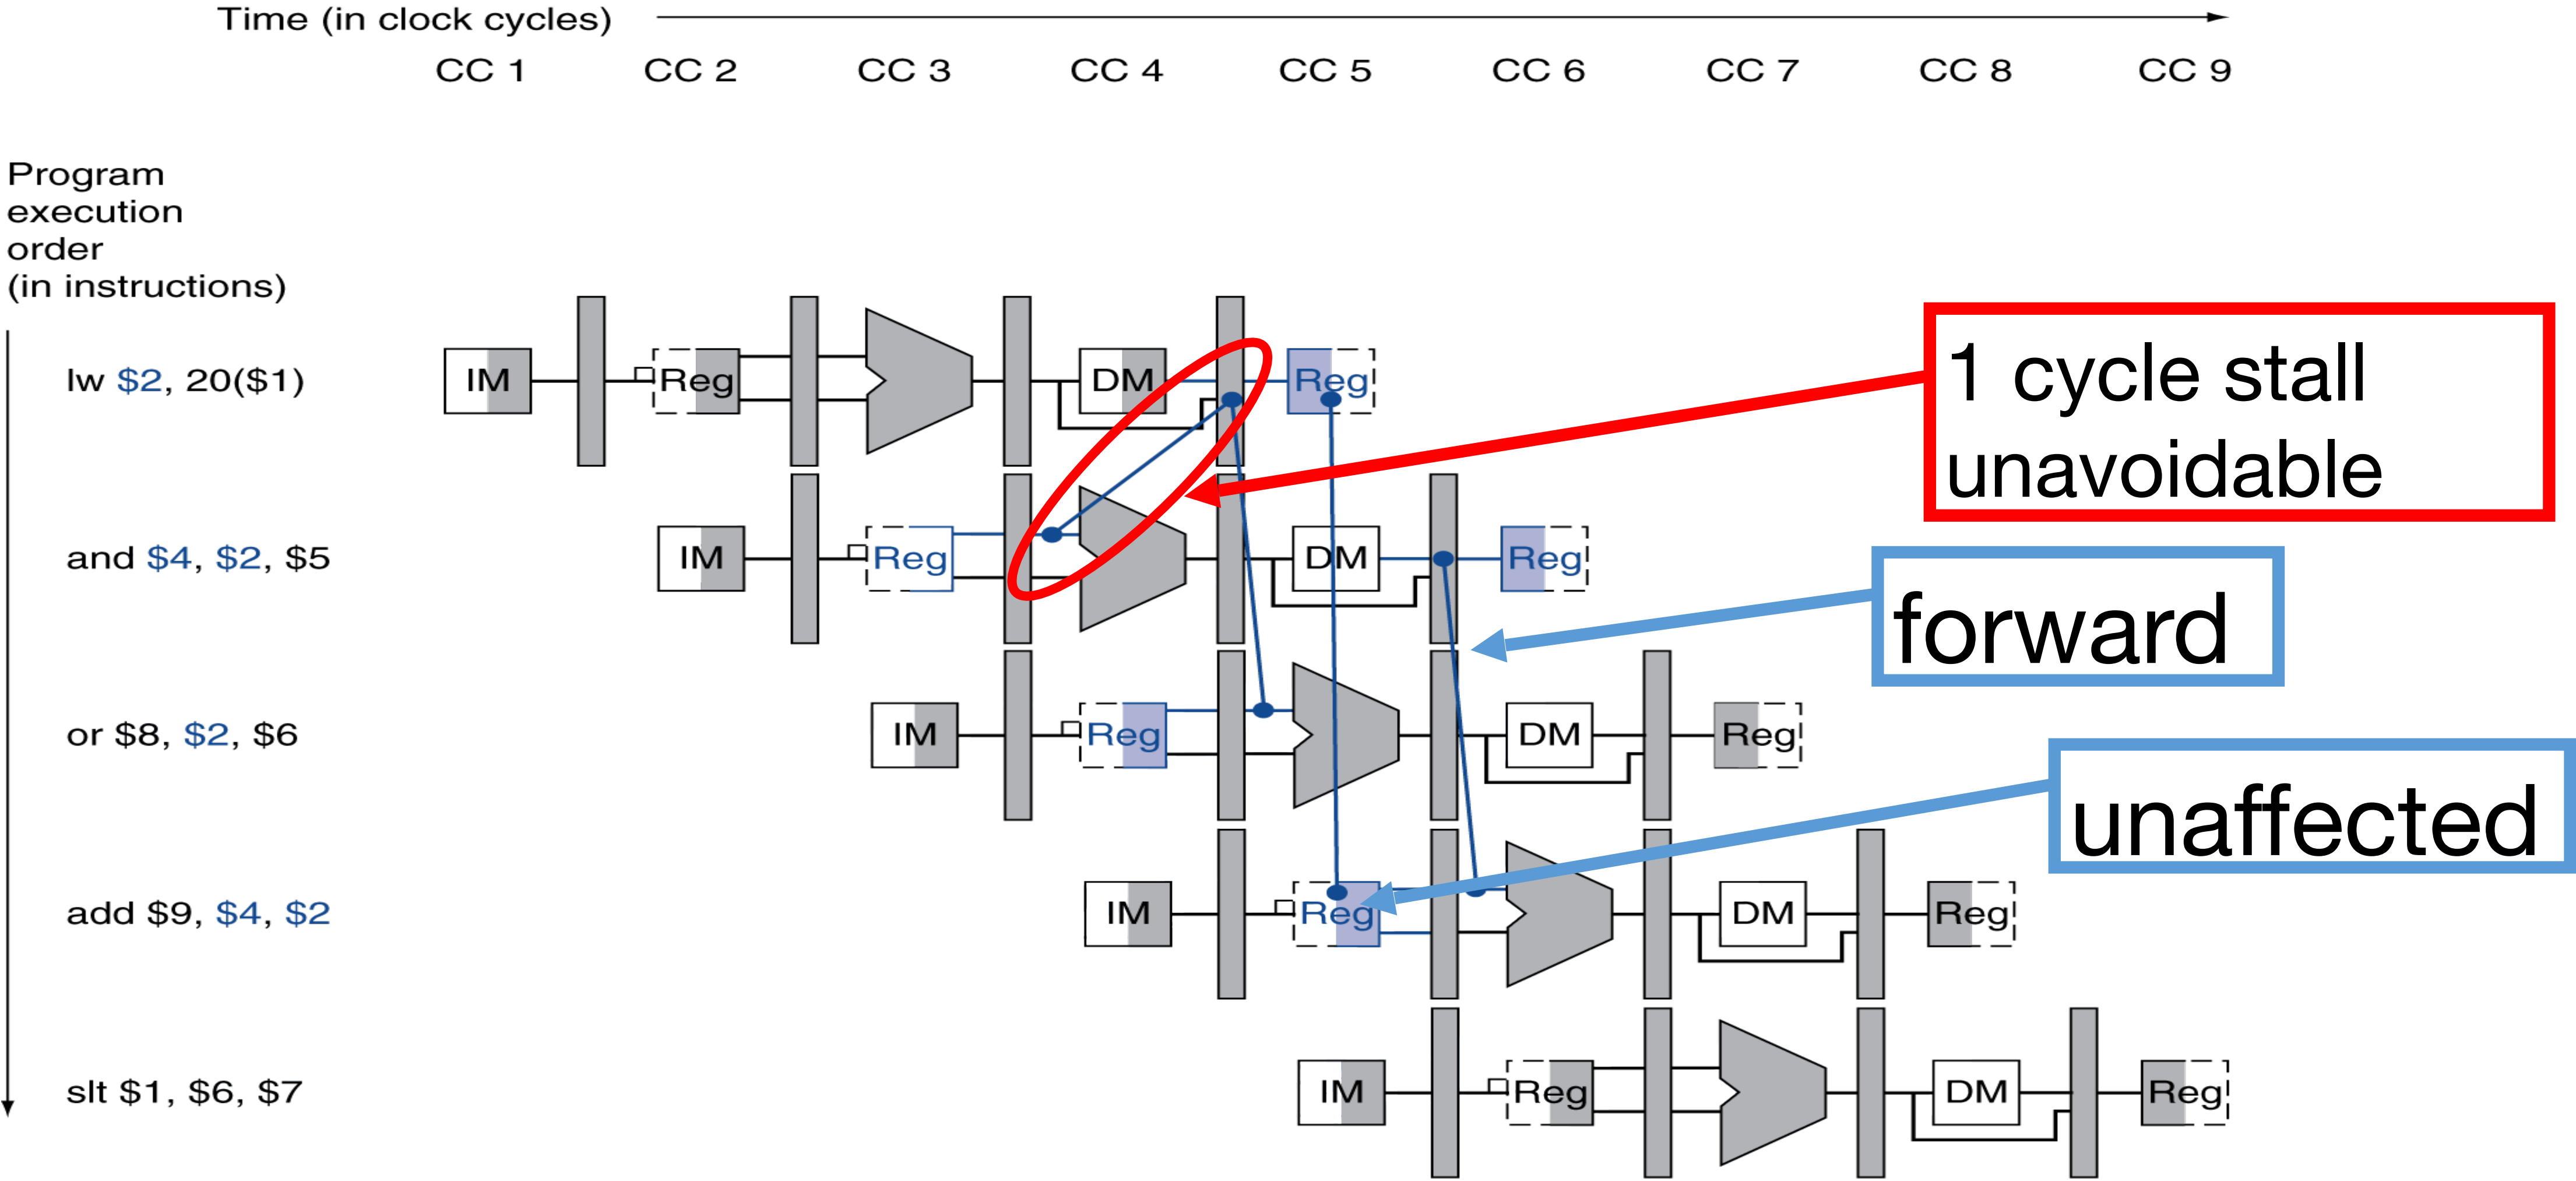
Example Forwarding Path



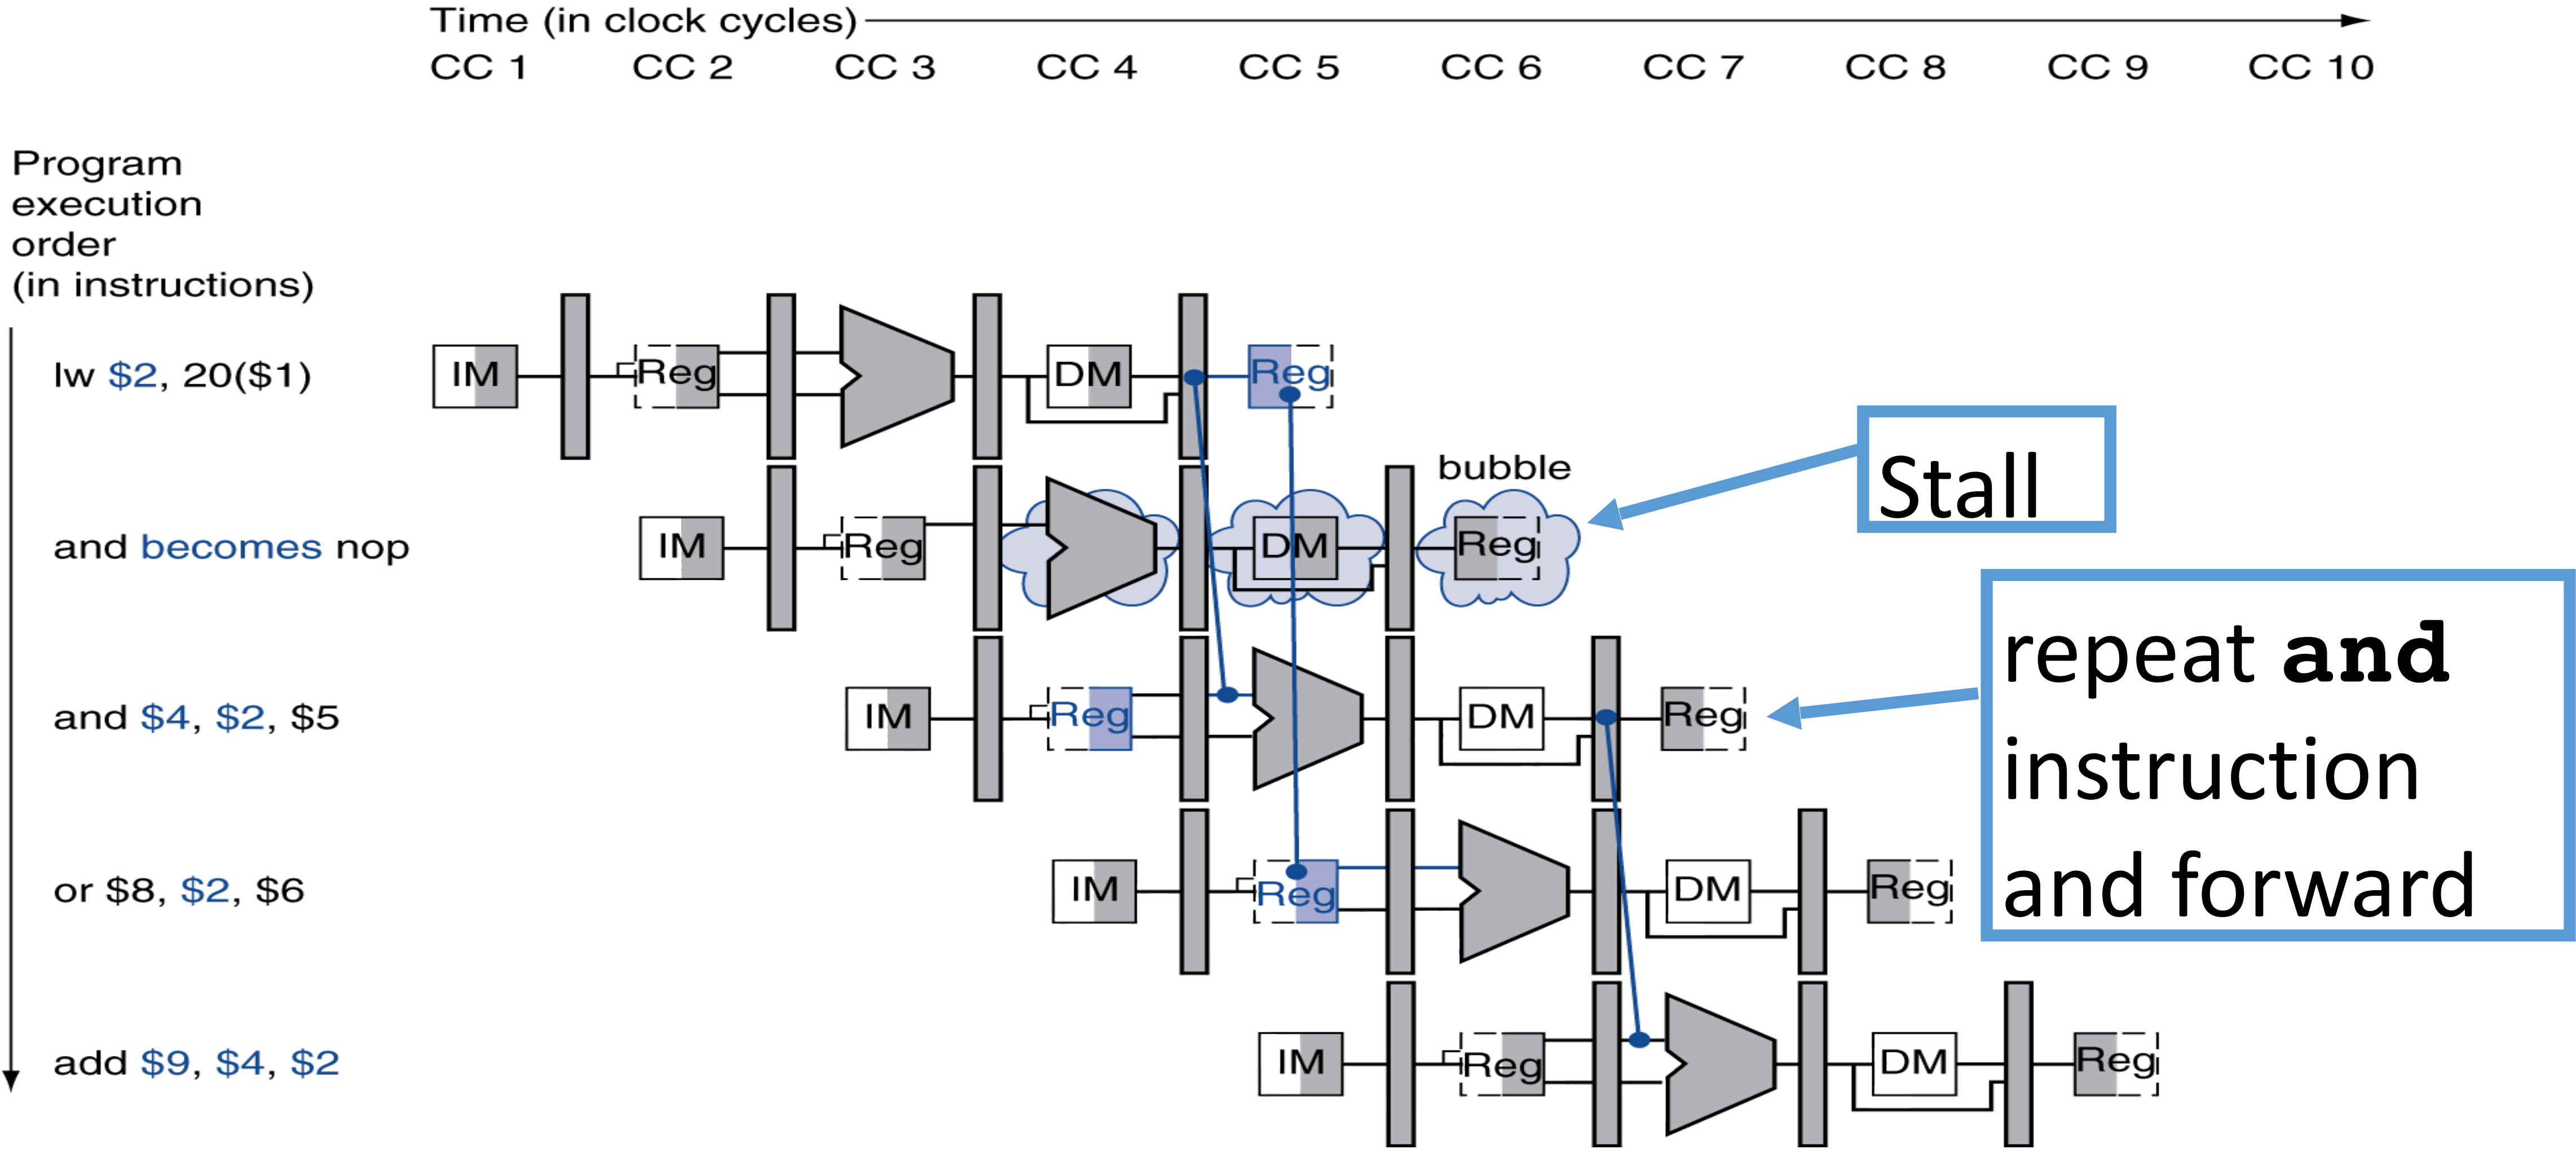
Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Load Data Hazard



Stall Pipeline



1w Data Hazard

- Slot after a load is called a *load delay slot*
 - If that instruction uses the result of the load, then the hardware will stall for one cycle
 - Equivalent to inserting an explicit **nop** in the slot
 - except the latter uses more code space
 - Performance loss!
- Idea:
 - Put unrelated instruction into load delay slot
 - No performance loss!

Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction!
- RISC-V code for $D=A+B$; $E=A+C$;

Original Order:

```
lw    t1, 0(t0)
lw    t2, 4(t0)
add   t3, t1, t2
sw    t3, 12(t0)
lw    t4, 8(t0)
add   t5, t1, t4
sw    t5, 16(t0)
```

Stall!

Stall!

13 cycles

Alternative:

```
lw    t1, 0(t0)
lw    t2, 4(t0)
lw    t4, 8(t0)
add   t3, t1, t2
sw    t3, 12(t0)
add   t5, t1, t4
sw    t5, 16(t0)
```

11 cycles

Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - **Control**
- Superscalar processors

Control Hazards


beq t0, t1, label

sub t2, s0, t5

or t6, s0, t3

xor t5, t1, s0

sw s0, 8(t3)



executed regardless of
branch outcome!

executed regardless of
branch outcome!!!

PC updated reflecting
branch outcome

Observation

- If branch not taken, then instructions fetched sequentially after branch are correct
- If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs

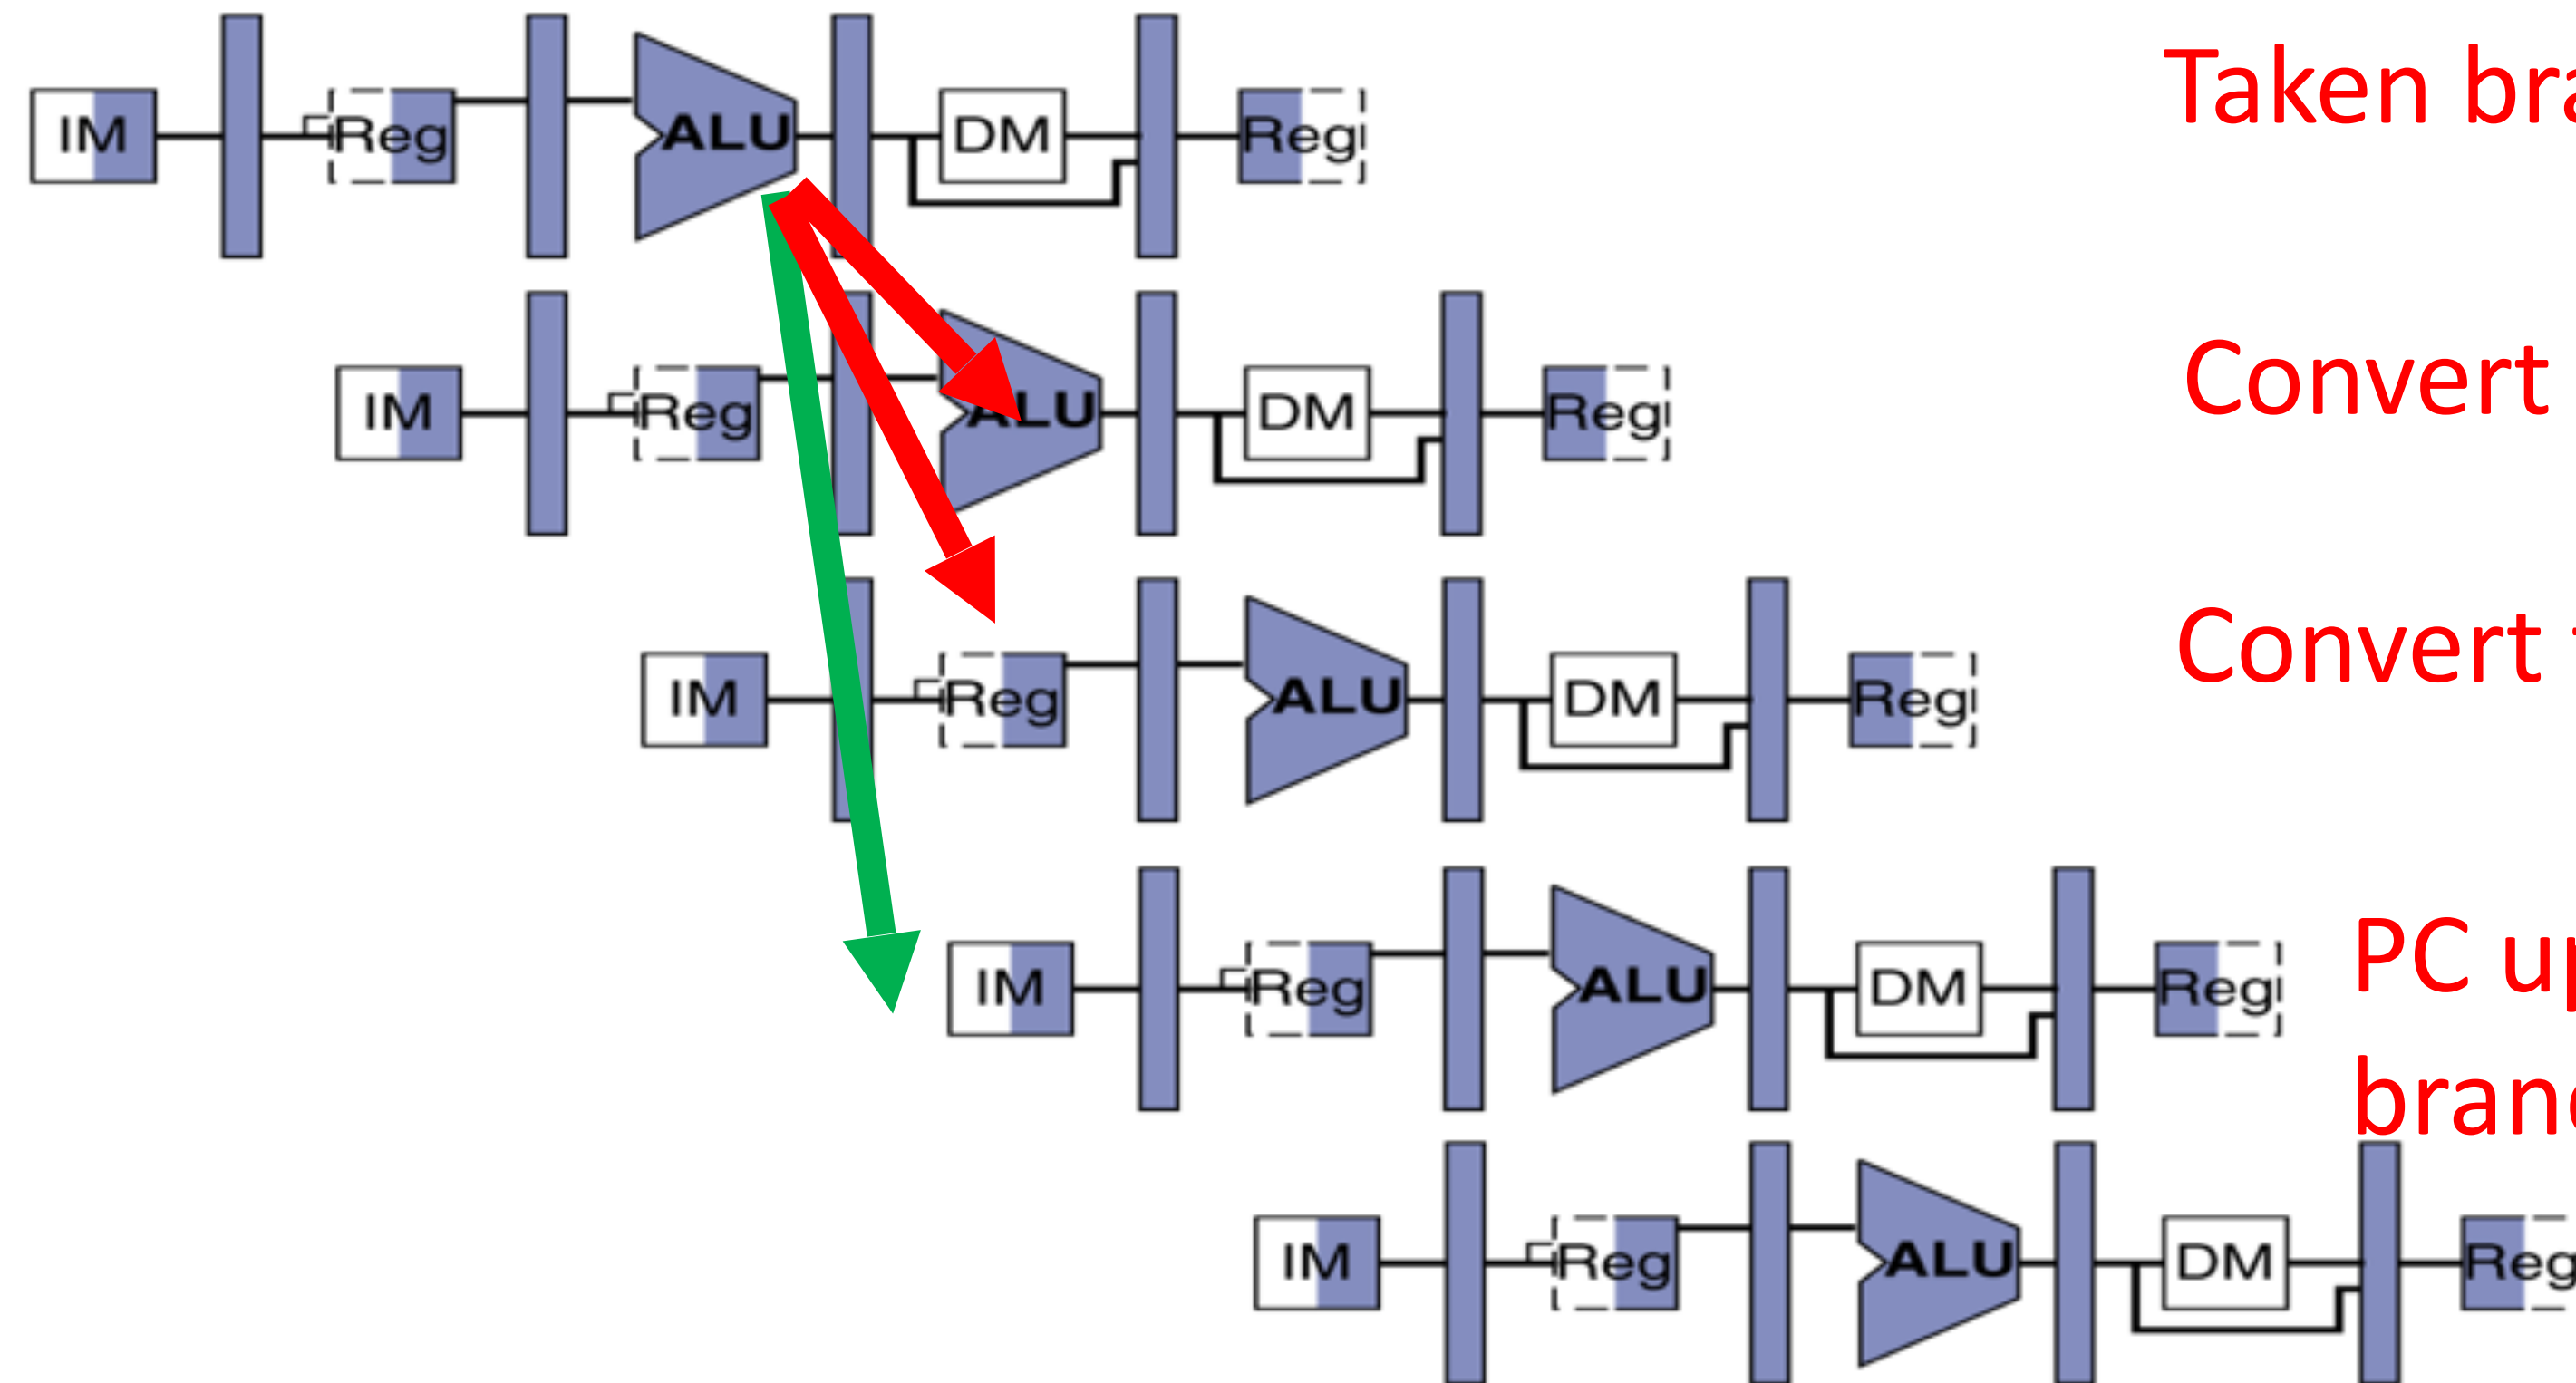
Kill Instructions after Branch if Taken

beq t0, t1, label

sub t2, s0, t5

or t6, s0, t3

label: xxxxxx



Taken branch

Convert to NOP

Convert to NOP

PC updated reflecting
branch outcome

Reducing Branch Penalties

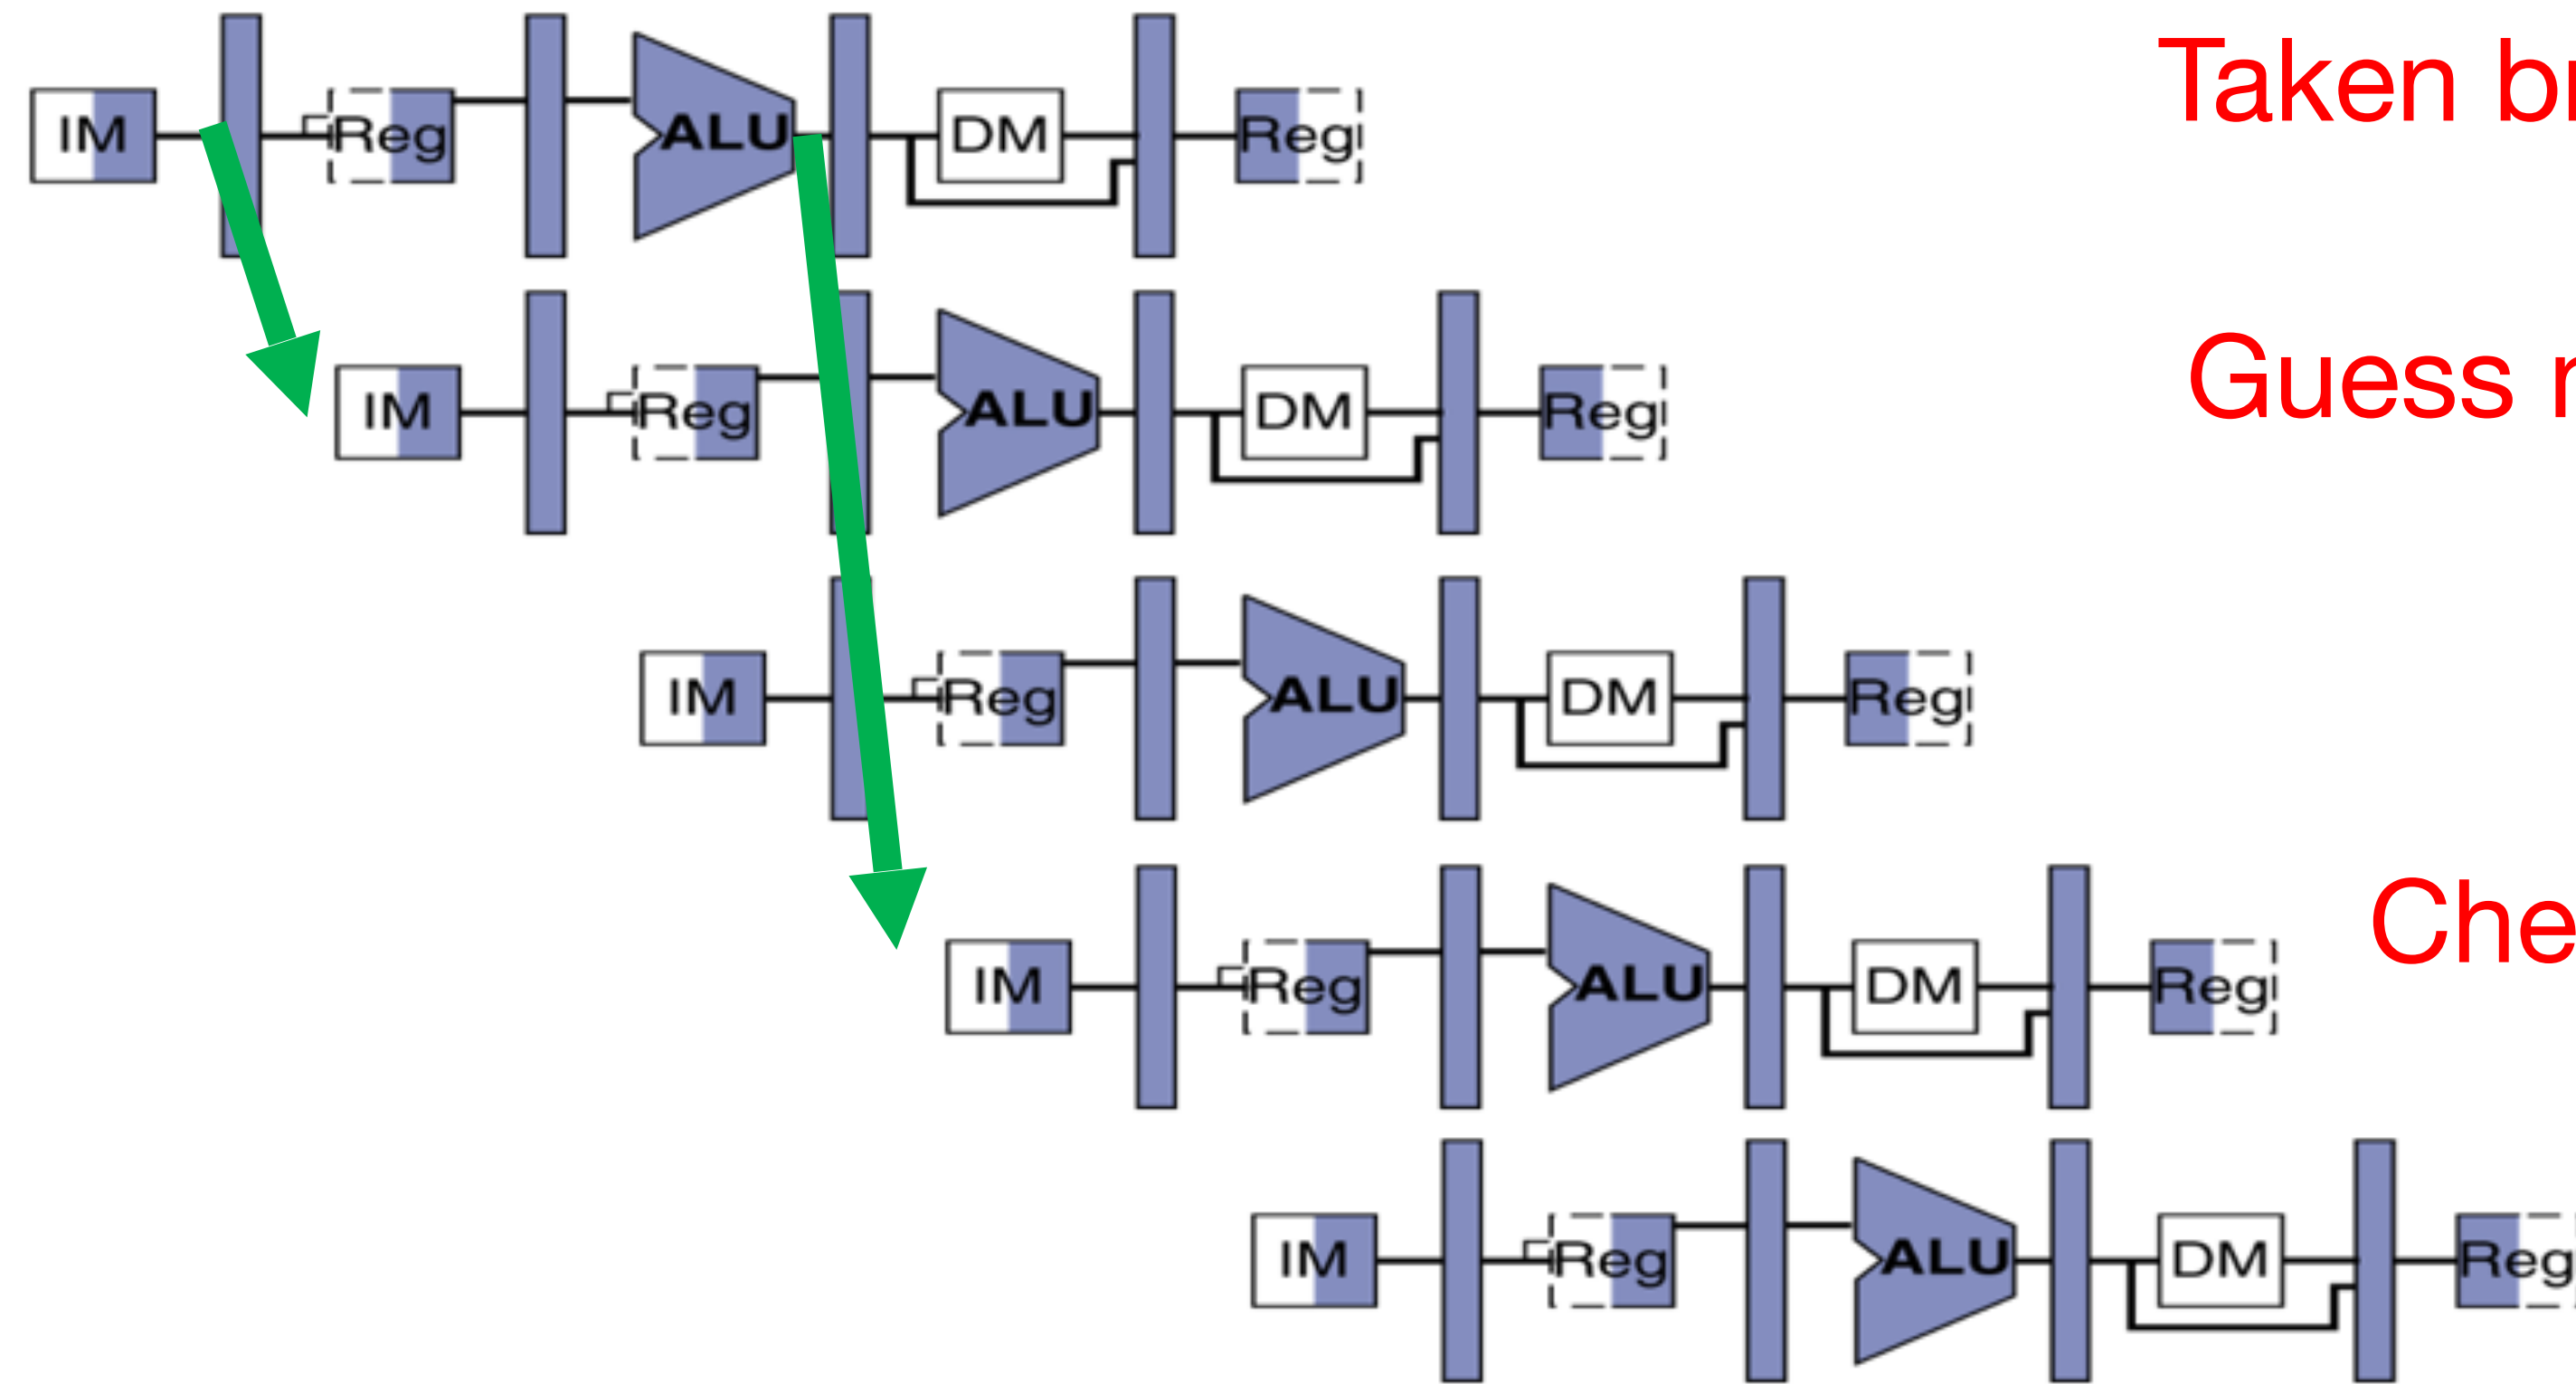
- Every taken branch in simple pipeline costs 2 dead cycles
- To improve performance, use “^{分支预测}branch prediction” to guess which way branch will go earlier in pipeline
- Only flush pipeline if branch prediction was incorrect

Branch Prediction

beq t0, t1, label

label:

.....



Taken branch

Guess next PC!

Check guess correct

Implementing Branch Prediction...

- This is a CS152 topic, but some ideas:
 - Branch prediction is critical for performance on deeper pipelines/superscalar as the "Misprediction penalty" is vastly higher
- Keep a branch prediction buffer/cache: Small memory addressed by the lowest bits of PC
- During instruction decode, if branch: Look up whether branch was taken last time?
 - If yes, compute $PC + \text{offset}$ and fetch that (or store actual branch target address from last time)
 - If no, stick with $PC + 4$
- If branch hasn't been seen before
 - Assume forward branches are not taken, backward branches are taken
- Update state on predictor with results of branch when it is finally calculated

Agenda

- RISC-V Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- **Superscalar processors** 超标量处理器.

Increasing Single Processor Core Performance

1. Clock rate

- Limited by technology and power dissipation

2. Pipelining

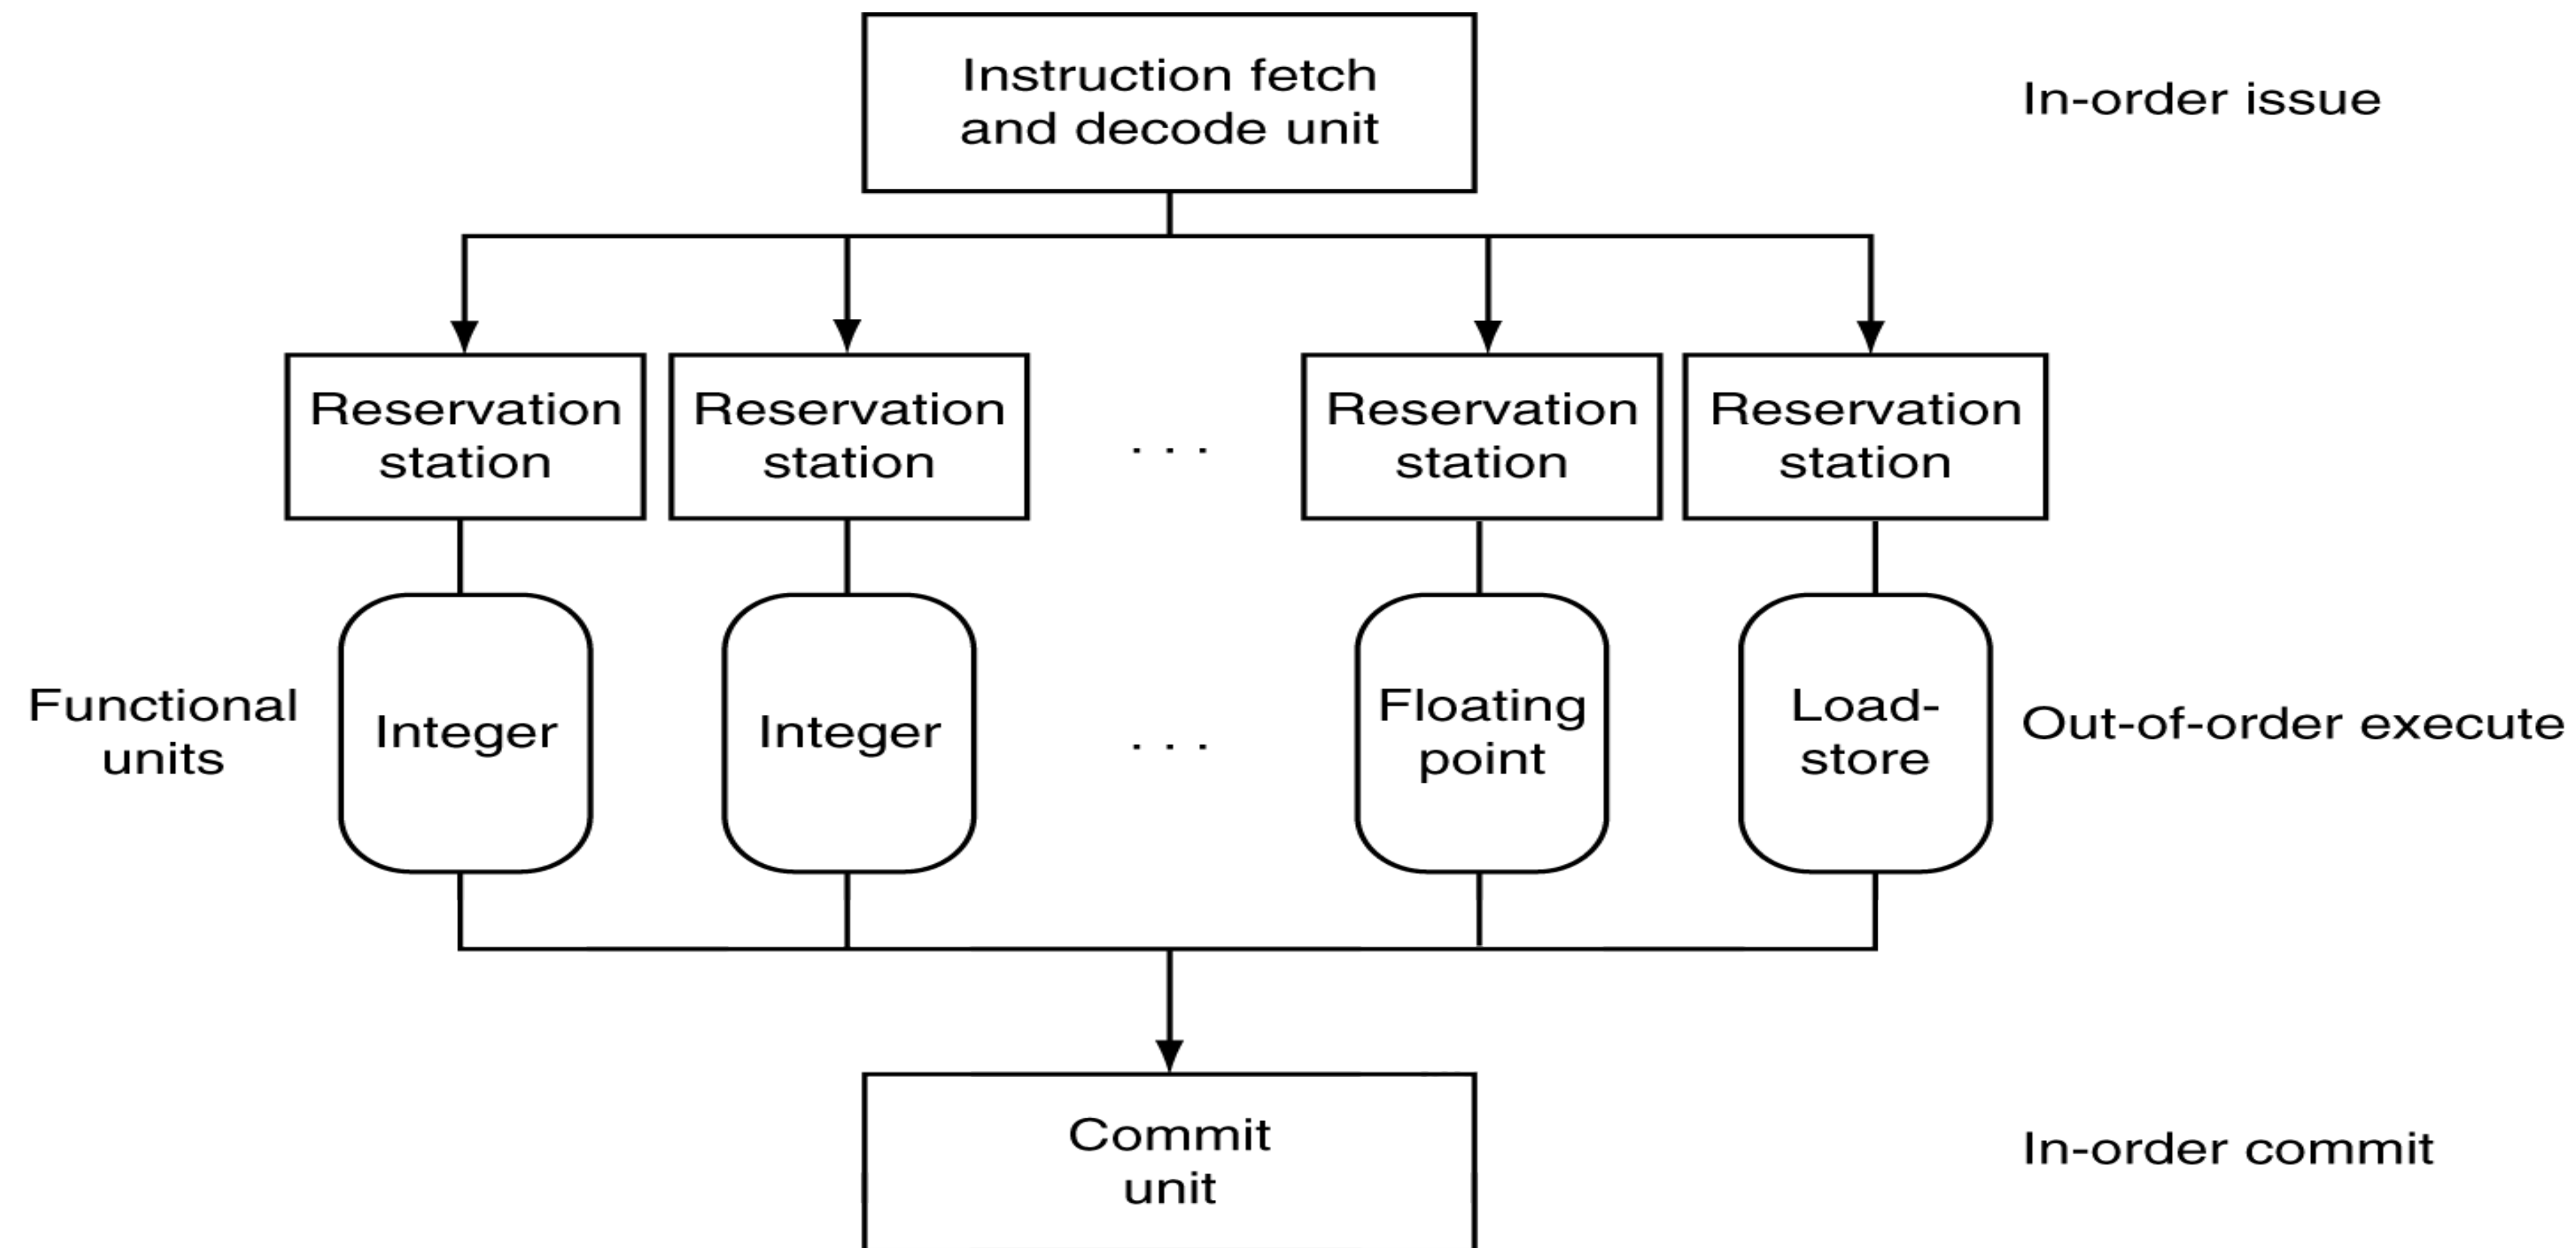
- “Overlap” instruction execution
- Deeper pipeline: 5 \Rightarrow 10 \Rightarrow 15 stages
 - Less work per stage \rightarrow shorter clock cycle
 - But more potential for hazards (CPI $>$ 1)
 - Depth limited by max clock rate and power dissipation

3. Multi-issue “superscalar” processor

Superscalar Processor

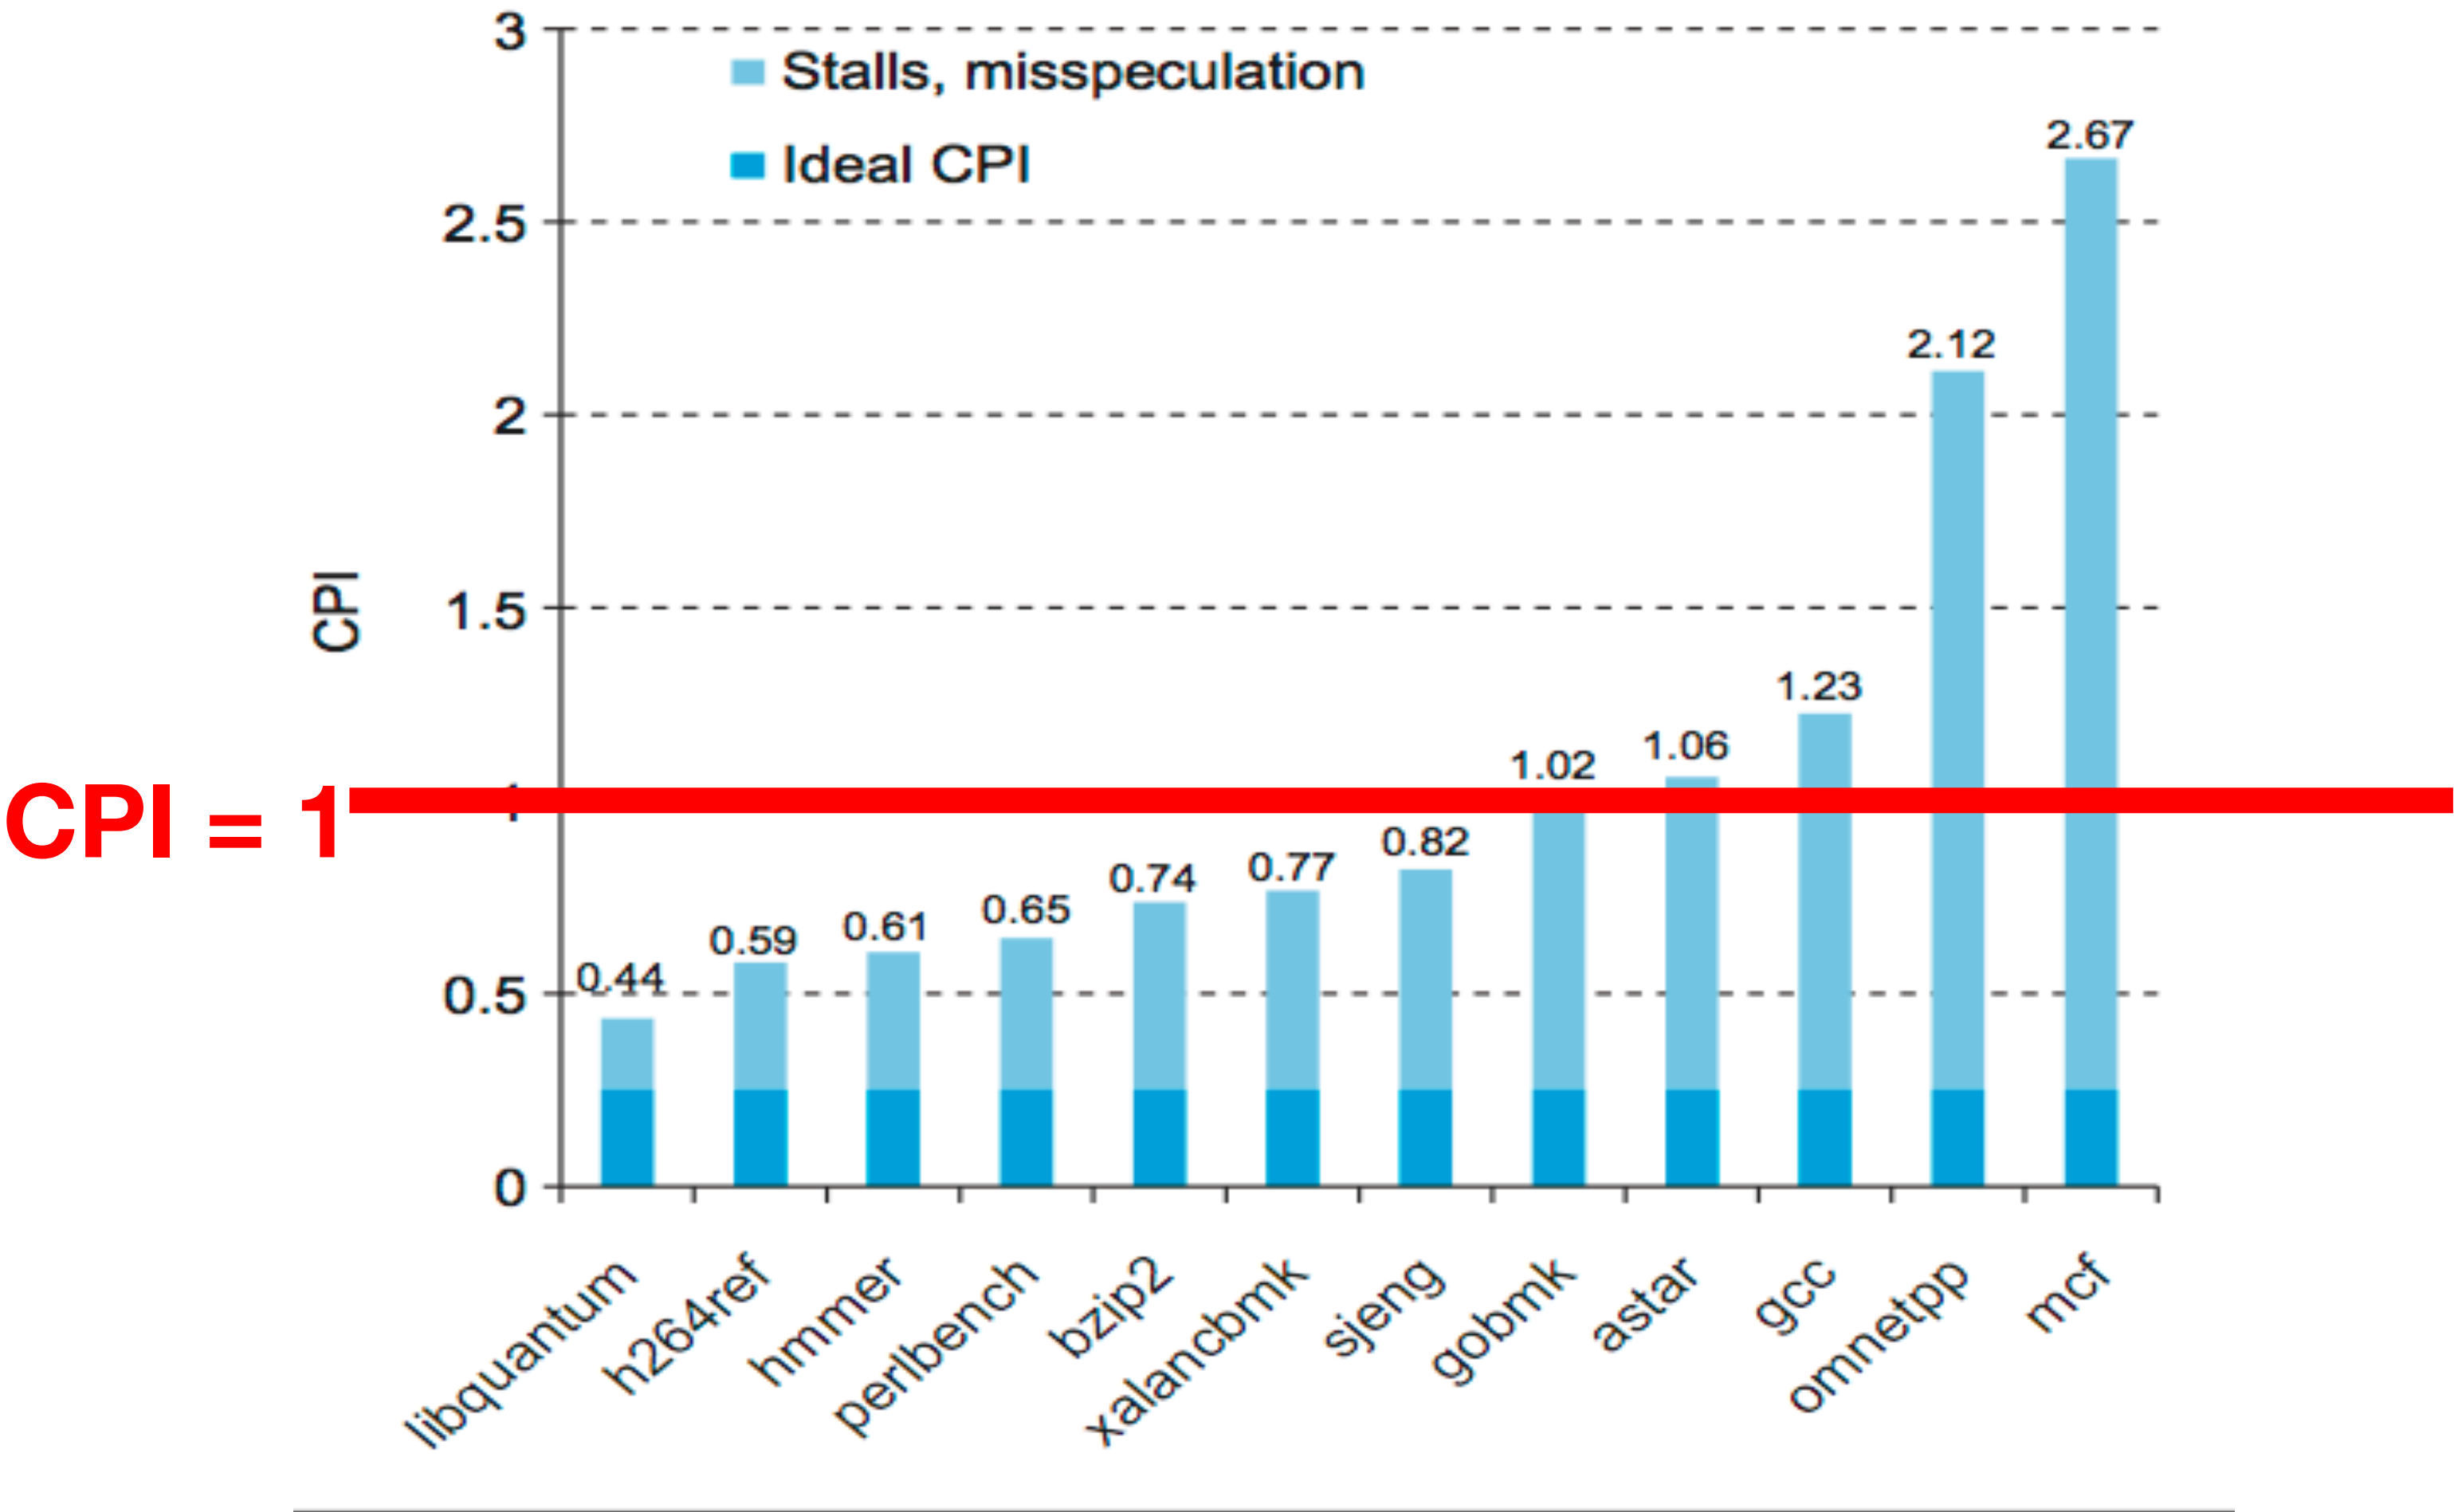
- Multiple issue “superscalar”
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - $CPI < 1$, so use Instructions Per Cycle (IPC)
 - E.g., 4GHz 4-way multiple-issue
 - 16 BIPS, peak $CPI = 0.25$, peak $IPC = 4$
 - Dependencies reduce this in practice
- “Out-of-Order” execution
 - Reorder instructions dynamically in hardware to reduce impact of hazards:
EG, memory/cache misses
- CS152 discusses these techniques!

Out Of Order Superscalar Processor



P&H p. 340

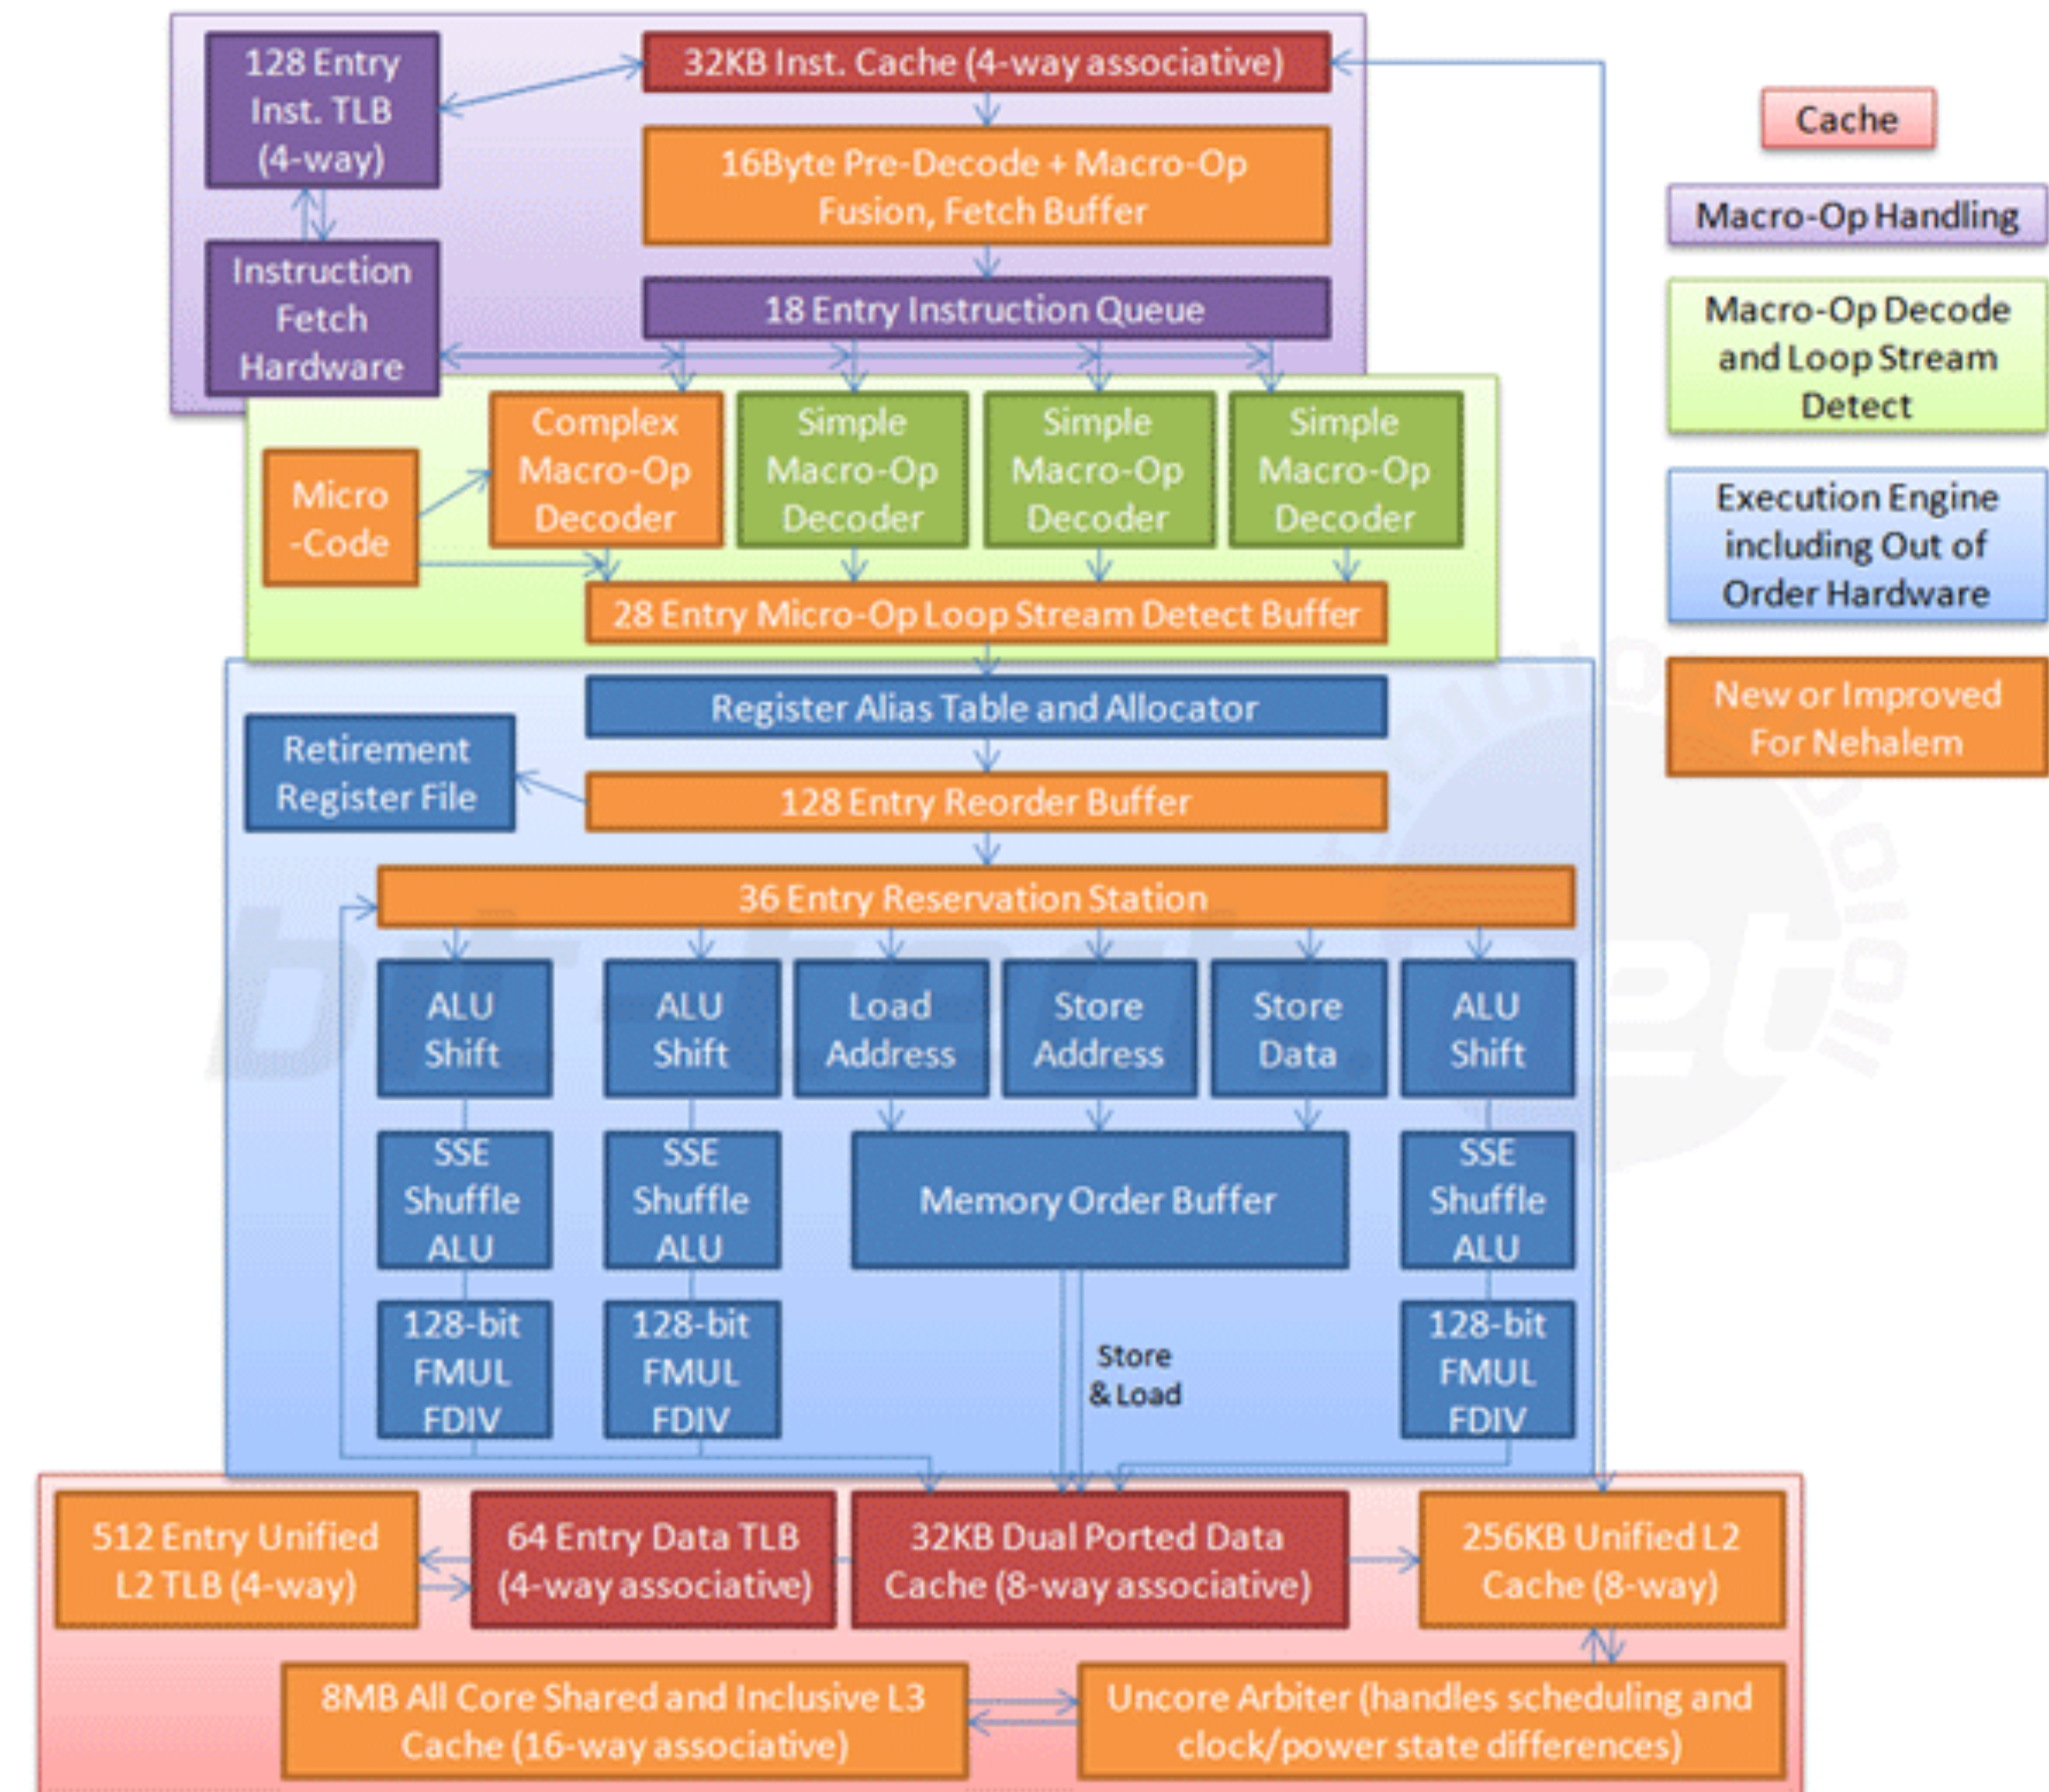
Benchmark: CPI of Intel Core i7



CPI of Intel Core i7 920 running SPEC2006 integer benchmarks.

And That Is A Beast...

- 6 separate functional units
 - 3x ALU
 - 3 for memory operations
- 20-24 stage pipeline
- Aggressive branch prediction and other optimizations
 - Massive out-of-order capability:
Can reorder up to 128 micro-operation instructions!
- And yet it still barely averages a 1 on CPI!



Pipelining and ISA Design

- RISC-V ISA designed for pipelining
 - All instructions are 32-bits in the RV-32 ISA
 - Easy to fetch and decode in one cycle
 - Variant additions add 16b and 64b instructions, but can tell by looking at just the first bytes what type it is
 - Versus x86: 1- to 15-byte instructions
 - Requires additional pipeline stages for decoding instructions
 - Few and regular instruction formats
 - Decode and read registers in one step
 - Load/store addressing
 - Calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle

In Conclusion

- Pipelining increases throughput by overlapping execution of multiple instructions
- All pipeline stages have same duration
 - Choose partition that accommodates this constraint
- Hazards potentially limit performance
 - Maximizing performance requires programmer/compiler assistance
- Superscalar processors use multiple execution units for additional instruction level parallelism
 - Performance benefit highly code dependent